



计算机视觉

**Computer Vision**

**Lecture 4: Linear Filters**

张 超

信息科学技术学院 智能科学系

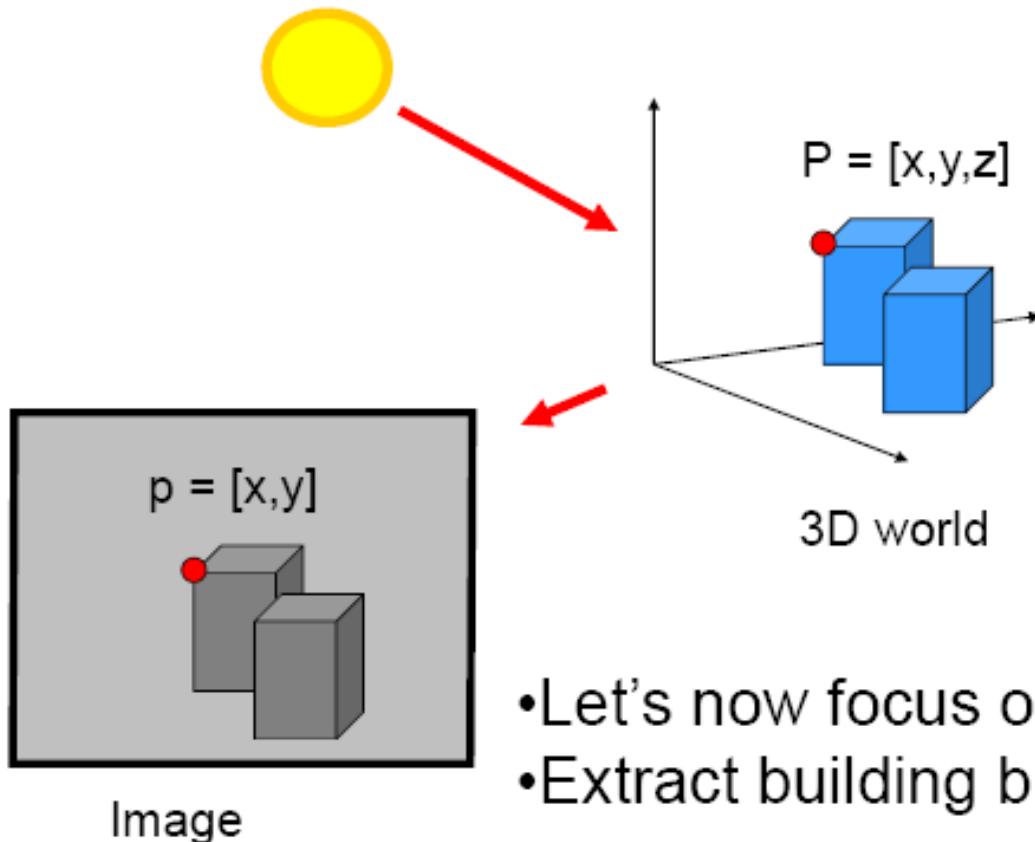
# Linear Filters

- Linear filters and convolution
- Fourier Transforms
- Multi-resolution representation

# Three views of filtering

- **Image filters in spatial domain**
  - Filter is a mathematical operation of a grid of numbers
  - Smoothing, sharpening, measuring texture
- **Image filters in the frequency domain**
  - Filtering is a way to modify the frequencies of images
  - Denoising, sampling, image compression
- **Templates and Image Pyramids**
  - Filtering is a way to match a template to the image
  - Detection, coarse-to-fine registration

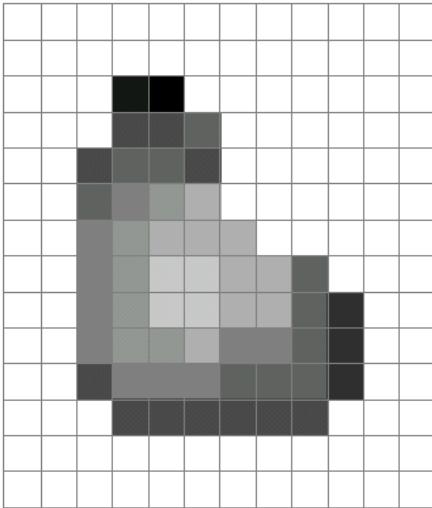
# From 3D to 2D



- Let's now focus on 2D
- Extract building blocks

# What is an image?

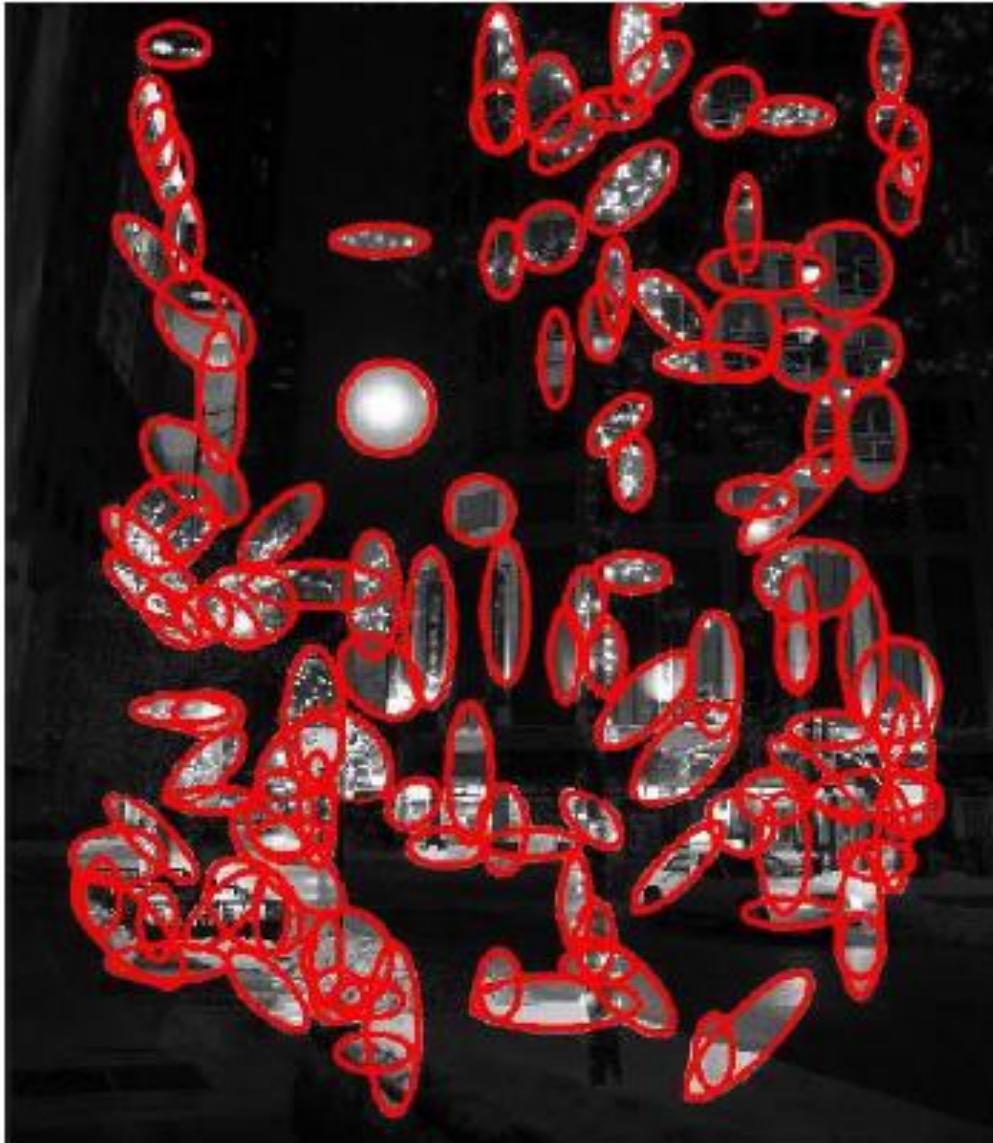
- A grid of intensity values



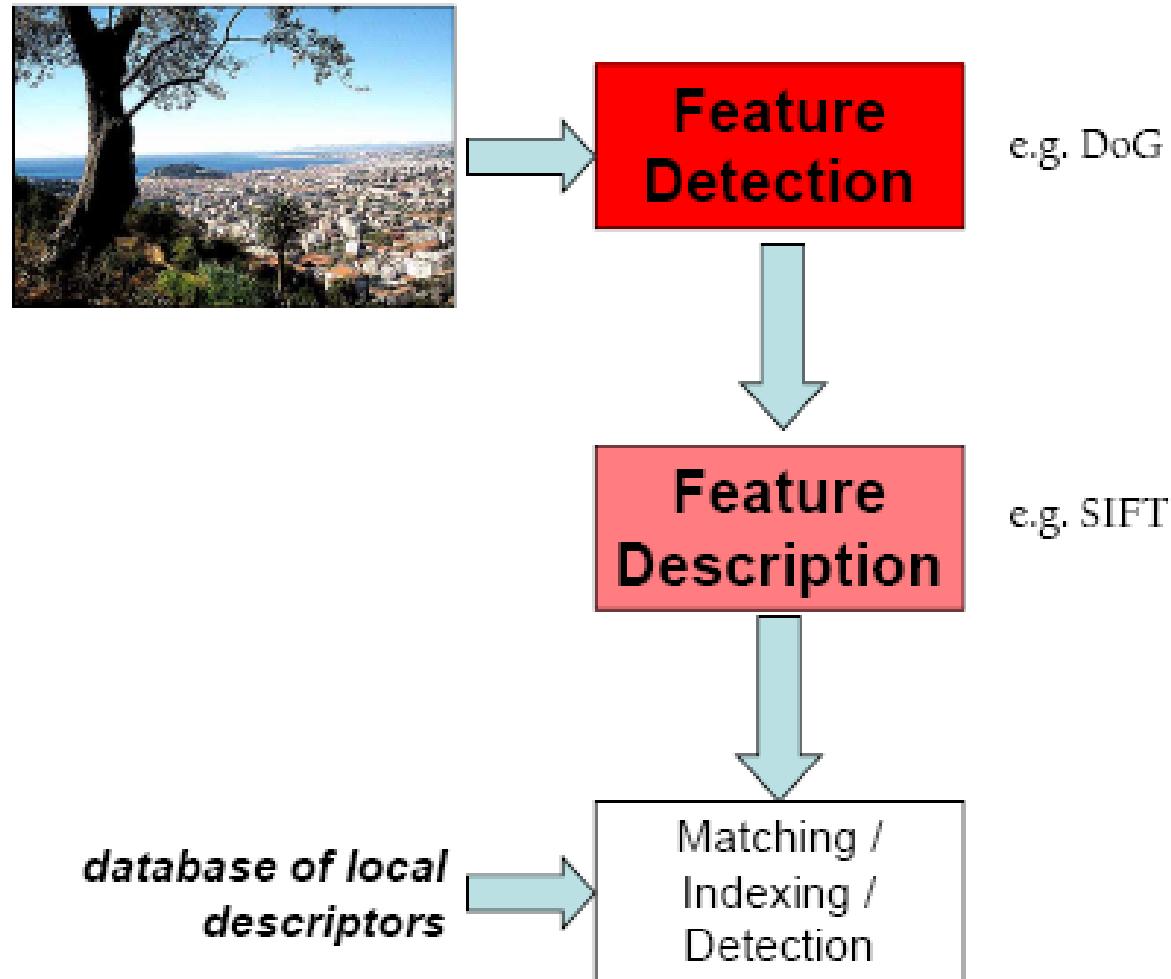
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	20	0	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255	255	255	255	255
255	255	127	145	145	175	127	127	95	95	47	255	255	255	255	255
255	255	74	127	127	127	127	95	95	95	47	255	255	255	255	255
255	255	255	74	255	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)

# Extract useful building blocks



# The big picture...

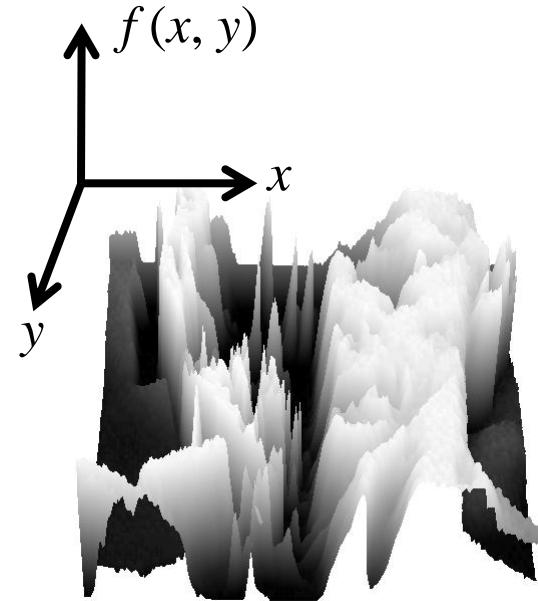


# What is an image?

- We can think of a (grayscale) image as a function,  $f$ , from  $\mathbf{R}^2$  to  $\mathbf{R}$ :
  - $f(x, y)$  gives the intensity at position  $(x, y)$



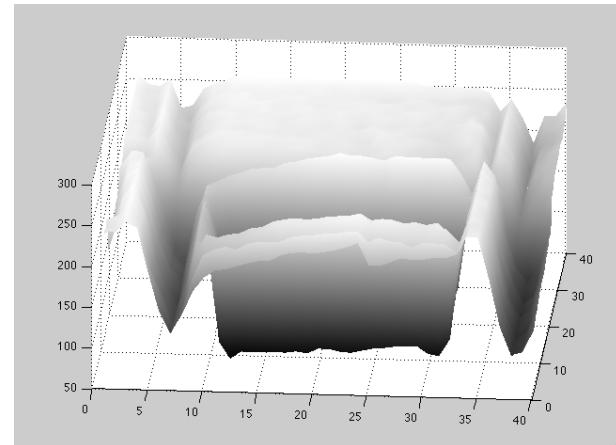
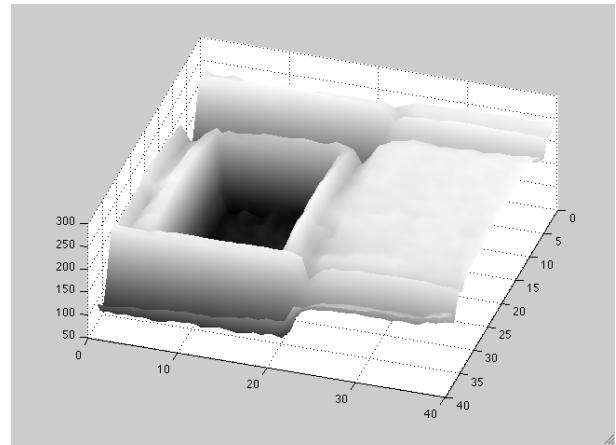
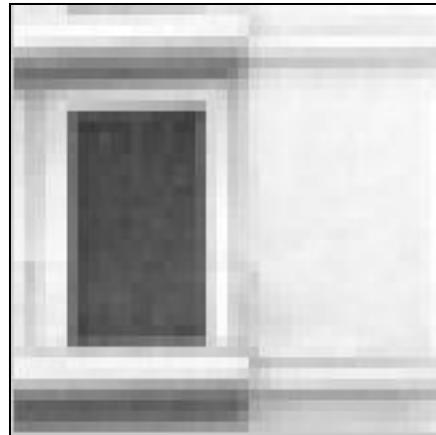
[snoop](#)



[3D view](#)

- A digital image is a discrete (sampled, quantized) version of this function

# The image as a “surface”



# Image transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, *convolution* (linear filtering)

# Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



- Take lots of images and average them!
- What's the next best thing?

Source: S. Seitz

# Linear filters and convolution

# Convolution

- **Each pixel in output image is**
  - weighted average of window of pixels in input image
  - weights stay the same
  - window centered on pixel
- **Important operation**
  - Example: smoothing by averaging
  - Example: smoothing by weighted average
  - Example: taking a derivative

# Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

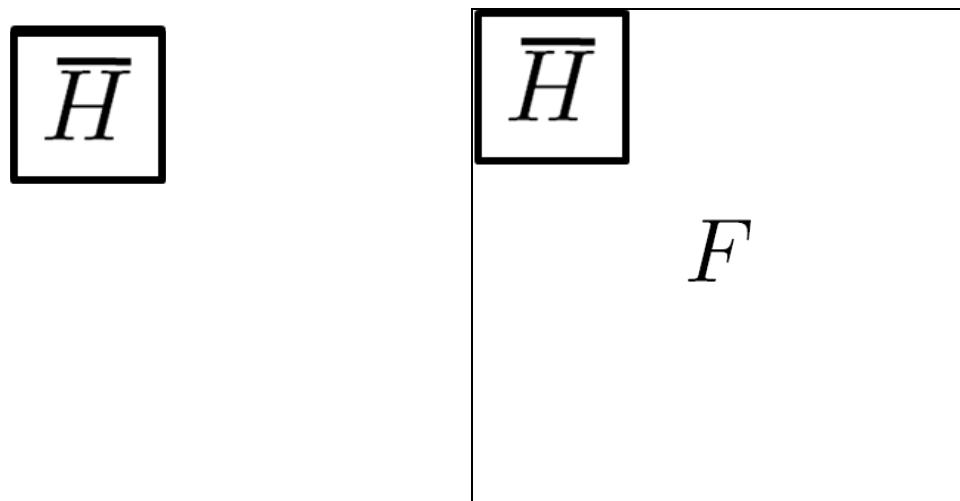
$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

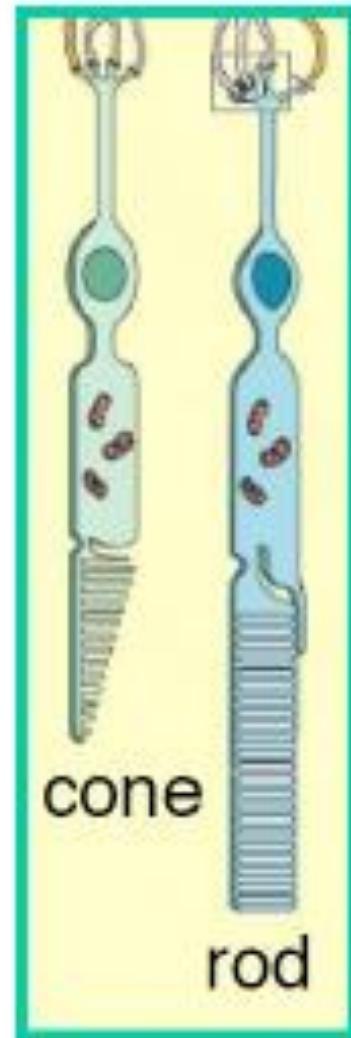
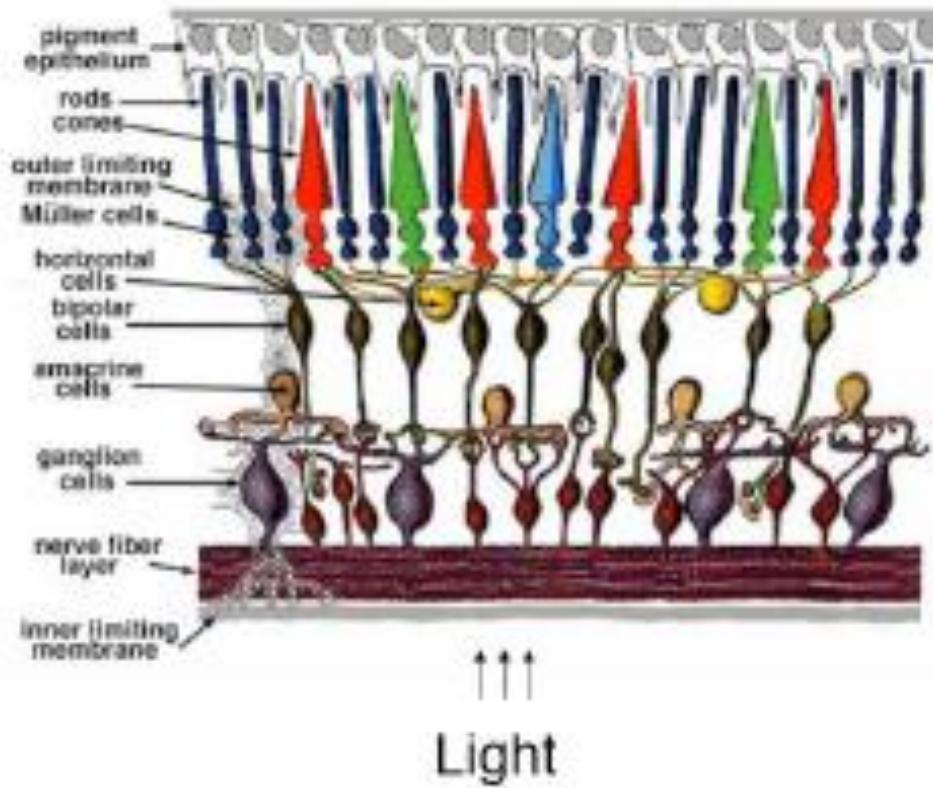
- MATLAB functions: `conv2`, `filter2`, `imfilter`

# Convolution

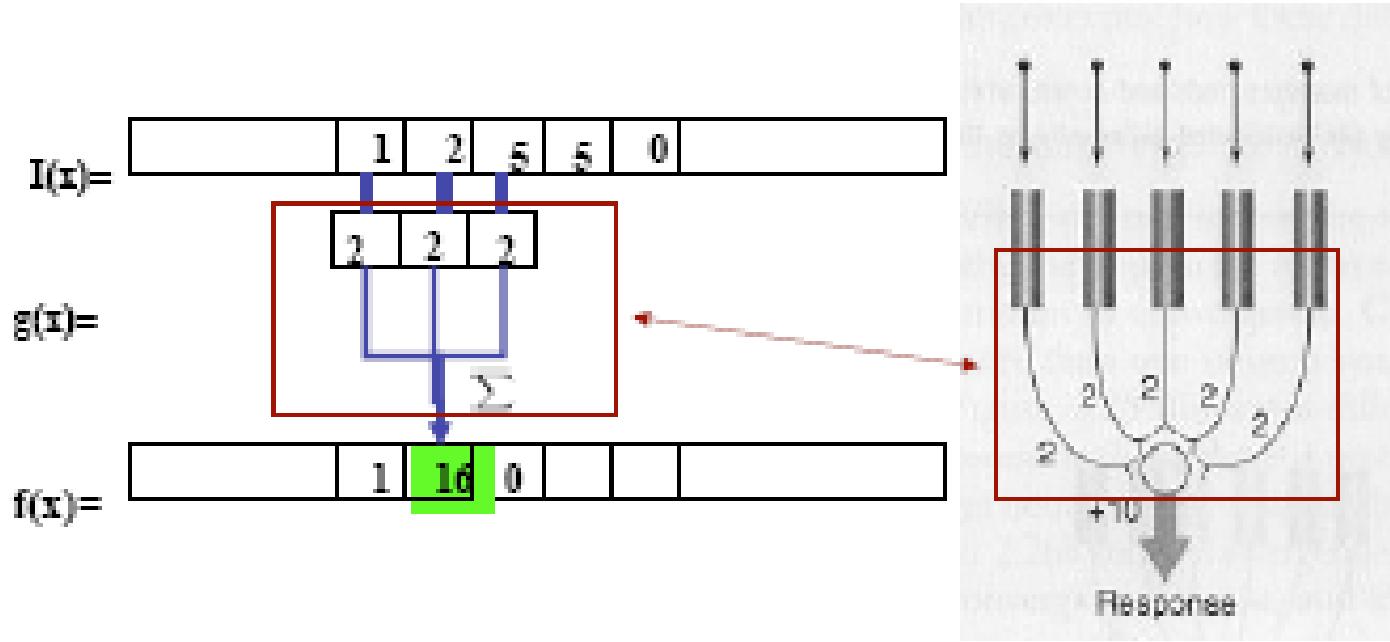


# Retina up-close

---

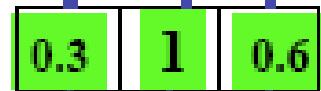


# Simple Neural Network

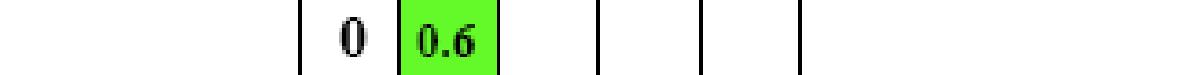


This leads to parallel computation

$I(x) =$  



$g(x) =$

$f(x) =$  

$f(x) =$  

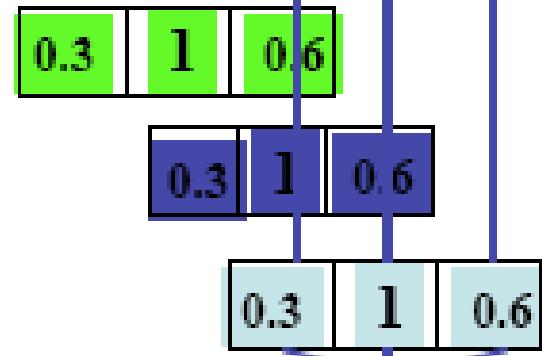


$g(x) =$  

$f(x) =$

	0	0	1	0	0	
--	---	---	---	---	---	--

$h(x) =$



$g(x) =$

	0	0.6	1	0.3		
--	---	-----	---	-----	--	--

# Filters

- Convolution is the same as applying a filter to an image
  - Weights are called the “filter kernel”
  - Notice I can re-index the filter kernel so it looks like a vector
  - If I do so, each output pixel is
    - the dot-product of that vector with
    - the vector of pixels centered at the output pixel
- Important insight
  - Filters respond strongly to image windows that look like them
    - positive response: just like  $H$
    - negative response: like  $H$ , but with contrast reversal

$$R_{ij} = \sum_{u,v} H_{i-u,j-v} F_{u,v}$$

↑  
Filter kernel

# Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function

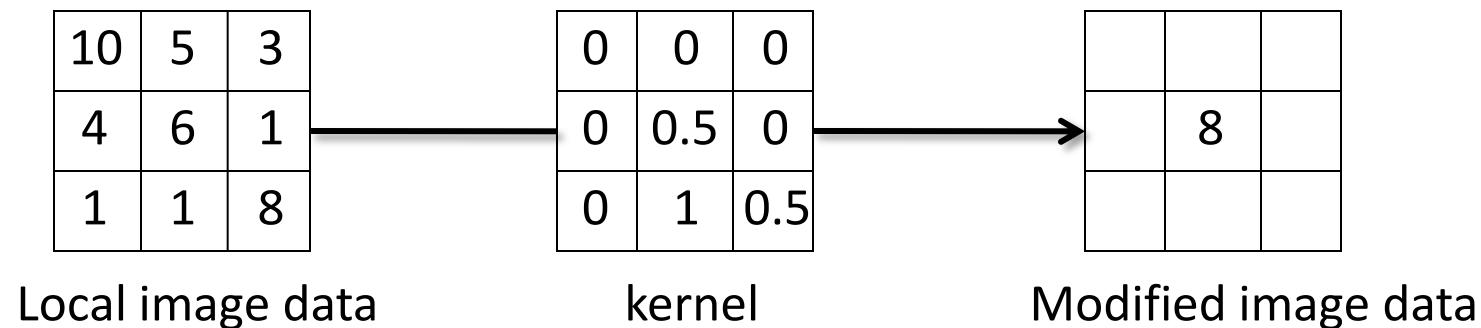


	7	

Modified image data

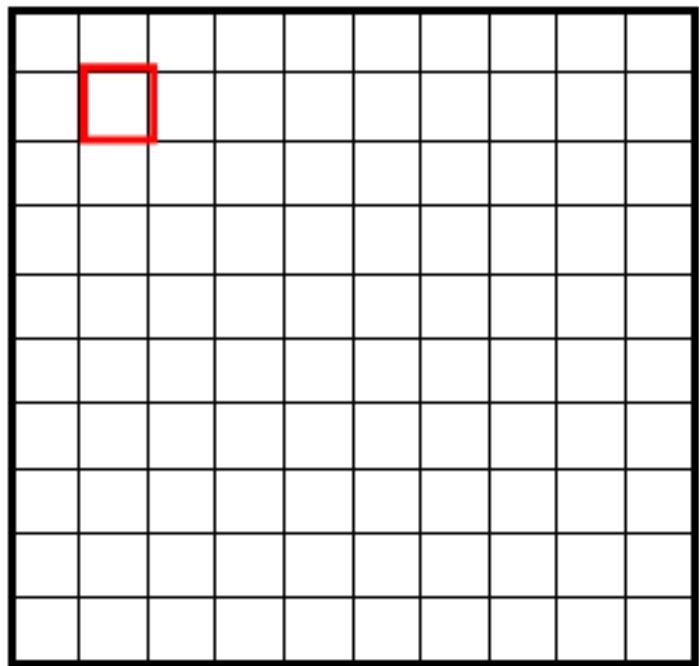
# Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	90	0	90	90	90	0
0	0	0	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0



$\frac{1}{9}$	1	1	1
	1	1	1
	1	1	1

$$\frac{1}{9}$$

$$\frac{1}{9}$$

$$\frac{1}{9}$$

$\frac{1}{9}$	1	1	1
	1	1	1
	1	1	1

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	10	20	30	30	30	30	20	10	
0	20	40	60	60	60	60	40	20	
0	30	60	90	90	90	90	60	30	
0	30	50	80	80	80	90	60	30	
0	30	50	80	80	80	90	60	30	
0	20	30	50	50	50	60	40	20	
10	20	30	30	30	30	30	20	10	
10	10	10	0	0	0	0	0	0	

## What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

# Key properties

- Linearity:  $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$
- Shift invariance: same behavior regardless of pixel location:  $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$
- Theoretical result: any linear shift-invariant operator can be represented as a convolution

# Properties in more detail

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$
- Identity: unit impulse  $e = [..., 0, 0, 1, 0, 0, ...]$ ,  
 $a * e = a$

# Linear filters: examples



\*

0	0	0
0	1	0
0	0	0

=



Original

Identical image

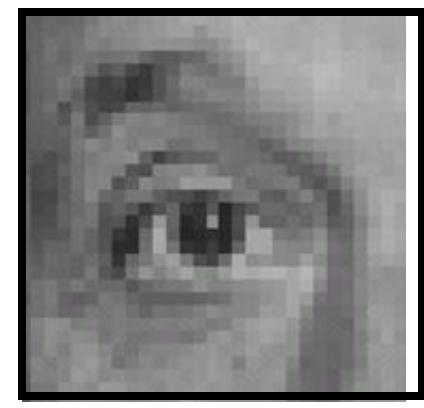
# Linear filters: examples



\*

0	0	0
1	0	0
0	0	0

=



Original

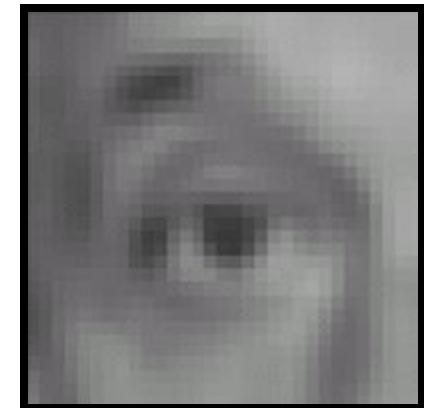
Shifted left  
By 1 pixel

# Linear filters: examples



Original

$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



Blur (with a mean filter)

# Linear filters: examples

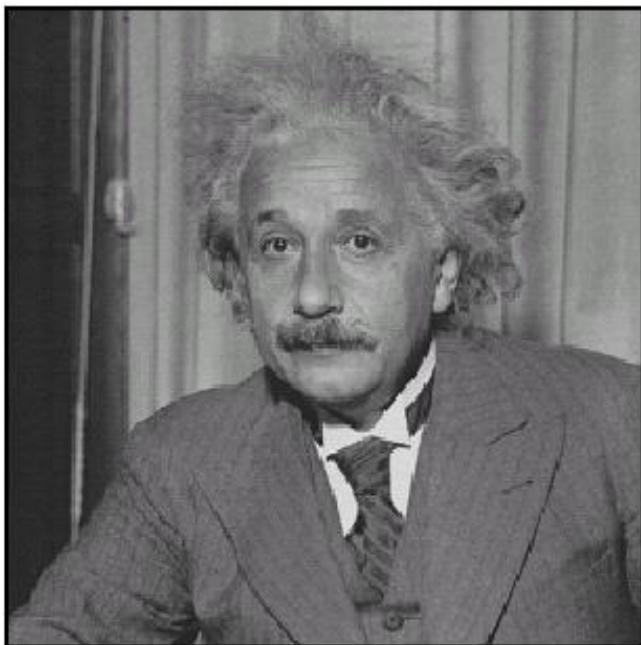


$$\text{Original} \quad * \left( \begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} - \frac{1}{9} \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right) =$$

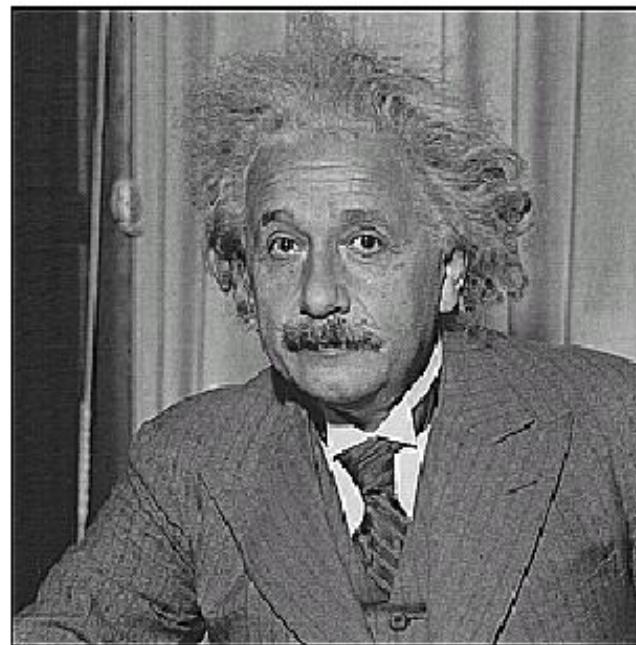


**Sharpening filter**  
(accentuates edges)

# Sharpening



**before**



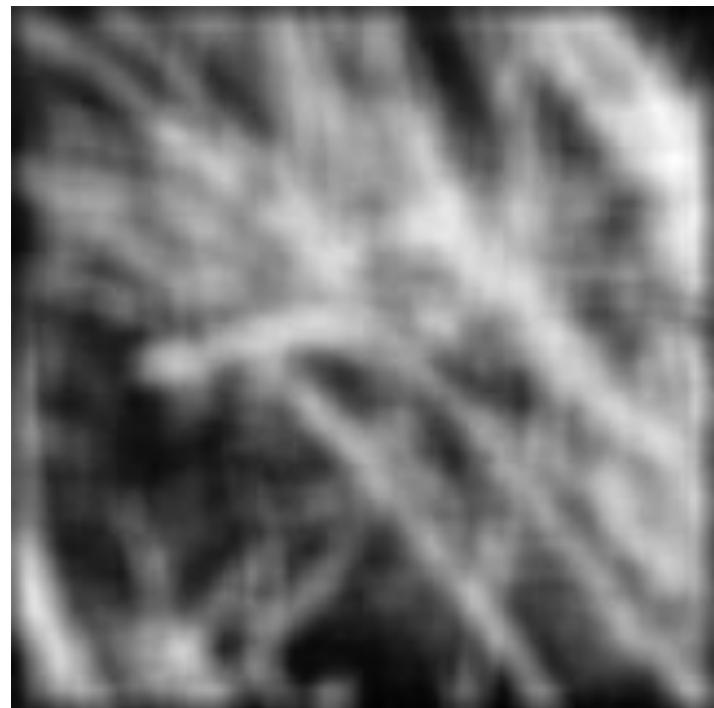
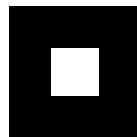
**after**

# Average Filter

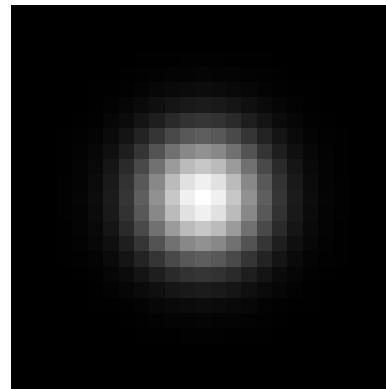
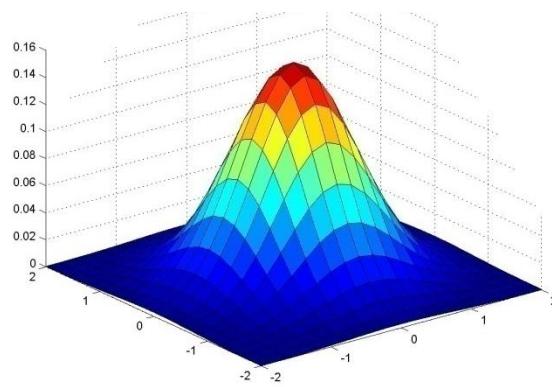
- Mask with positive entries, that sum 1.
- Replaces each pixel with an average of its neighborhood.
- If all weights are equal, it is called a BOX filter.

$$F = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

# Smoothing with box filter revisited

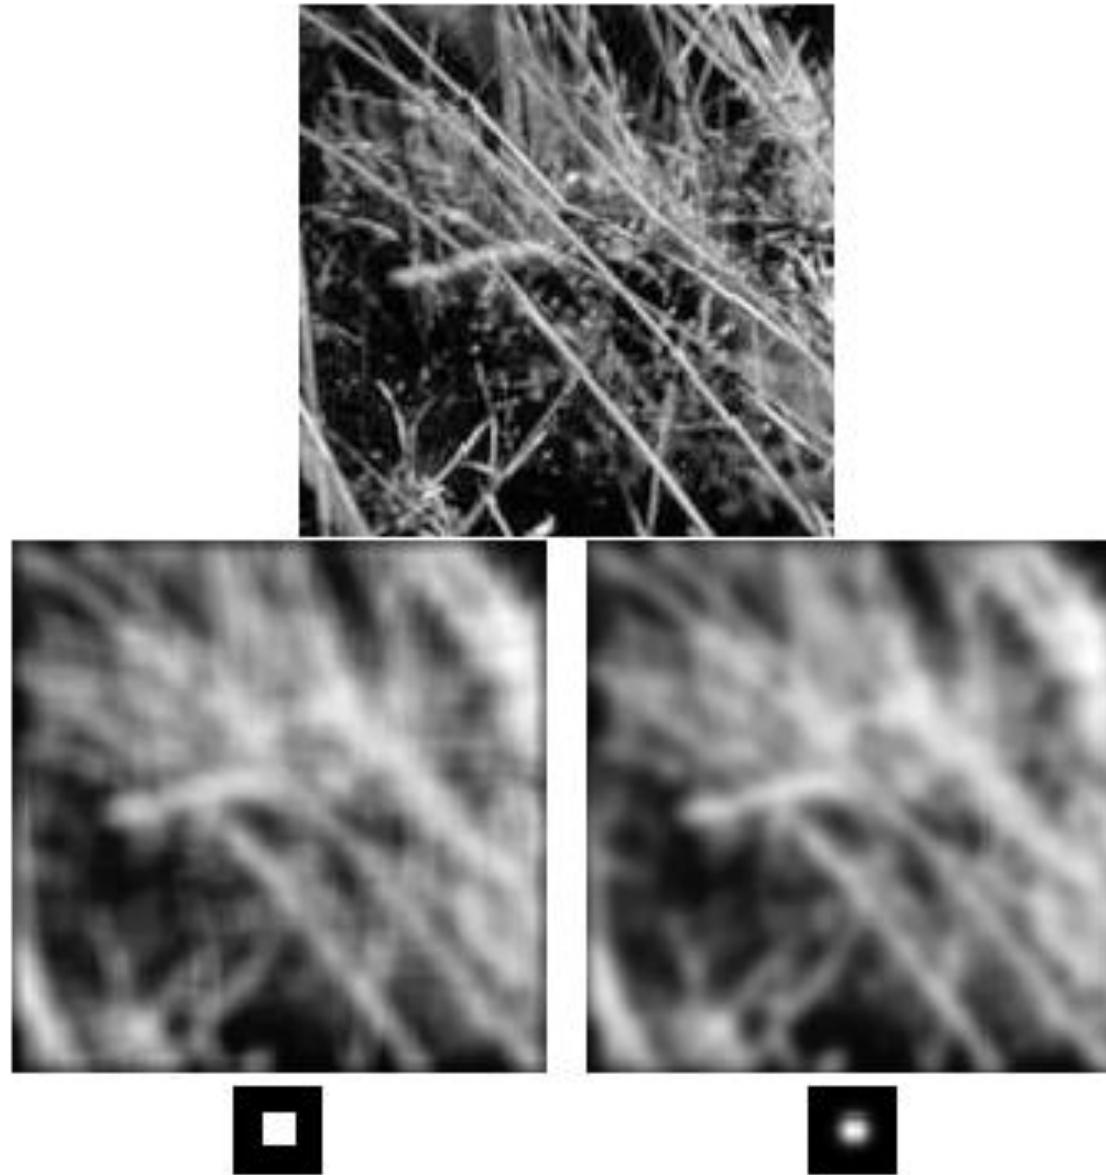


# Gaussian Kernel



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Mean vs. Gaussian filtering



# Gaussian filter

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian



- Convolving two times with Gaussian kernel of width  $\sigma$  = convolving once with kernel of width  $\sigma\sqrt{2}$

# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Separability example

2D convolution  
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix} \cdots$$

Perform convolution  
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 & & \\ 18 & & \\ 18 & & \end{matrix}$$

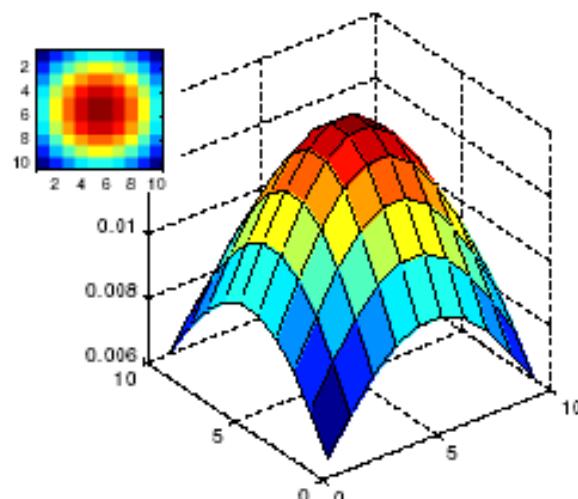
Followed by convolution  
along the remaining column:

# Why is separability useful?

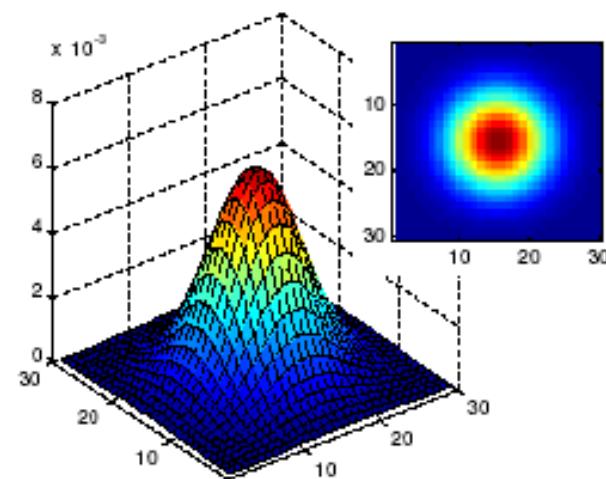
- What is the complexity of filtering an  $n \times n$  image with an  $m \times m$  kernel?
  - $O(n^2 m^2)$
- What if the kernel is separable?
  - $O(n^2 m)$

# Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
  - Note, Gaussian function has infinite support, but discrete filters use finite kernels



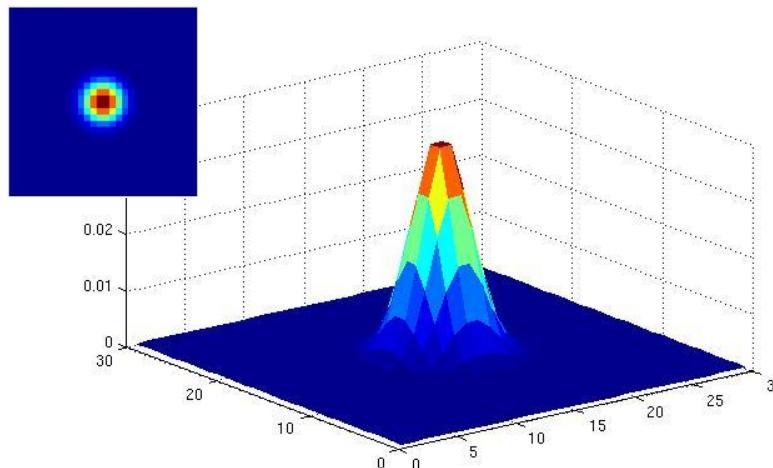
$\sigma = 5$  with  
10 x 10  
kernel



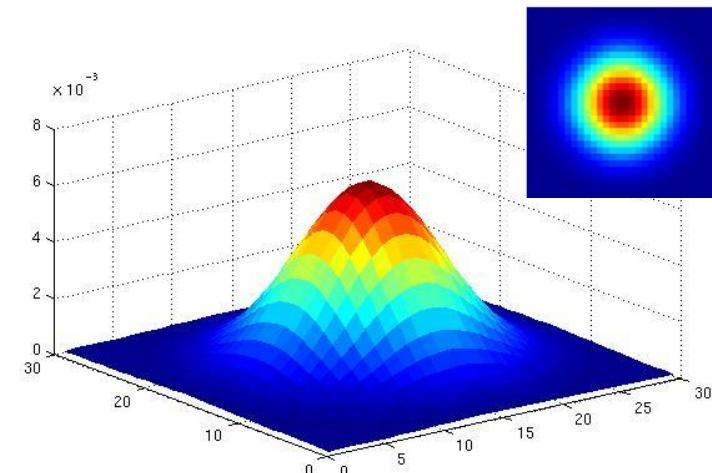
$\sigma = 5$  with  
30 x 30  
kernel

# Gaussian filters

- What parameters matter here?
- **Variance of Gaussian:** determines extent of smoothing



$\sigma = 2$  with  
30 x 30  
kernel

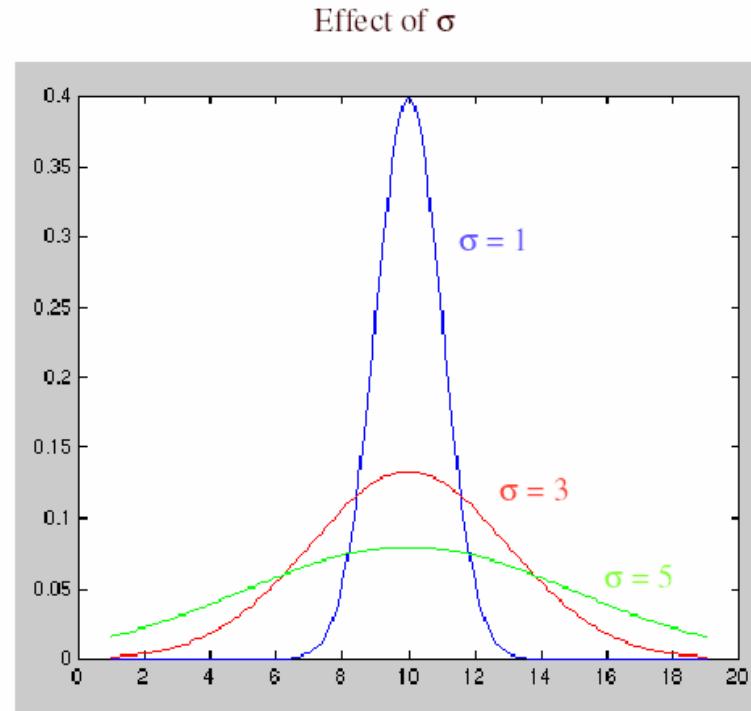


$\sigma = 5$  with  
30 x 30  
kernel

# Gaussian filters

How big should the filter be?

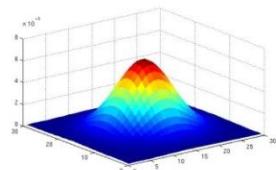
- Values at edges should be near zero ← important!
- Rule of thumb for Gaussian: set filter half-width to about  $3 \sigma$



# Matlab

```
>> hsize = 10;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```

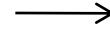
```
>> mesh(h);
```



```
>> imagesc(h);
```



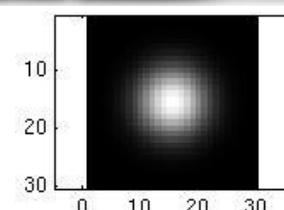
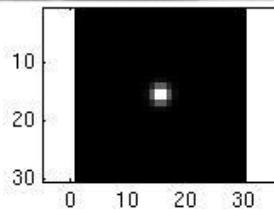
```
>> outim = imfilter(im, h); % correlation  
>> imshow(outim);
```



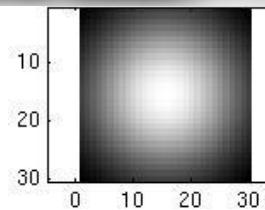
**outim**

# Smoothing with a Gaussian

Parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...



```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

# Noise



Original



Salt and pepper noise



Impulse noise

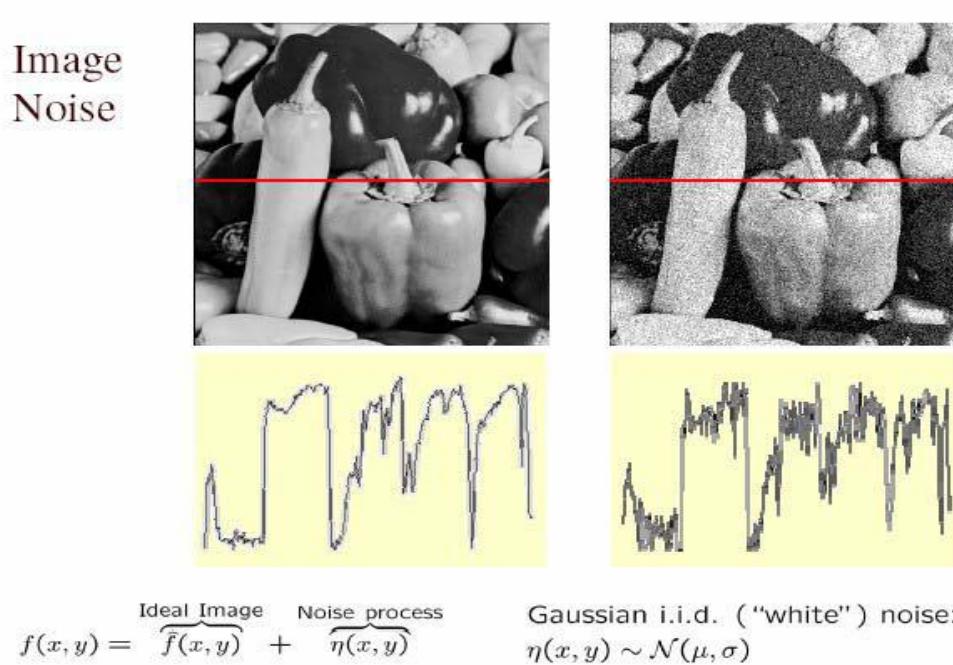


Gaussian noise

- **Salt and pepper noise:** contains random occurrences of black and white pixels
- **Impulse noise:** contains random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution

# Gaussian noise

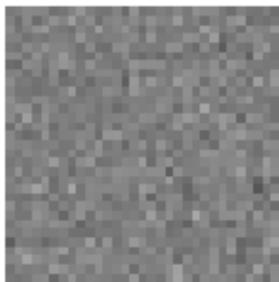
- Mathematical model: sum of many independent factors
- Good for small standard deviations
- Assumption: independent, zero-mean noise



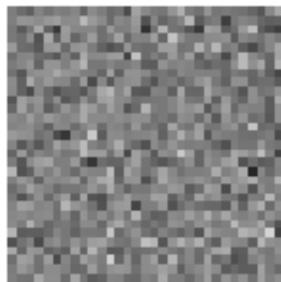
Source: M. Hebert

# Reducing Gaussian noise

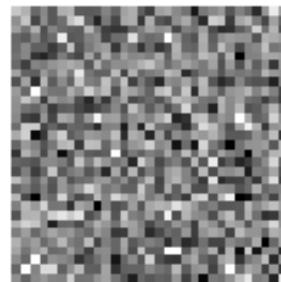
$\sigma=0.05$



$\sigma=0.1$



$\sigma=0.2$



no  
smoothing



$\sigma=1$  pixel



$\sigma=2$  pixels



Smoothing with larger standard deviations suppresses noise, but also blurs the image

# Reducing salt-and-pepper noise

3x3



5x5



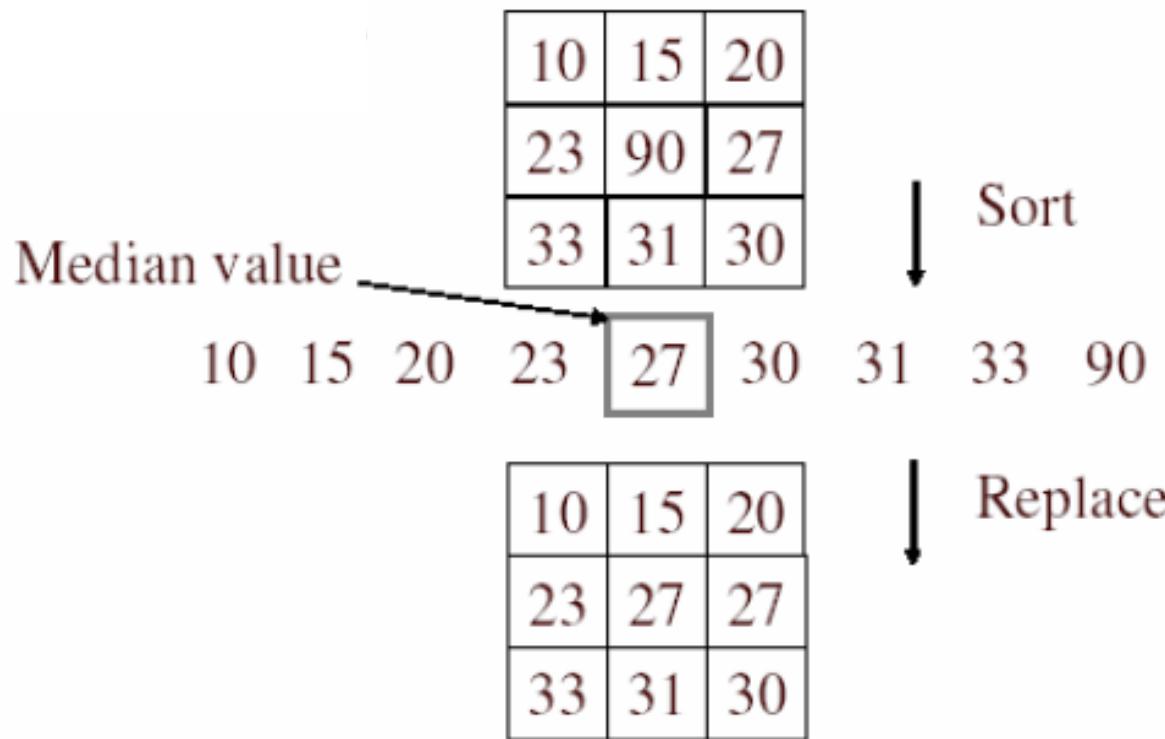
7x7



- What's wrong with the results?

# Alternative idea: Median filtering

- A median filter operates over a window by selecting the median intensity in the window

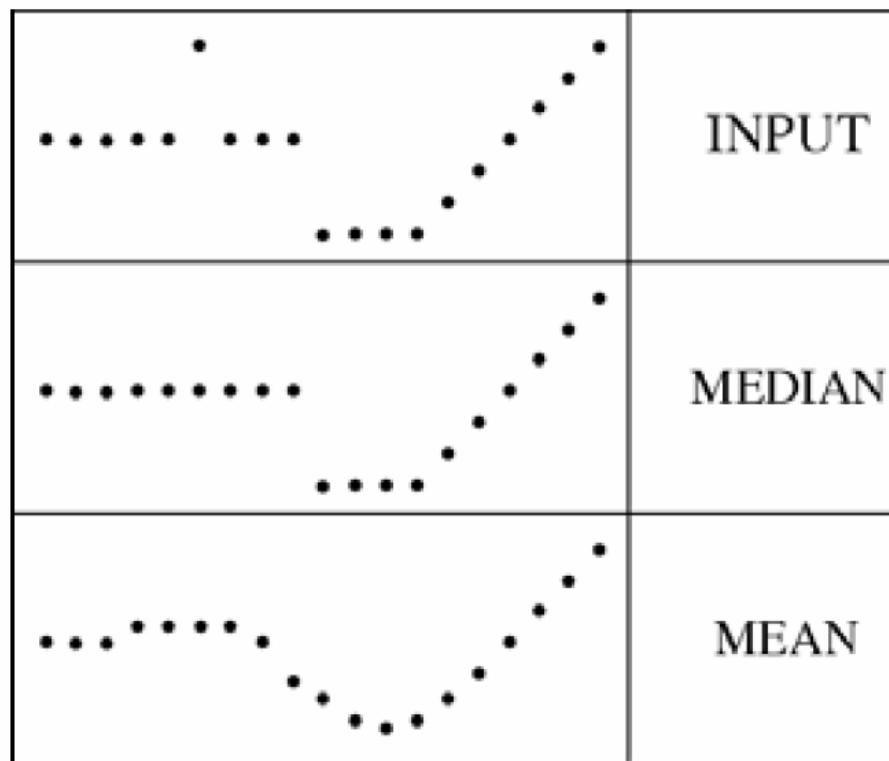


- Is median filtering linear?

# Median filter

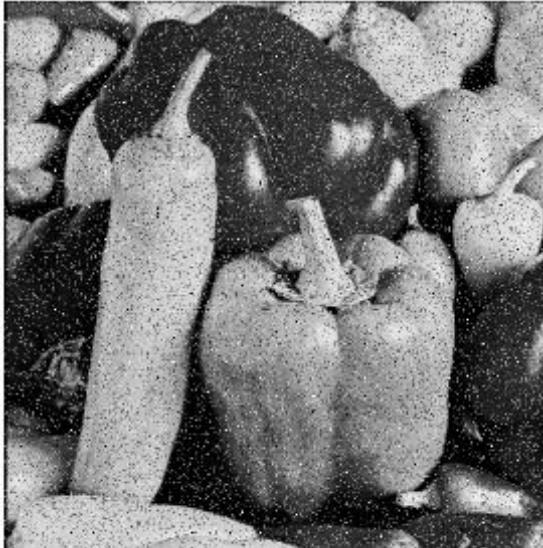
- What advantage does median filtering have over Gaussian filtering?
  - Robustness to outliers

filters have width 5 :

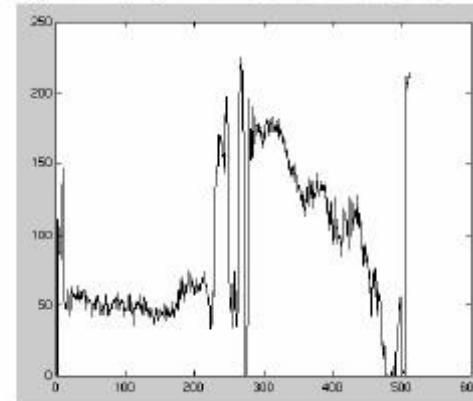
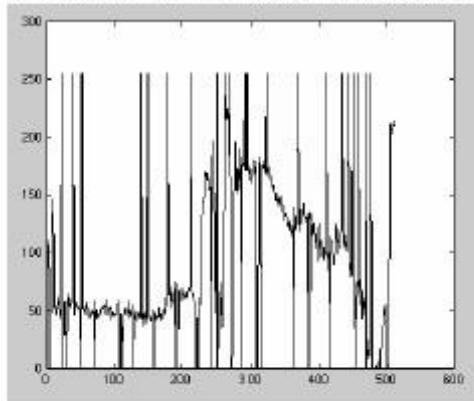


# Median filter

Salt-and-pepper noise



Median filtered



- MATLAB: [medfilt2\(image, \[h w\]\)](#)

# Gaussian vs. median filtering

3x3



5x5



7x7

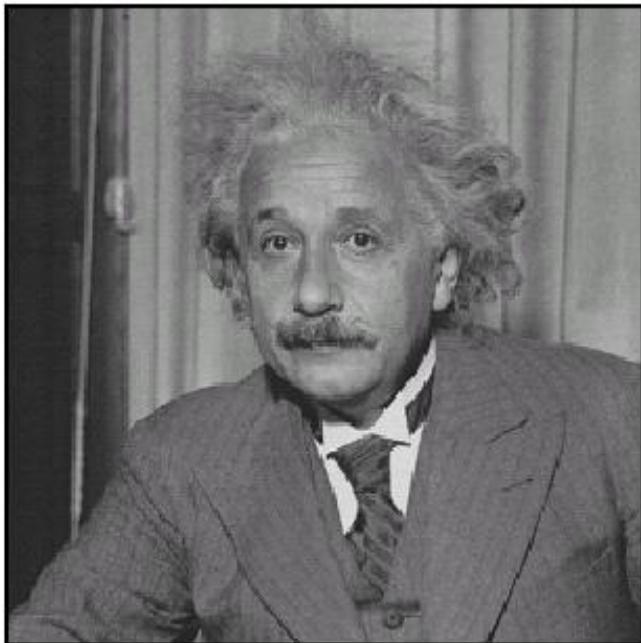


Gaussian

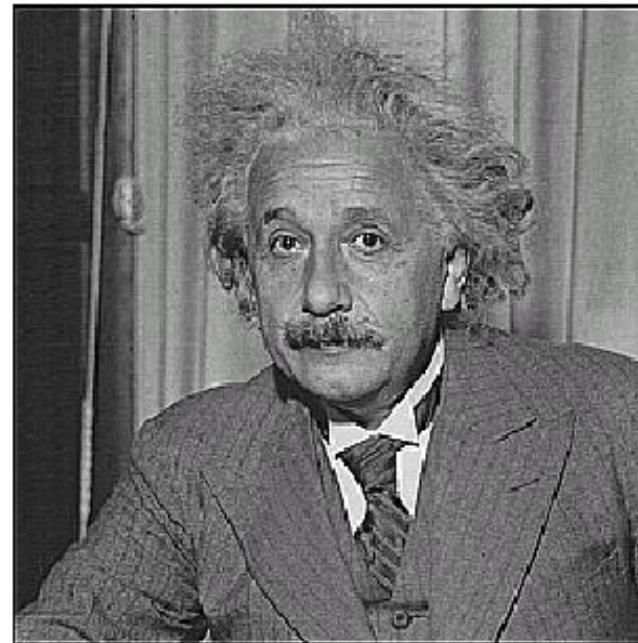
Median



# Sharpening revisited



**before**



**after**

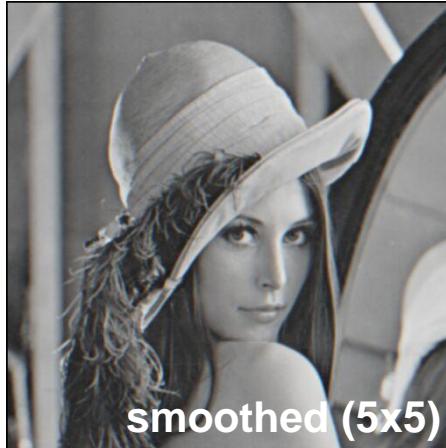
# Sharpening revisited

- What does blurring take away?



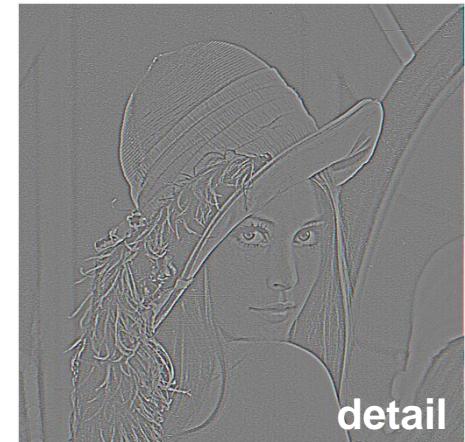
original

-



smoothed (5x5)

=



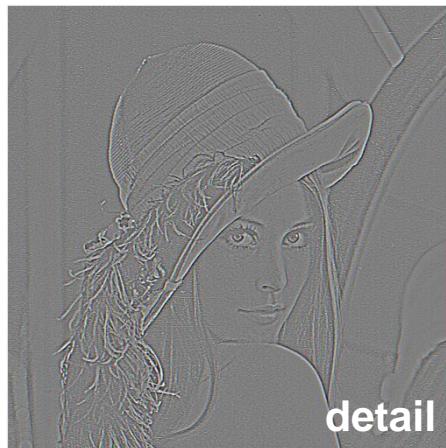
detail

Let's add it back:



original

+  $\alpha$



detail

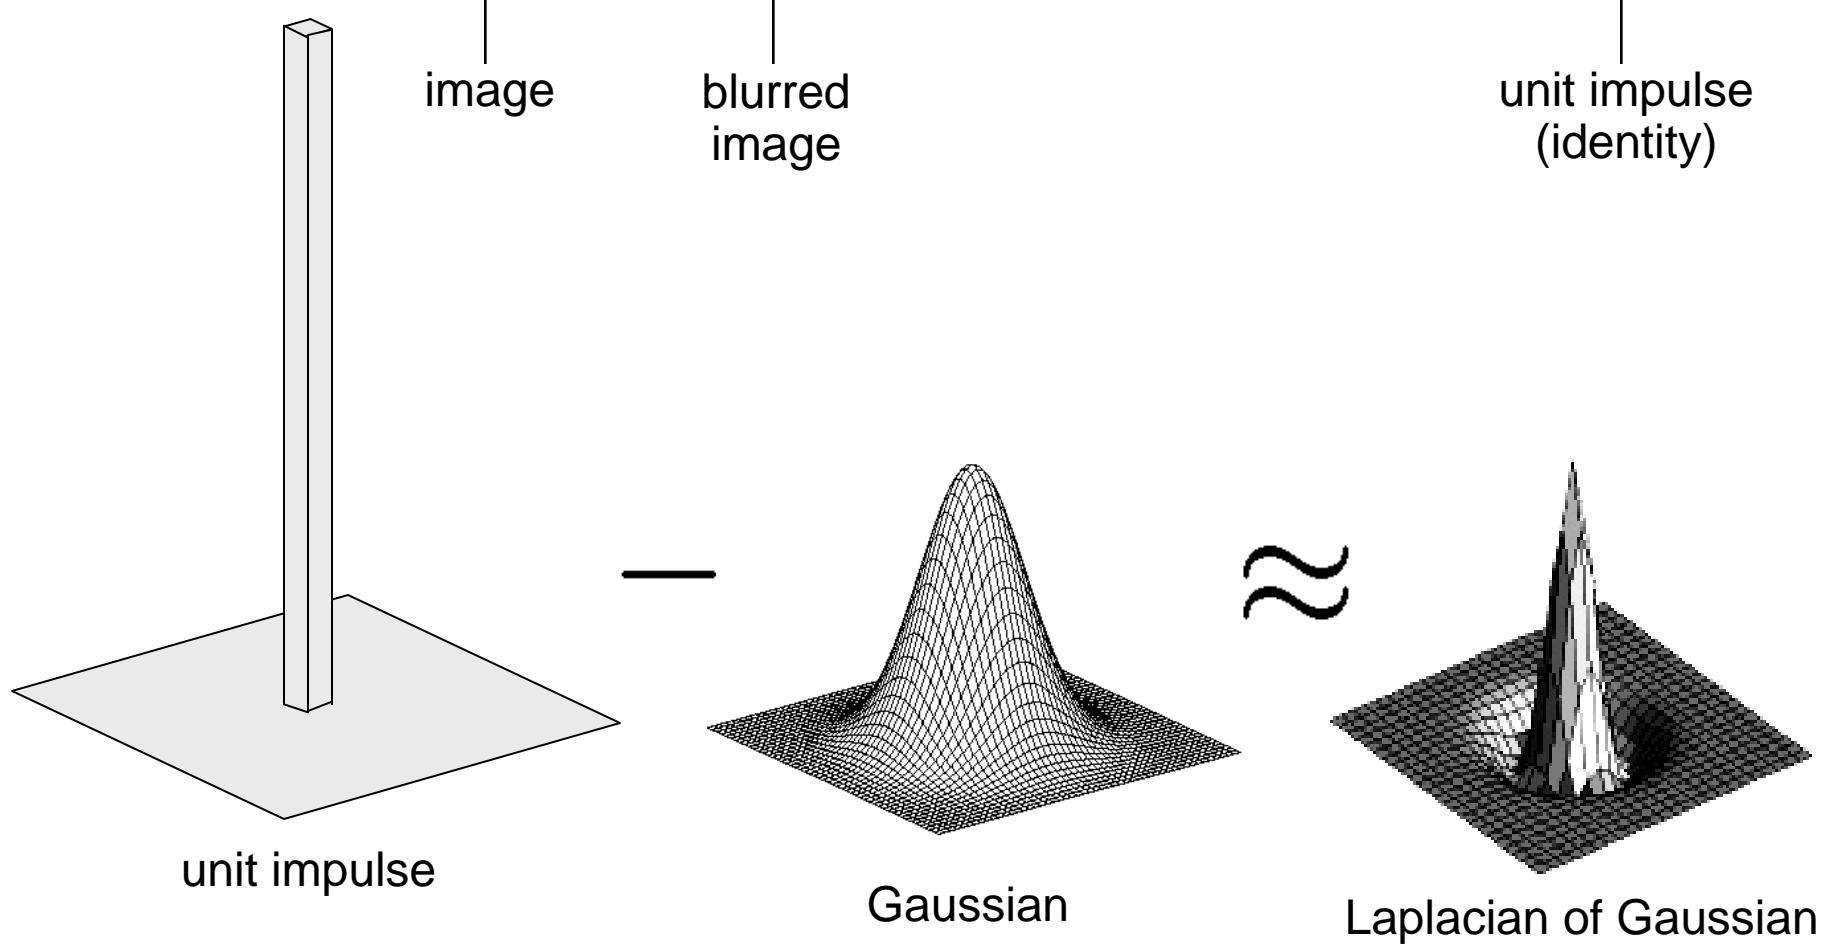
=



sharpened

# Unsharp mask filter

$$f + \alpha(f - f * g) = (1 + \alpha)f - \alpha f * g = f * ((1 + \alpha)e - \alpha g)$$



# Sharpen filter



# Finding Simple Patterns with Filters

- A filter is a pattern detector
  - strong filter responses suggest the image pattern “looks like” the filter
- It's not a very good one, because
  - it's linear, so brighter image patches give stronger filter responses
  - it can respond to constant image patches
    - if the mean of the filter weights isn't zero

# Normalized Correlation

- **Strategy (Normalized correlation):**
  - use a zero mean filter (subtract the average of the filter weights)
  - for each image patch, subtract the mean
  - divide the filter response by the magnitude of the result
  - **Output**
    - value close to 1 --- image pattern like the filter kernel
    - value close to -1 --- image pattern like contrast-reversed filter kernel
- **This is a great way to build simple, fast pattern detectors**
  - that require little computation

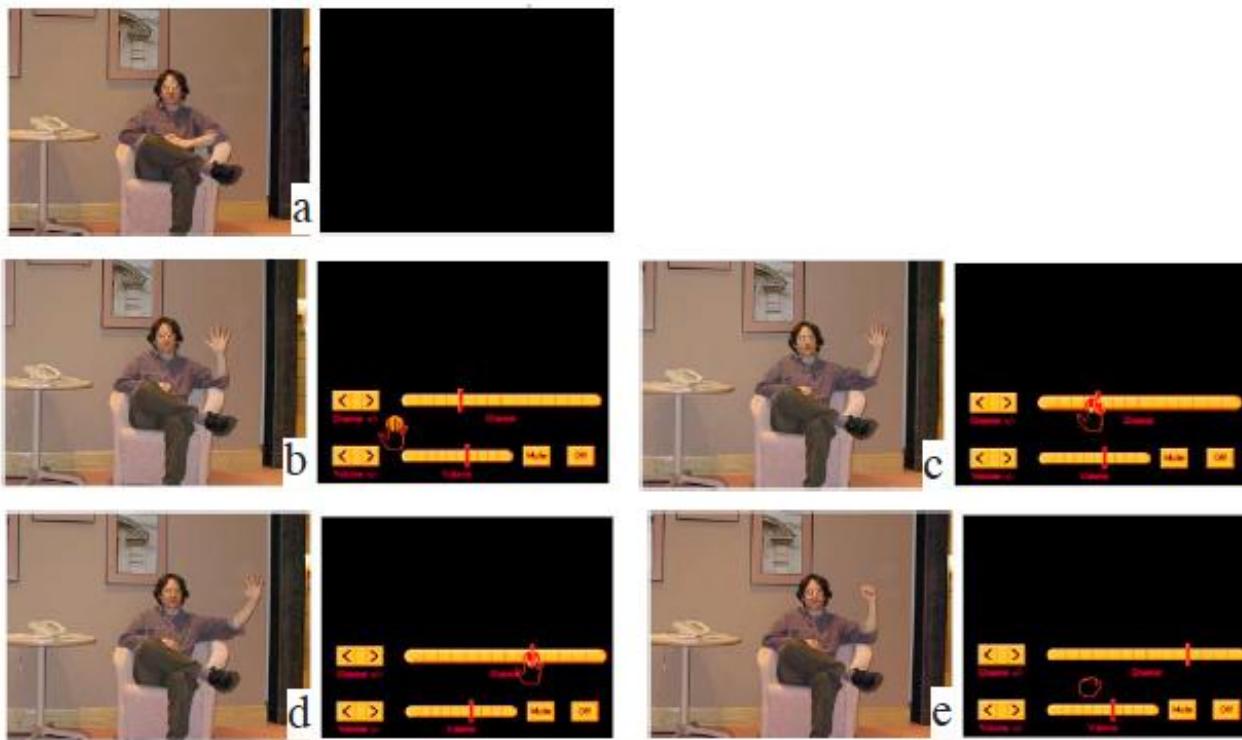
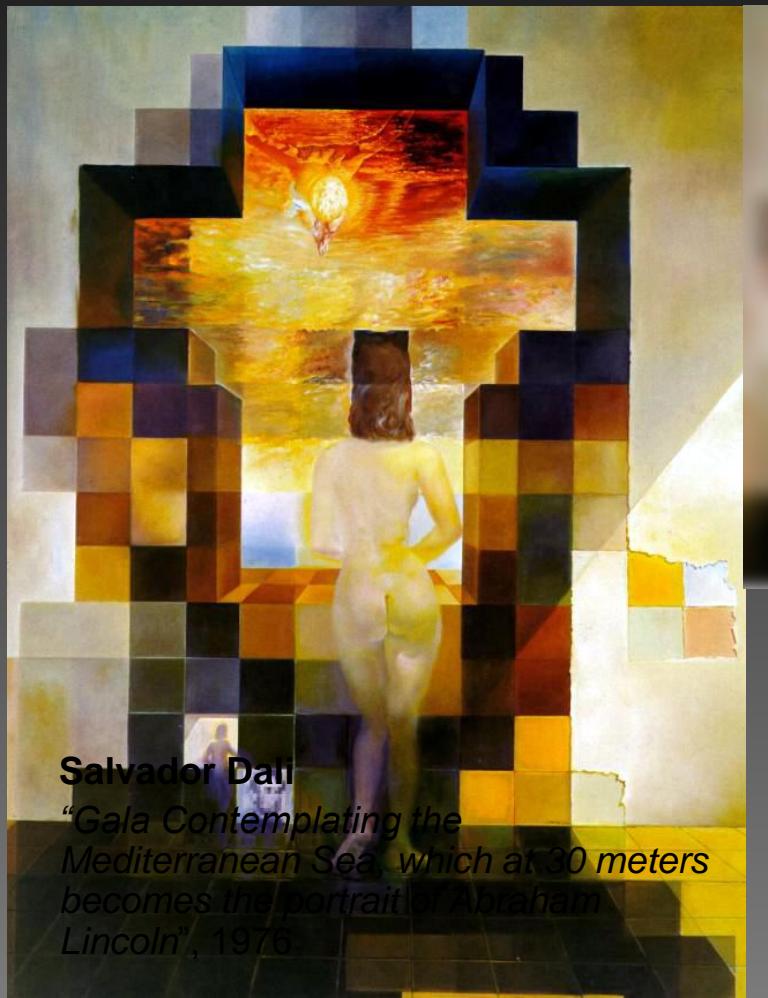


FIGURE 4.16: Examples of Freeman *et al.*'s system controlling a television set. Each state is illustrated with what the television sees on the left and what the user sees on the right. In (a), the television is asleep, but a process is watching the user. An open hand causes the television to come on and show its user interface panel (b). Focus on the panel tracks the movement of the user's open hand in (c), and the user can change channels by using this tracking to move an icon on the screen in (d). Finally, the user displays a closed hand in (e) to turn off the set. *This figure was originally published as Figure 12 of "Computer Vision for Interactive Computer Graphics," W. Freeman et al., IEEE Computer Graphics and Applications, 1998 © IEEE, 1998.*

# Fourier transform

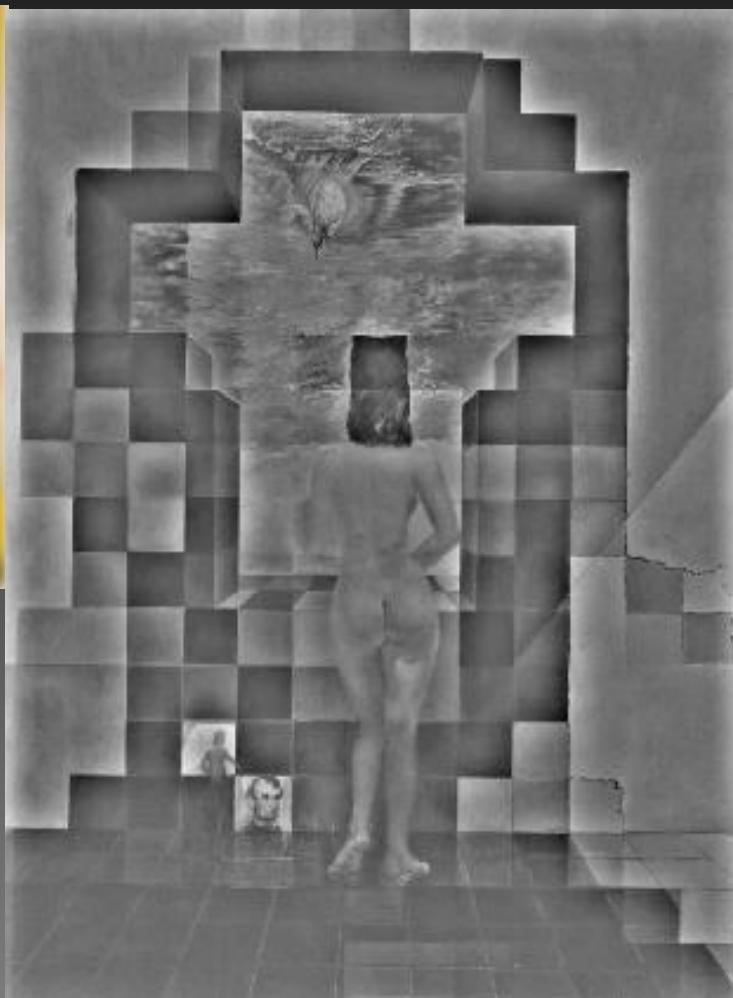
# Fourier Transform

Fourier analysis is a method by which any two dimensional luminance image can be analyzed into the sum of a set of sinusoidal gratings that differ in **spatial frequency, orientation, amplitude and phase**.



Salvador Dalí

"Gala Contemplating the  
Mediterranean Sea, which at 30 meters  
becomes the portrait of Abraham  
Lincoln", 1976



# Linear image transformations

- In analyzing images, it's often useful to make a change of basis.

transformed image

$$\vec{F} = \vec{U}\vec{f}$$

Vectorized image

Fourier transform, or  
Wavelet transform, or  
Steerable pyramid transform

# Self-inverting transforms

Same basis functions are used for the inverse transform

$$\vec{f} = U^{-1} \vec{F}$$

$$= U^+ \vec{F}$$



U transpose and complex conjugate

# An example of such a transform: the Discrete Fourier transform

Forward transform

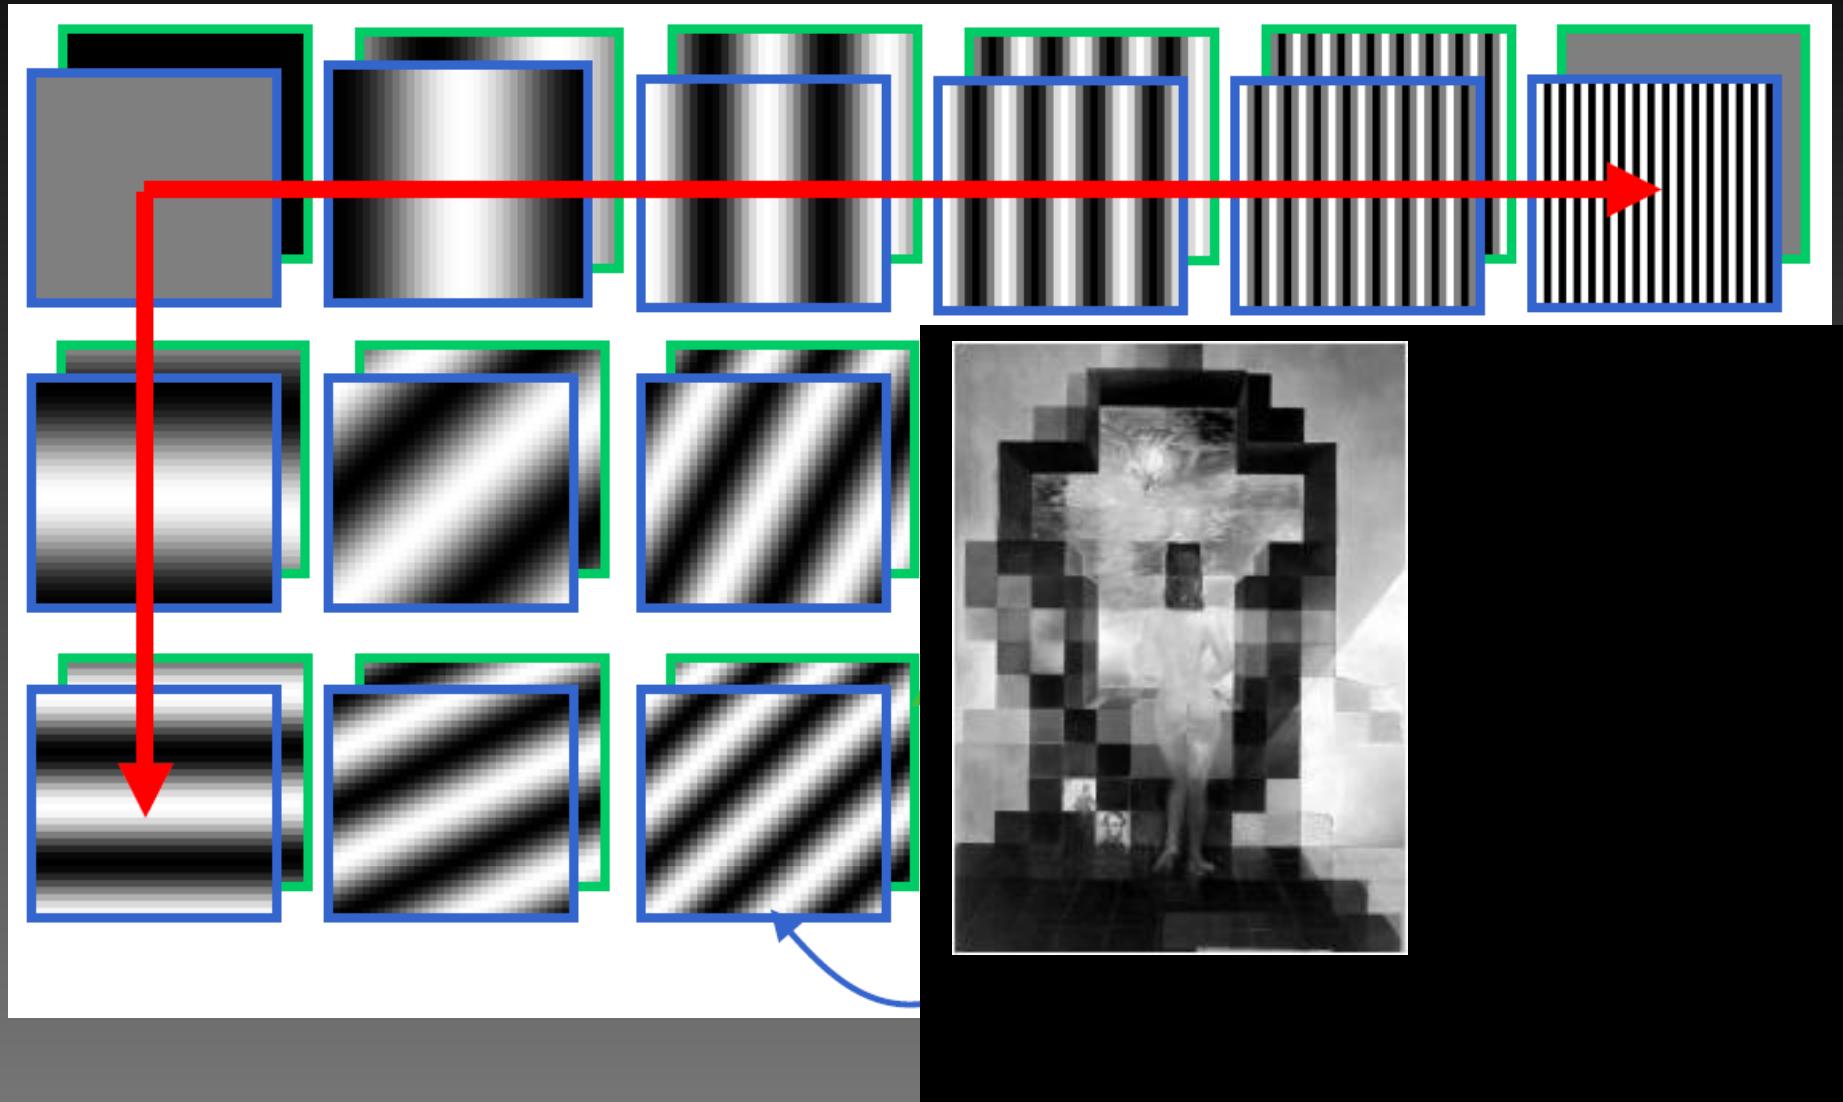
$$F[m, n] = \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} f[k, l] e^{-\pi i \left( \frac{km}{M} + \frac{ln}{N} \right)}$$

Inverse transform

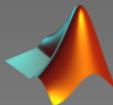
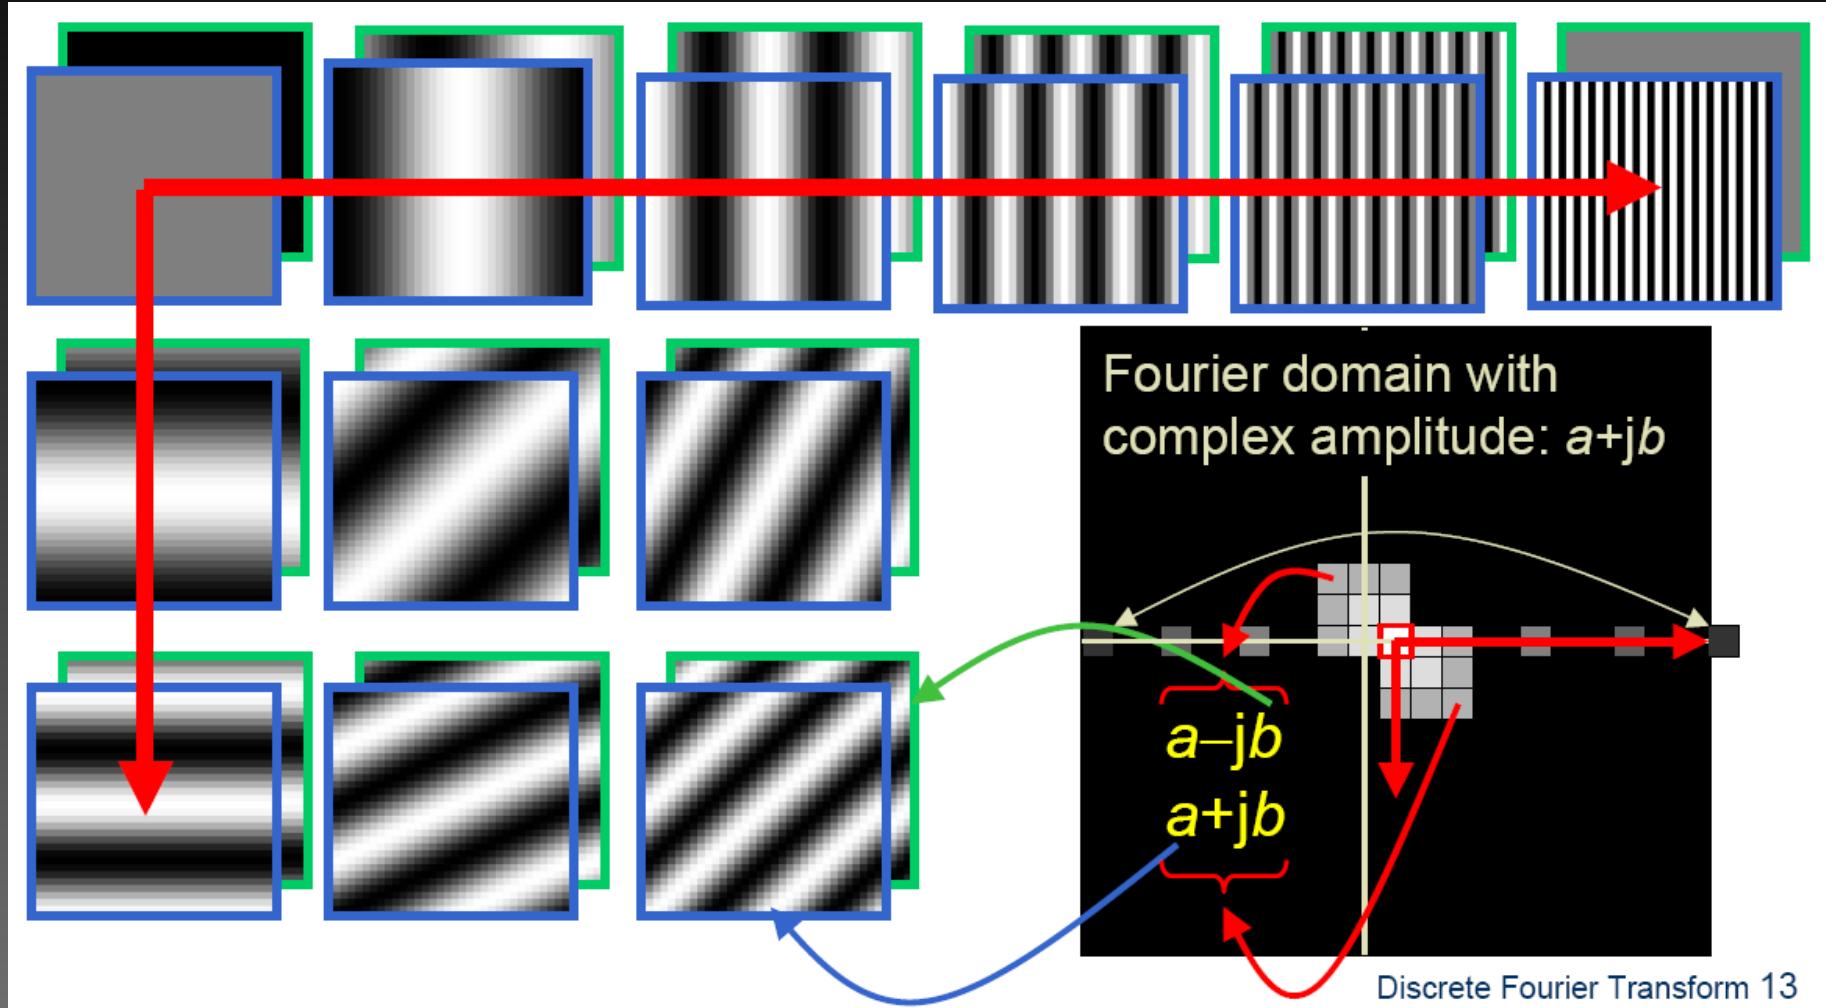
$$f[k, l] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} F[m, n] e^{+\pi i \left( \frac{km}{M} + \frac{ln}{N} \right)}$$

# Sinusoidal gratings as the “primitives” of an image

A nice set of basis: Teases away fast vs. slow changes in the image.



# Sinusoidal gratings as the “primitives” of an image

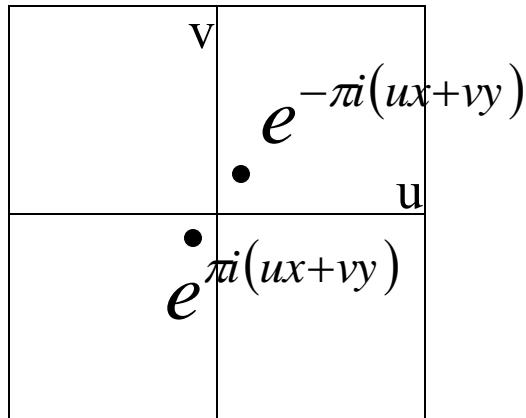


```
imagesc(log(abs(fftshift(fft2(im)))));
```

Slide from A. Efros

To get some sense of what basis elements look like, we plot a basis element --- or rather, its real part ---

as a function of  $x, y$  for some fixed  $u, v$ . We get a function that is constant when  $(ux+vy)$  is constant. The magnitude of the vector  $(u, v)$  gives a frequency, and its direction gives an orientation. The function is a sinusoid with this frequency along the direction, and constant perpendicular to the direction.



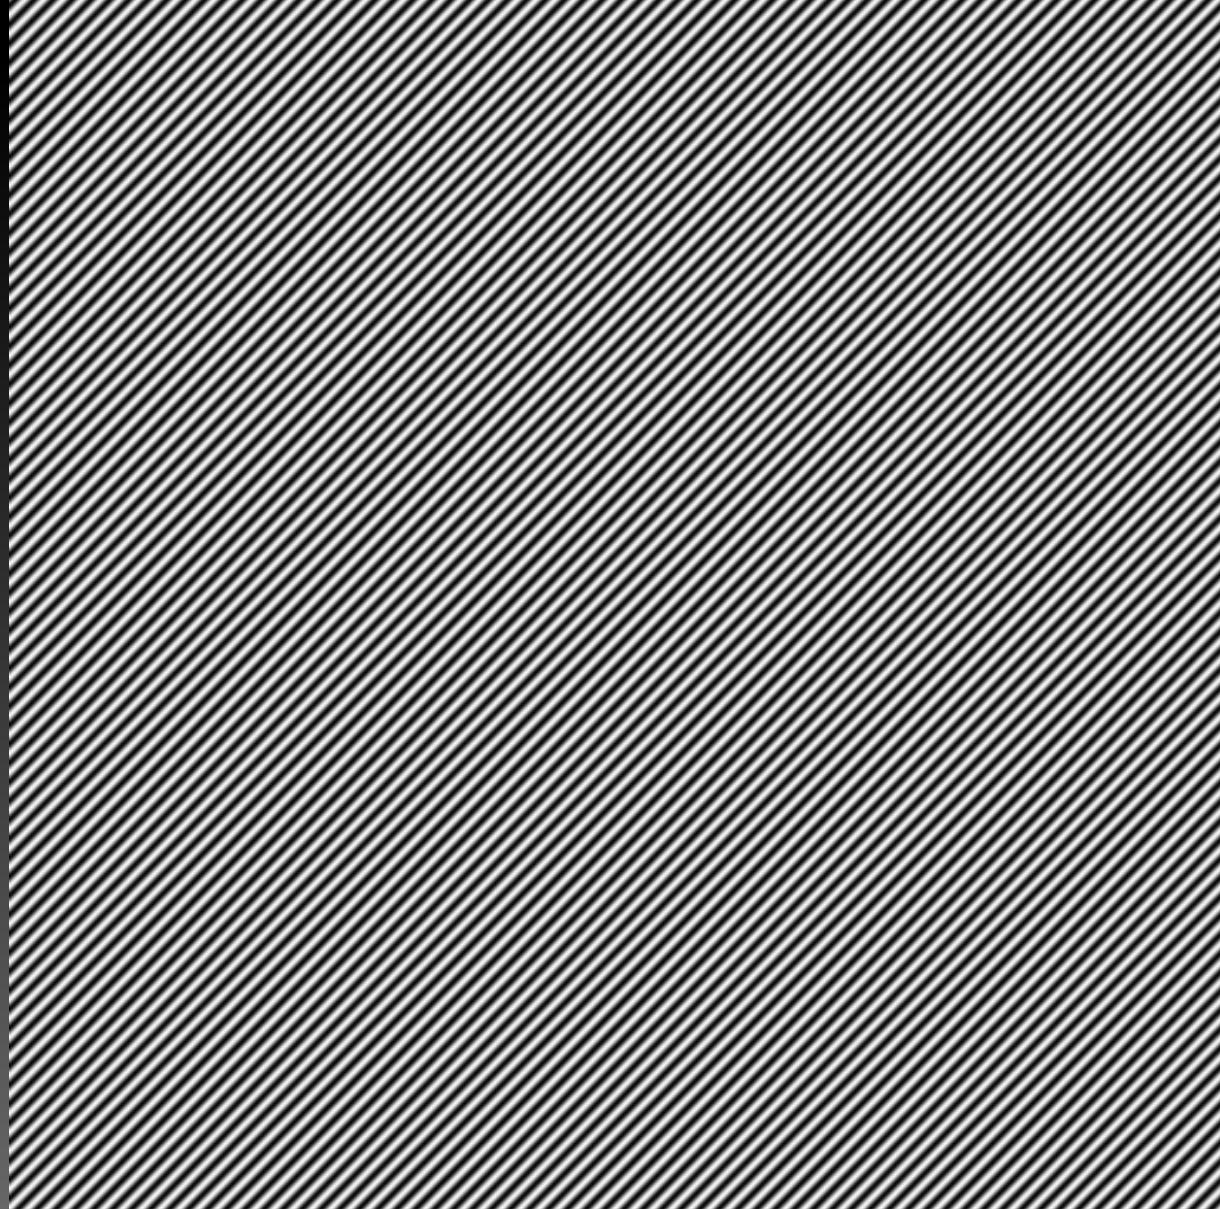
Here  $u$  and  $v$   
are larger  
than in the  
previous  
slide.

	$e^{-\pi i(ux+vy)}$
•	
	$e^{\pi i(ux+vy)}$

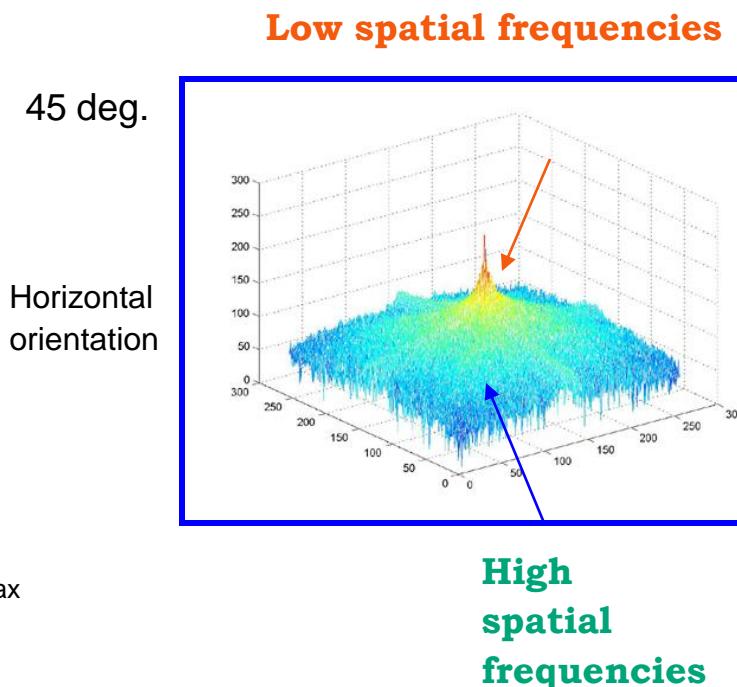
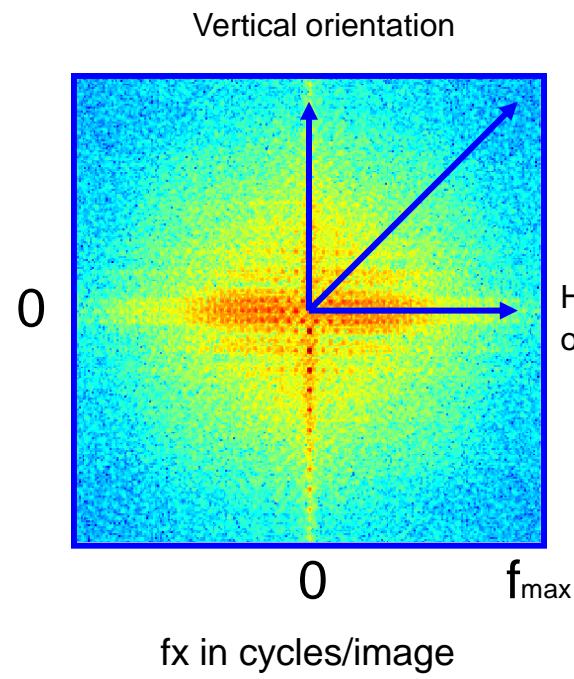
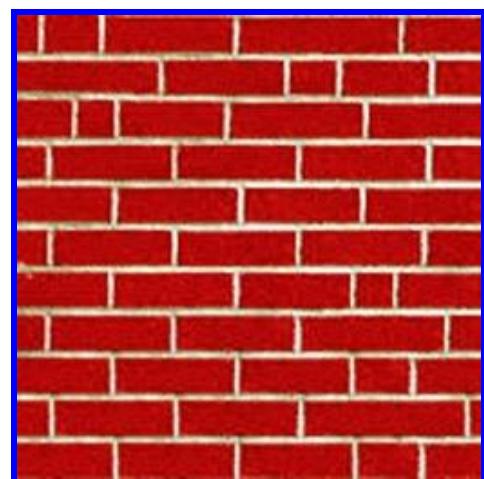


And larger still...

$$e^{-\pi i(ux+vy)}$$
$$\begin{array}{|c|c|} \hline & v & \\ \hline \bullet & & \\ \hline & u & \\ \hline \end{array}$$
$$e^{\pi i(ux+vy)}$$

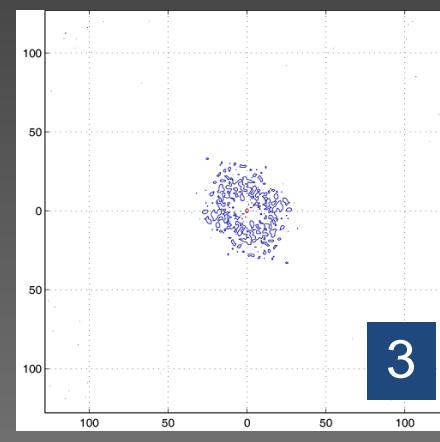
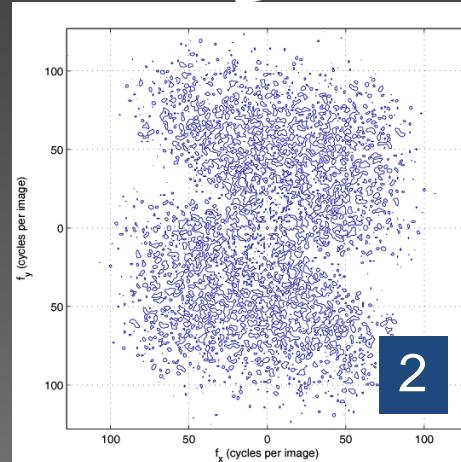
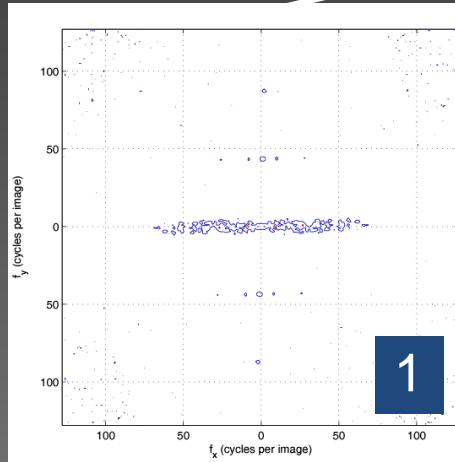
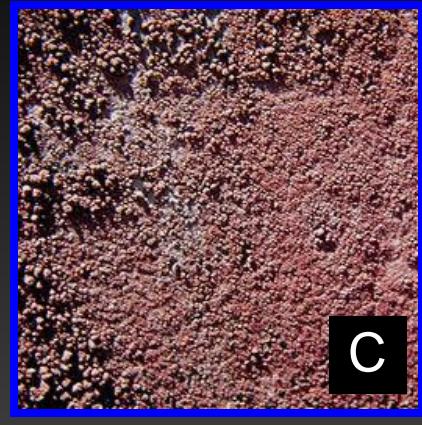
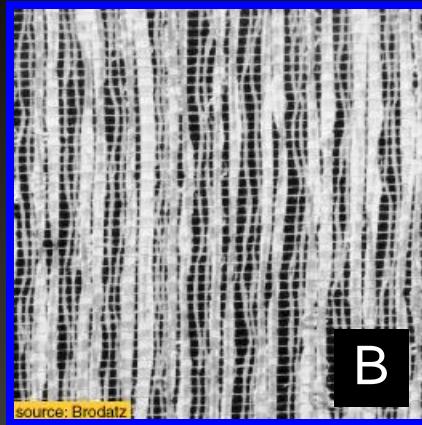


# How to interpret a Fourier Spectrum



Log power spectrum

# Fourier Amplitude Spectrum



# Two examples of image synthesis with Fourier basis

**First:** randomly sample the Fourier coefficients of an image and reconstruct from those.

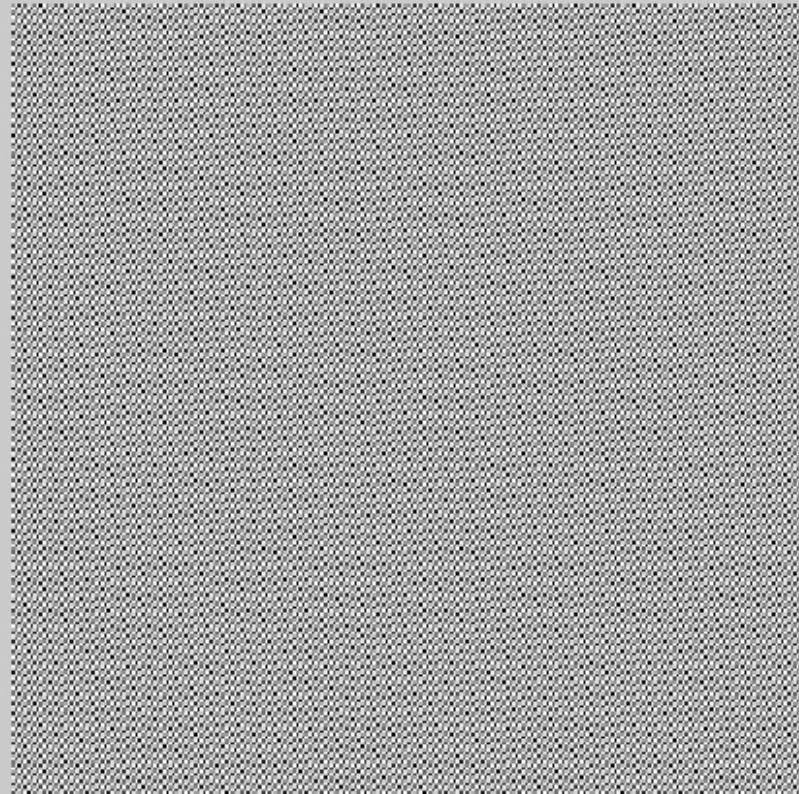
**Second:** sample Fourier coefficients in descending order of amplitude.

2

2



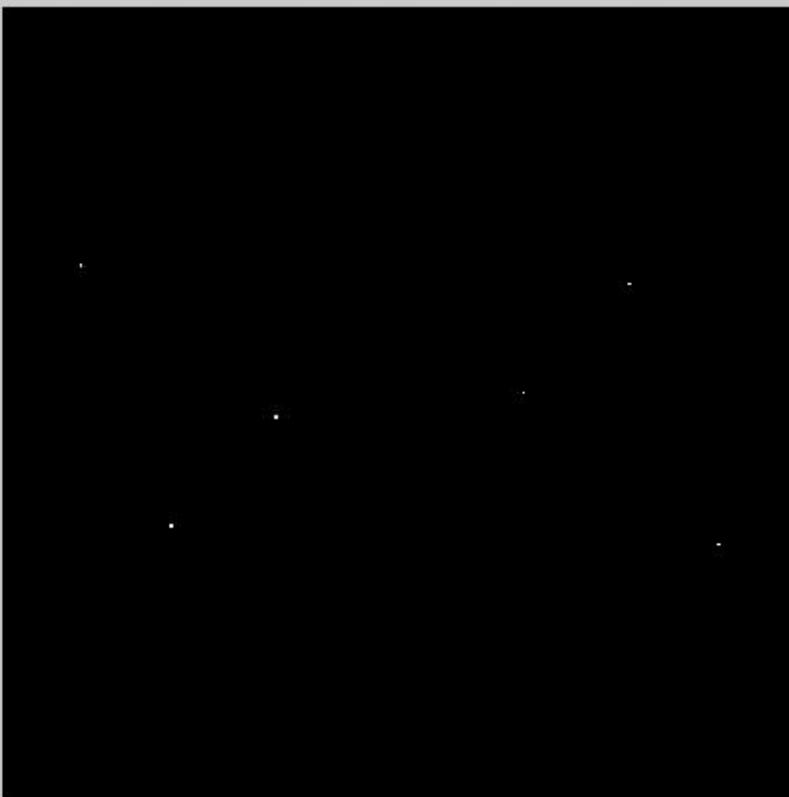
#1: Range [0, 1]  
Dims [256, 256]



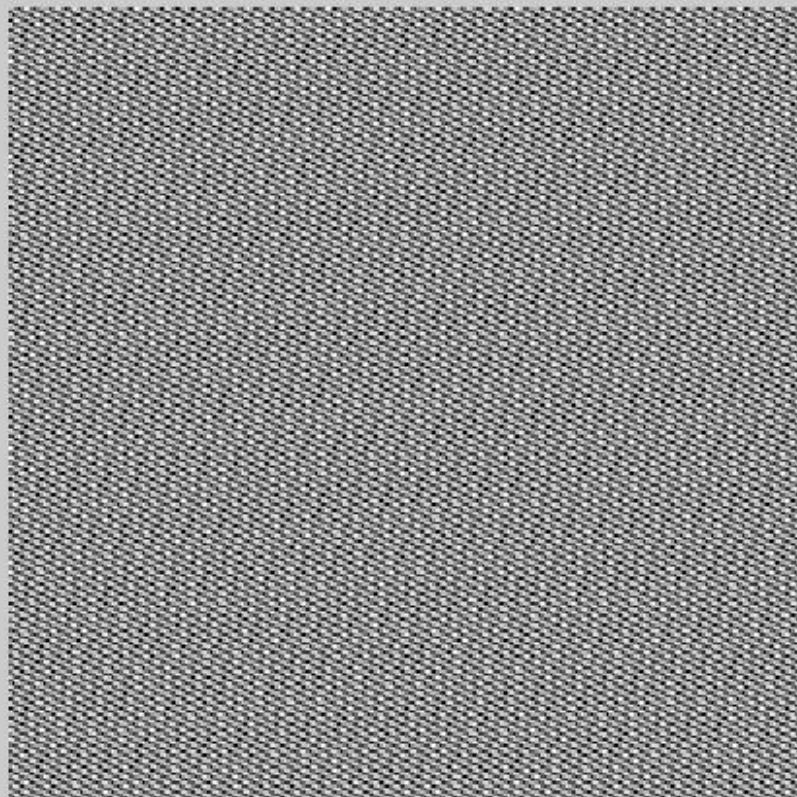
#2: Range [0.000109, 0.0267]  
Dims [256, 256]

# 6

6



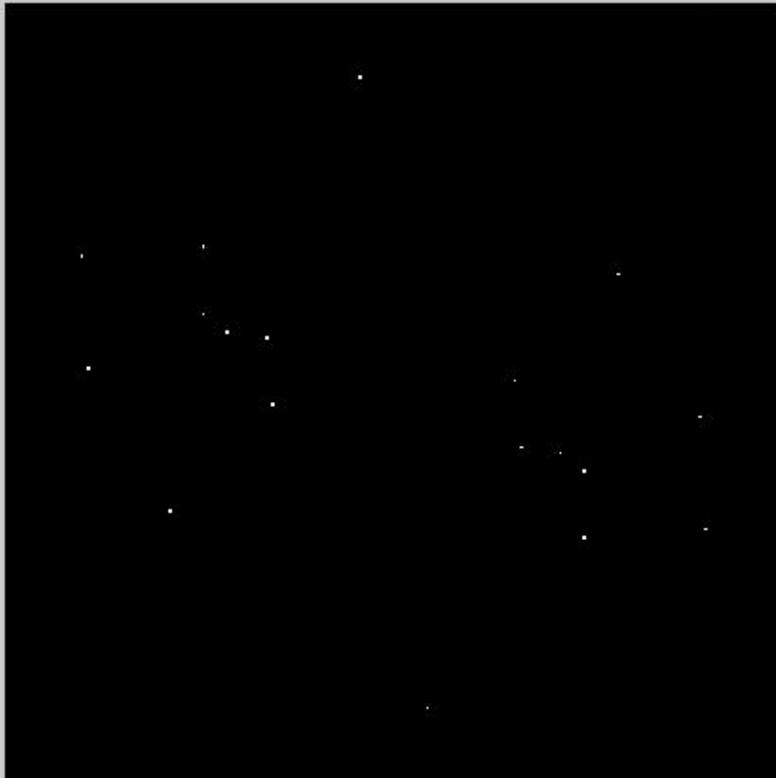
#1: Range [0, 1]  
Dims [256, 256]



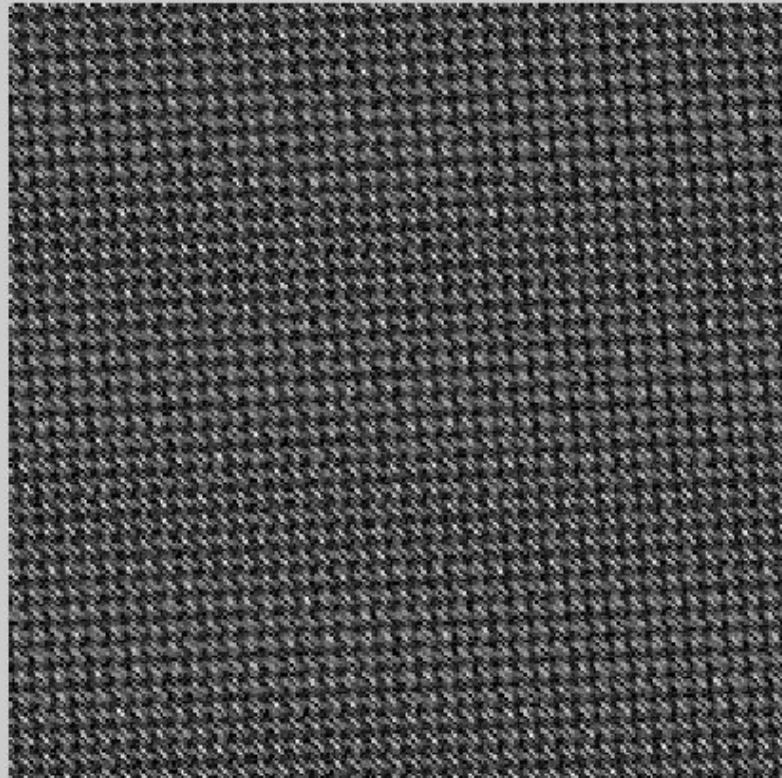
#2: Range [1.89e-007, 0.226]  
Dims [256, 256]

# 18

18



#1: Range [0, 1]  
Dims [256, 256]



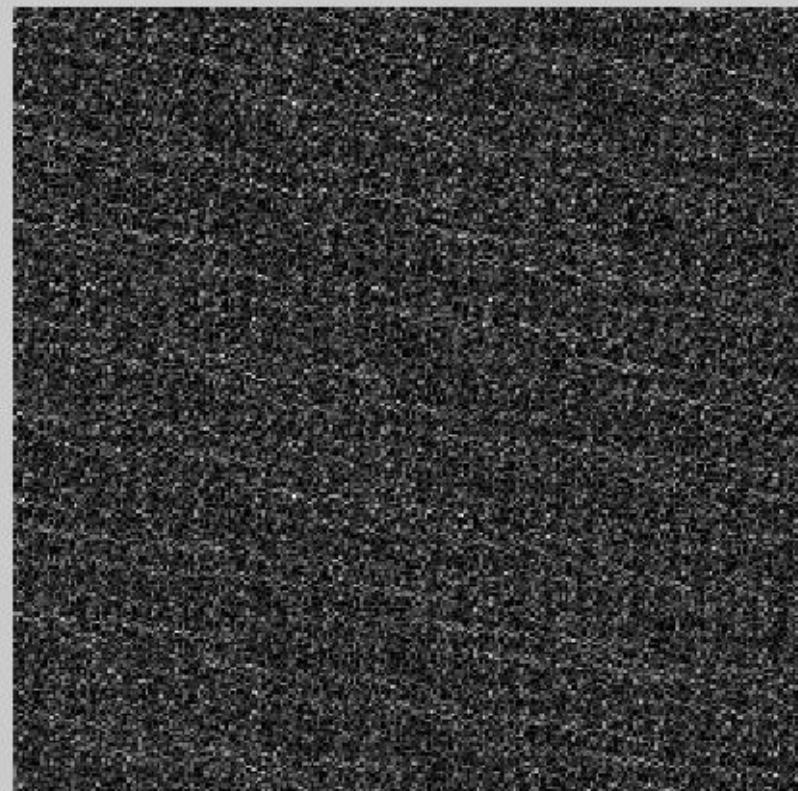
#2: Range [4.79e-007, 0.503]  
Dims [256, 256]

# 136

136



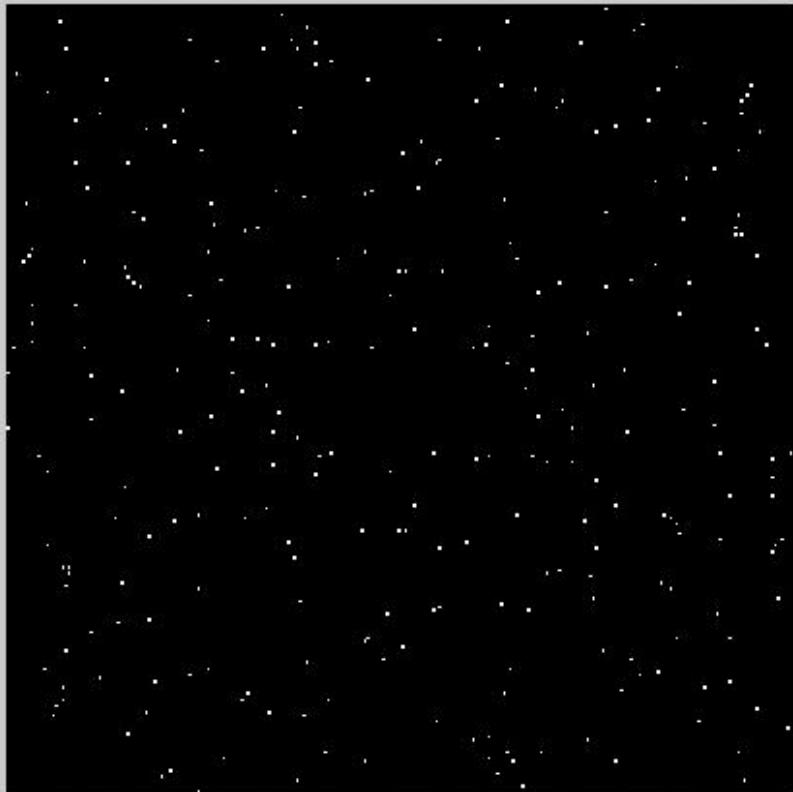
#1: Range [0, 1]  
Dims [256, 256]



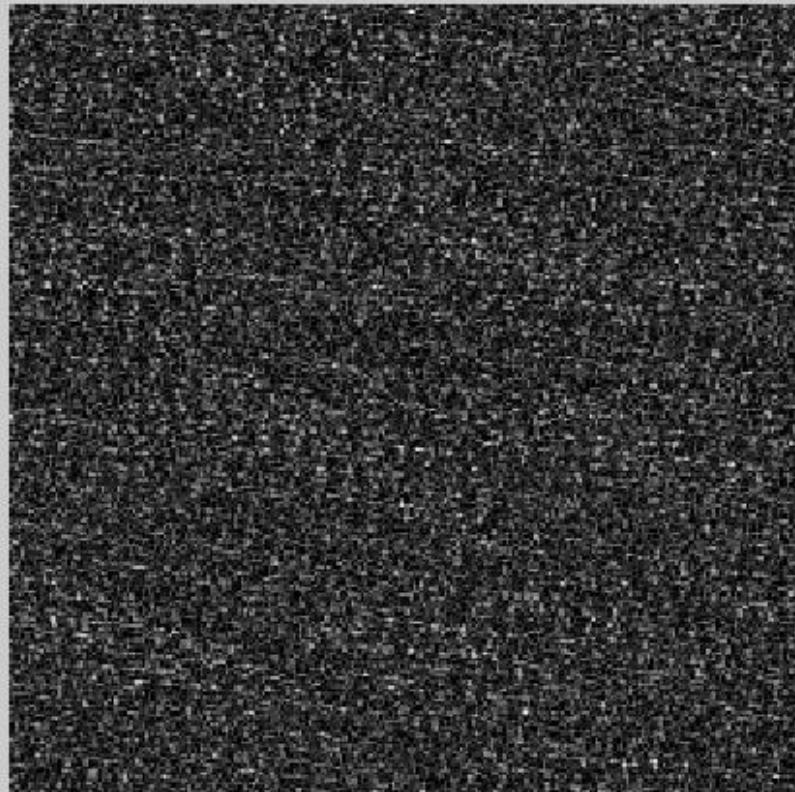
#2: Range [8.25e-006, 3.48]  
Dims [256, 256]

# 282

282



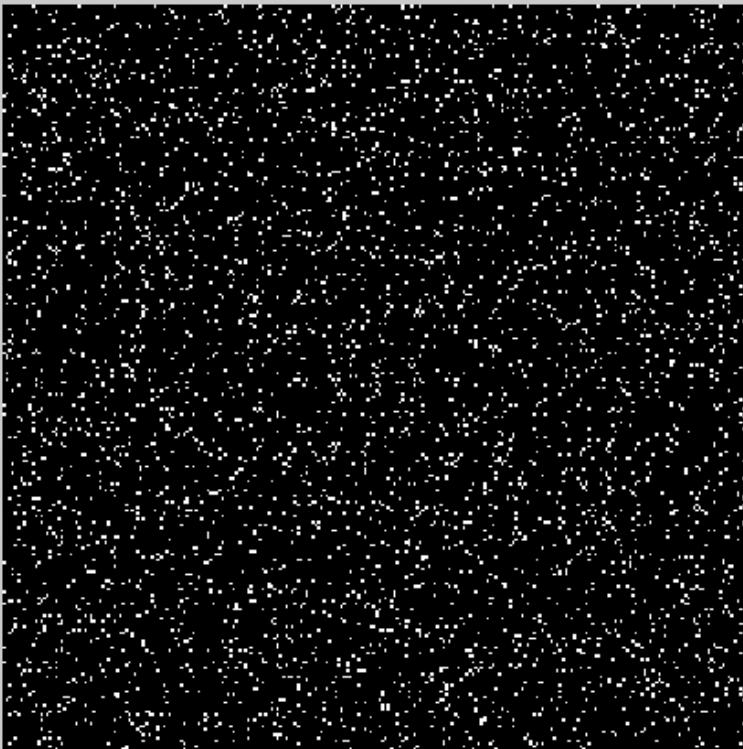
#1: Range [0, 1]  
Dims [256, 256]



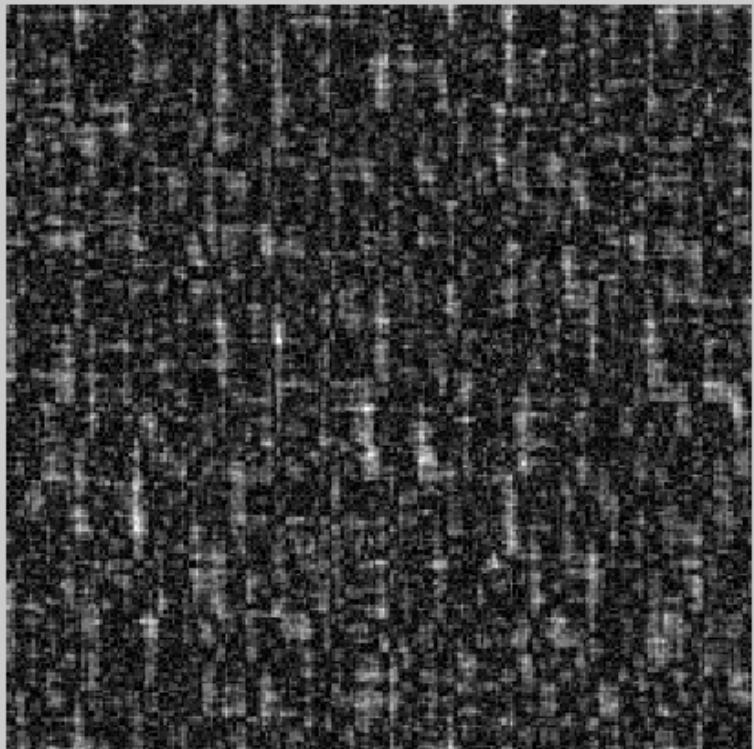
#2: Range [1.39e-005, 5.88]  
Dims [256, 256]

# 4052.

4052



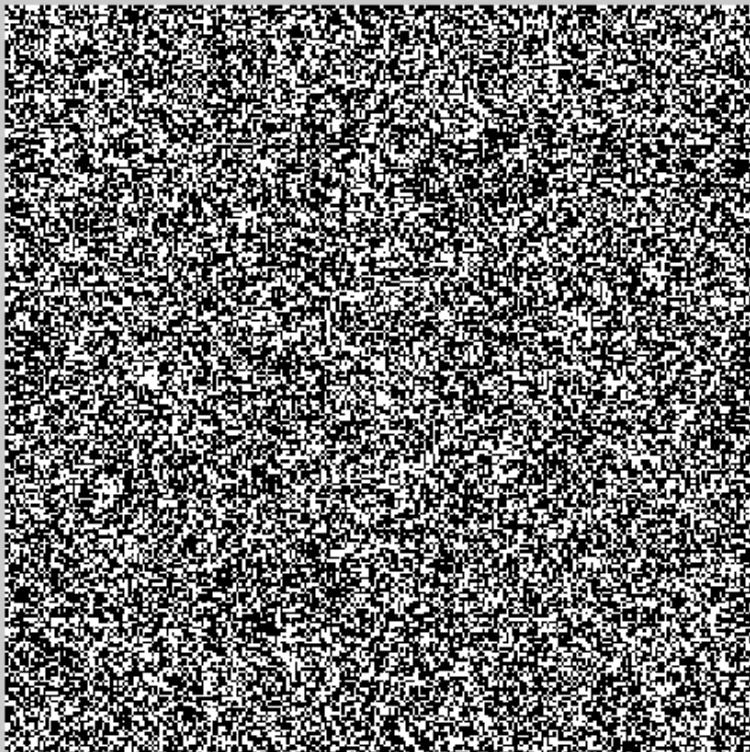
#1: Range [0, 1]  
Dims [256, 256]



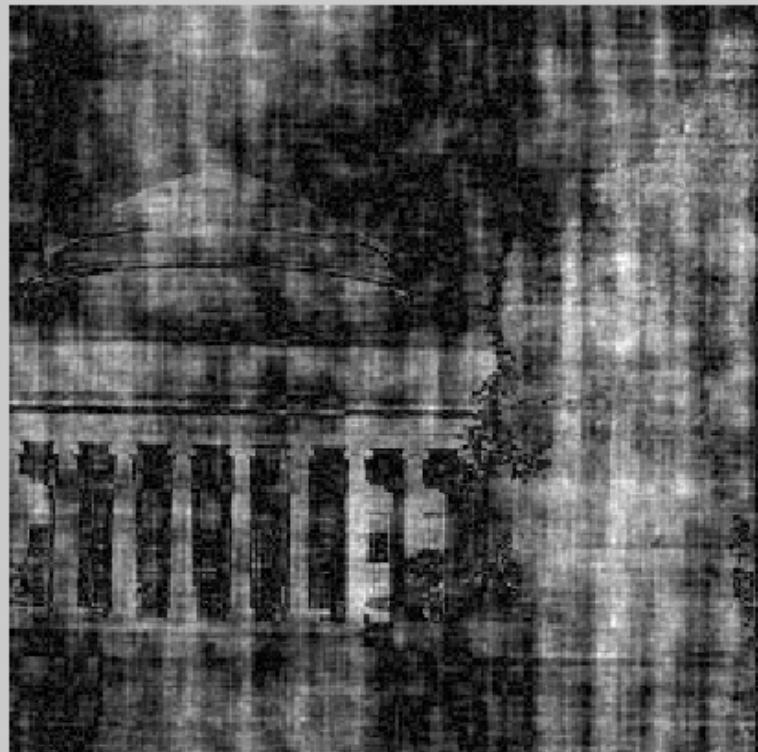
#2: Range [0.000556, 37.7]  
Dims [256, 256]

# 28743

28743



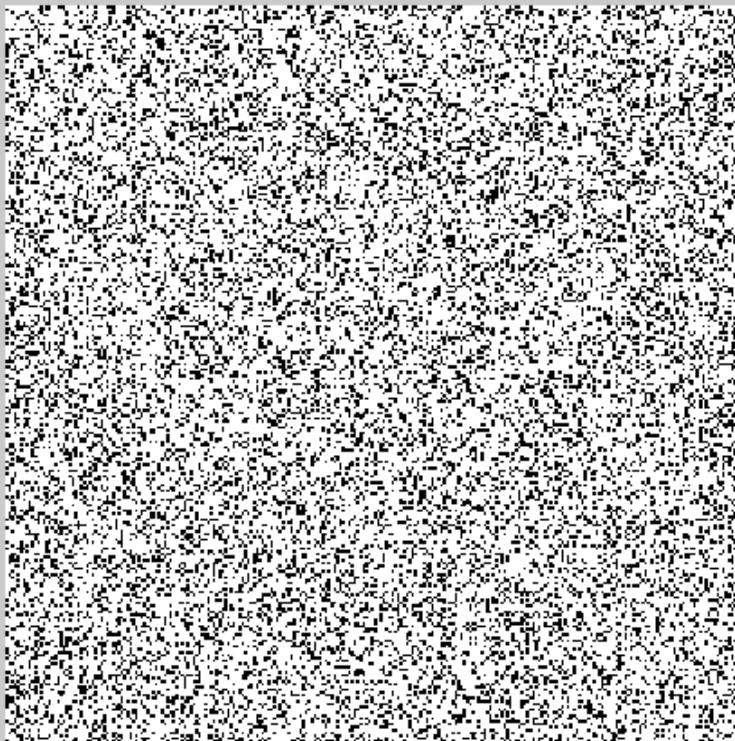
#1: Range [0, 1]  
Dims [256, 256]



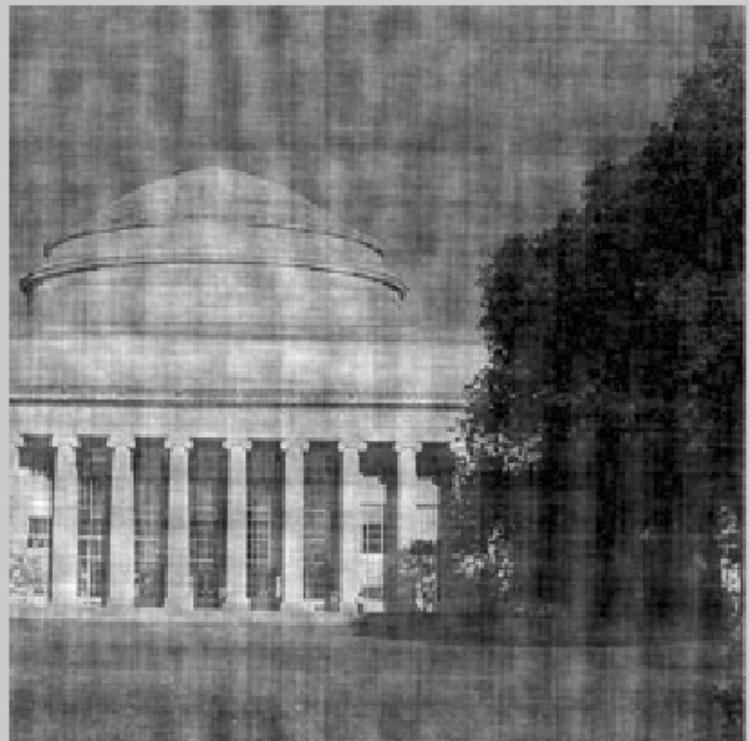
#2: Range [0.00109, 146]  
Dims [256, 256]

# 49190.

49190



#1: Range [0, 1]  
Dims [256, 256]



#2: Range [0.00758, 294]  
Dims [256, 256]

# 65536.

65536.



#1: Range [0.5, 1.5]  
Dims [256, 256]



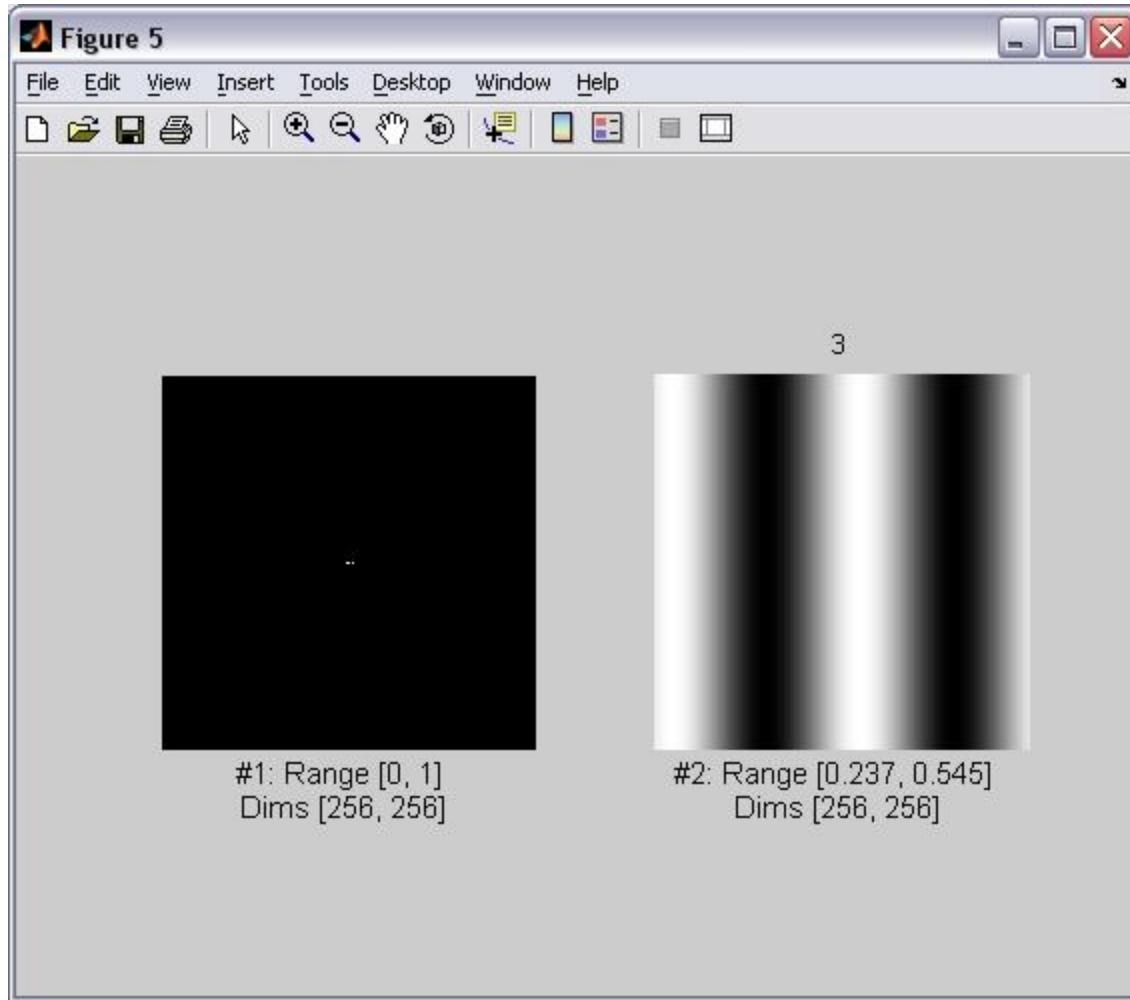
#2: Range [4.43e-015, 256]  
Dims [256, 256]

# Two examples of image synthesis with Fourier basis

First: randomly sample the Fourier coefficients of an image and reconstruct from those.

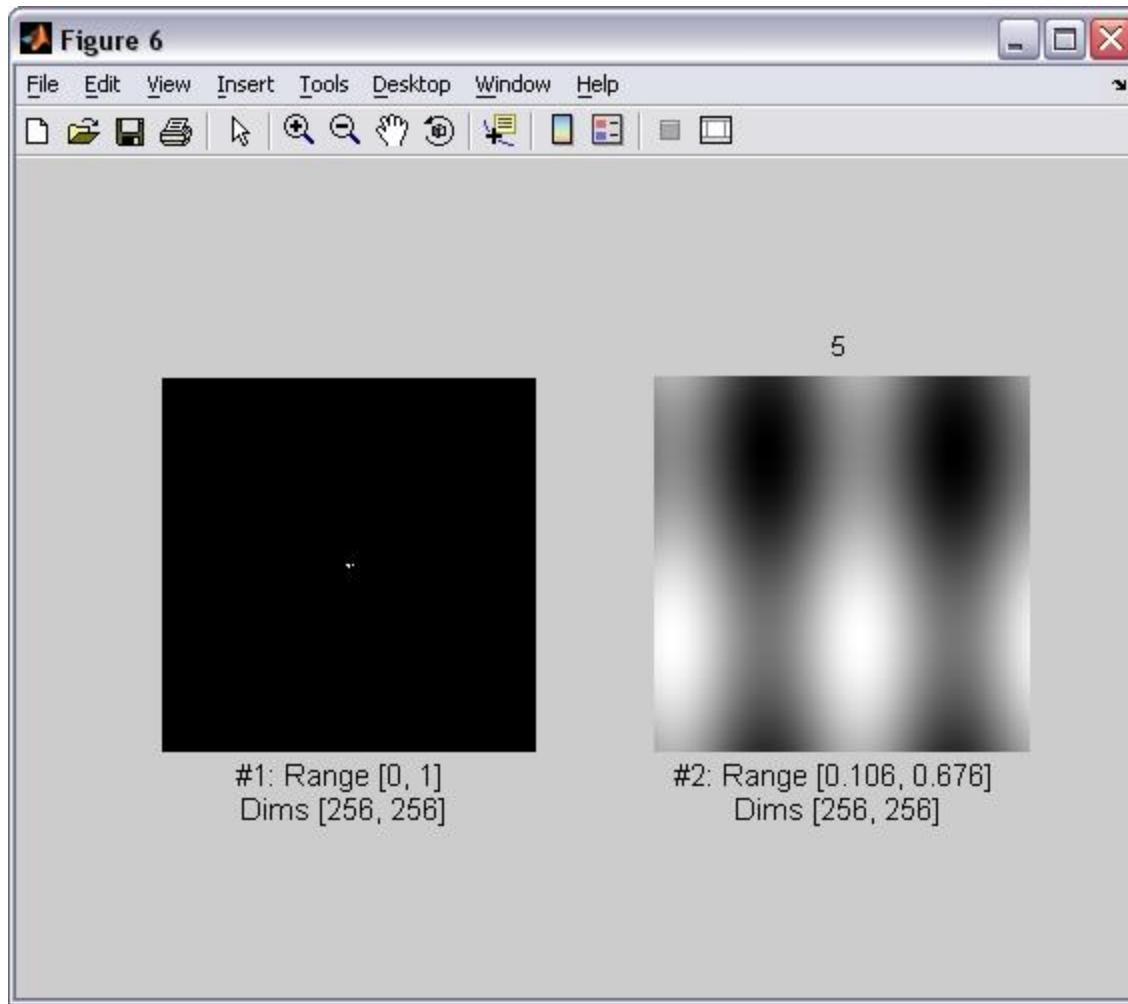
Second: sample Fourier coefficients in descending order of amplitude.

# 3

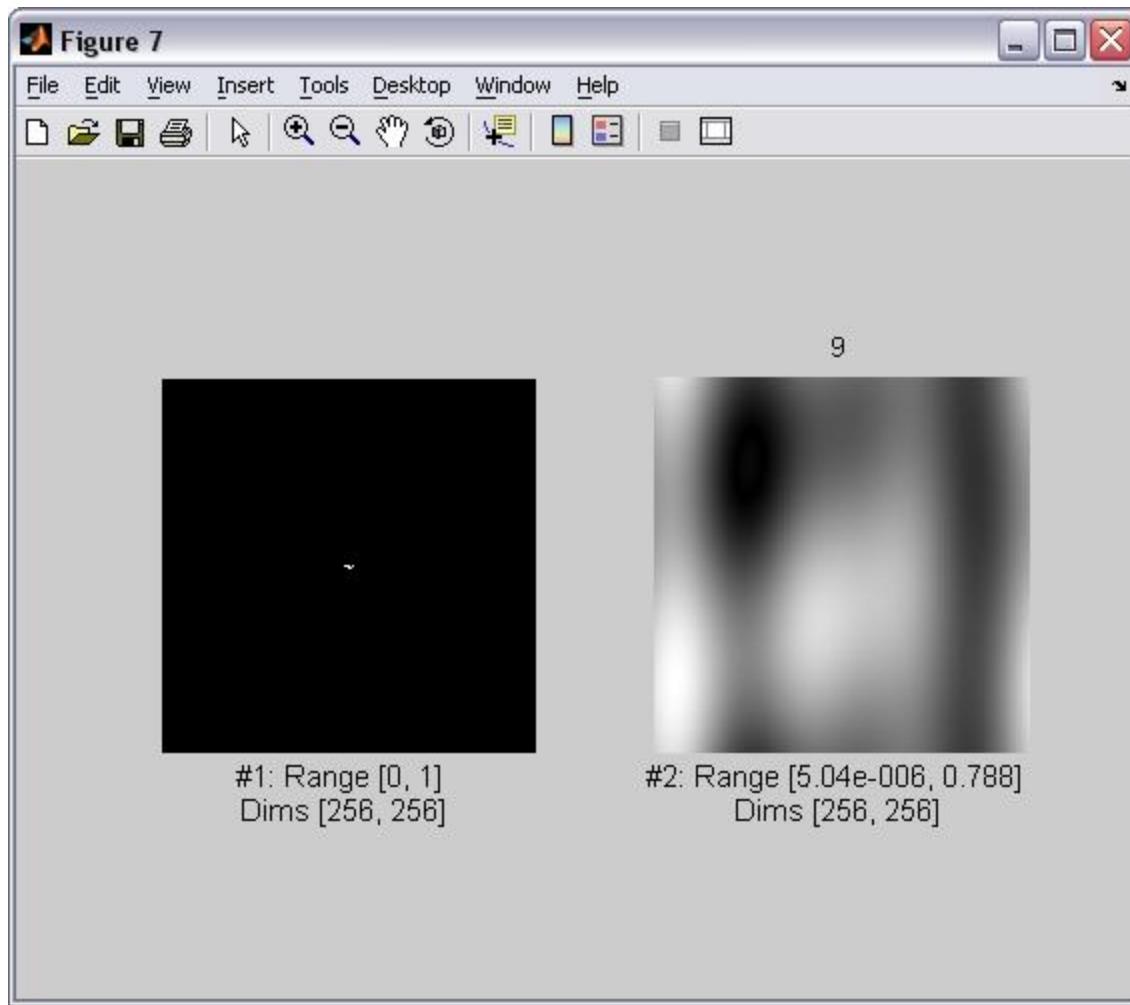


Now, an analogous sequence of images, but selecting Fourier components in descending order of magnitude.

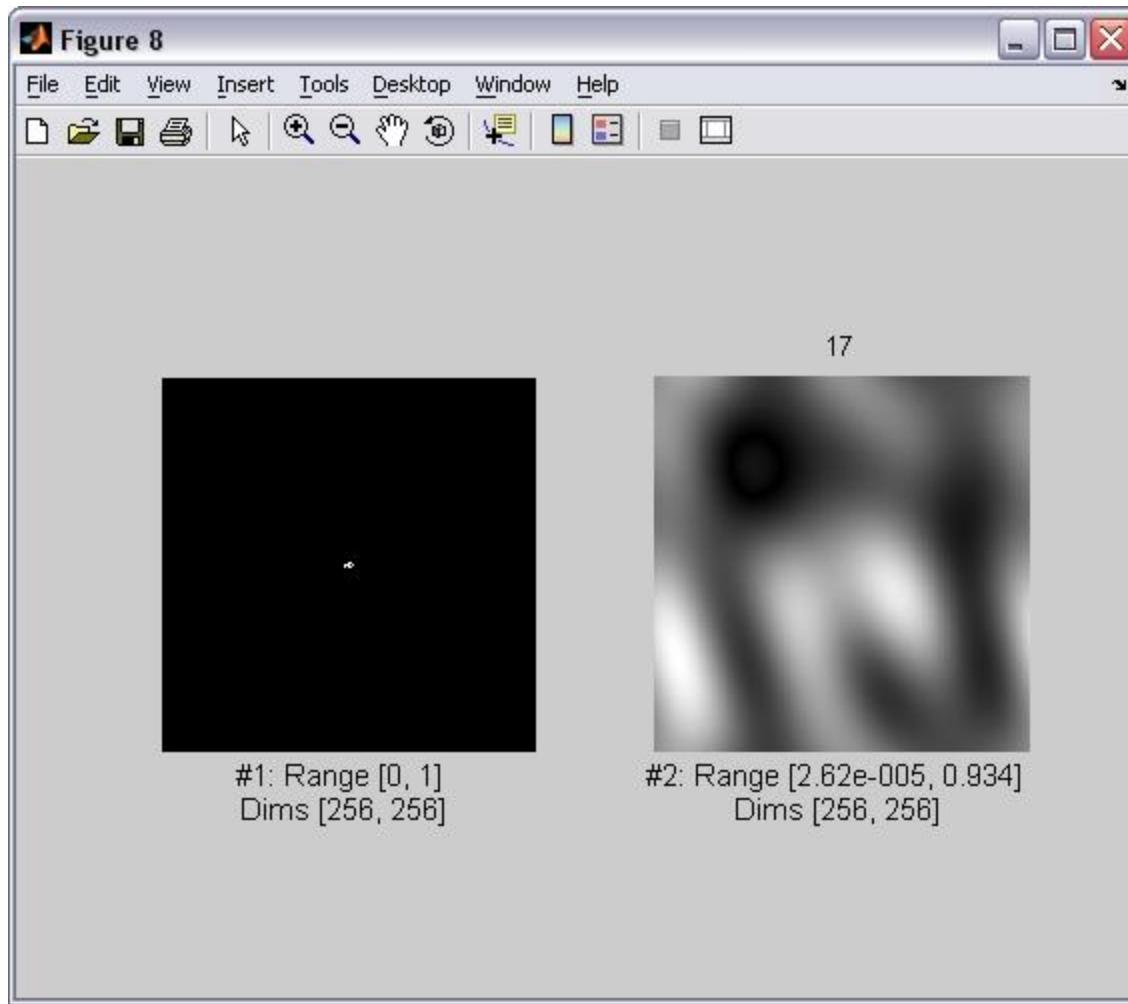
5



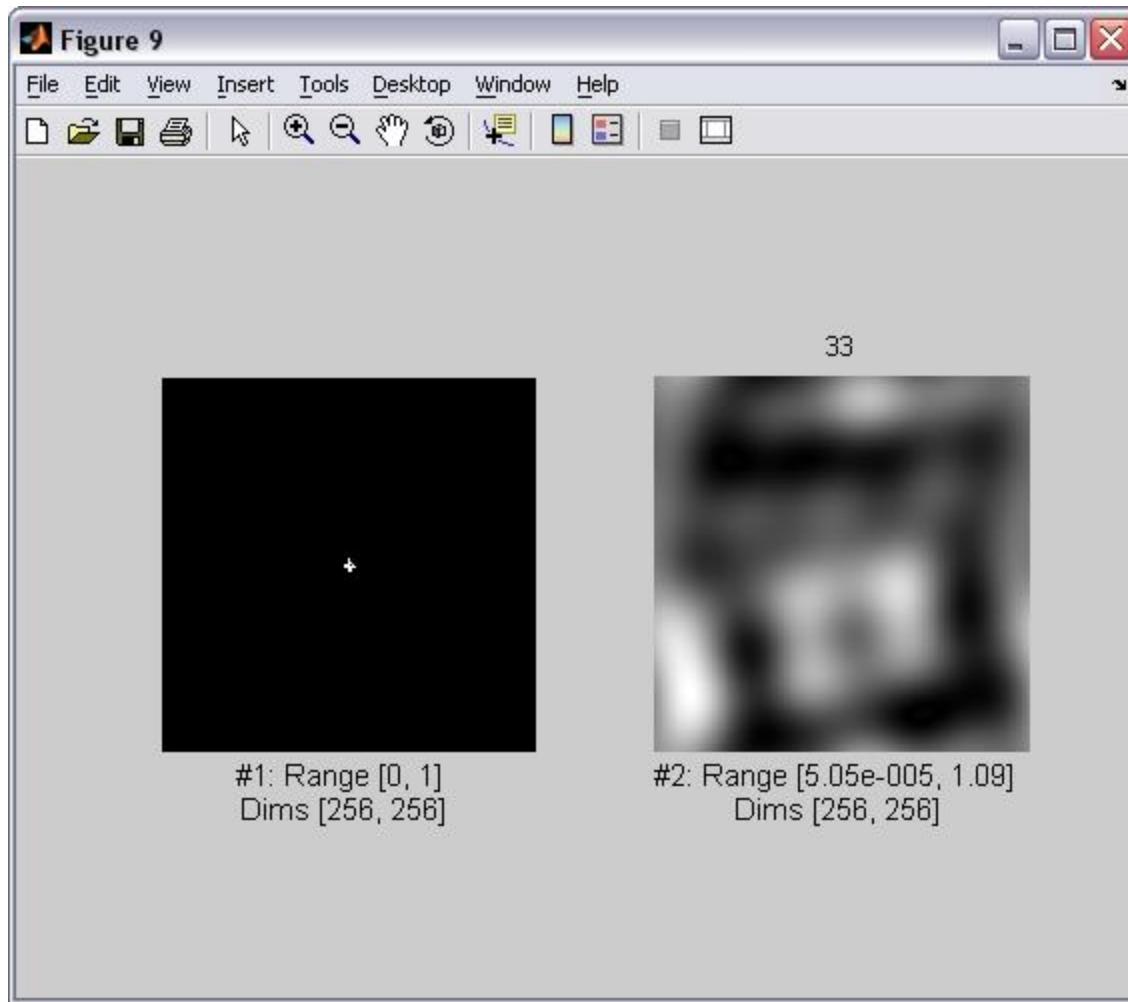
# 9



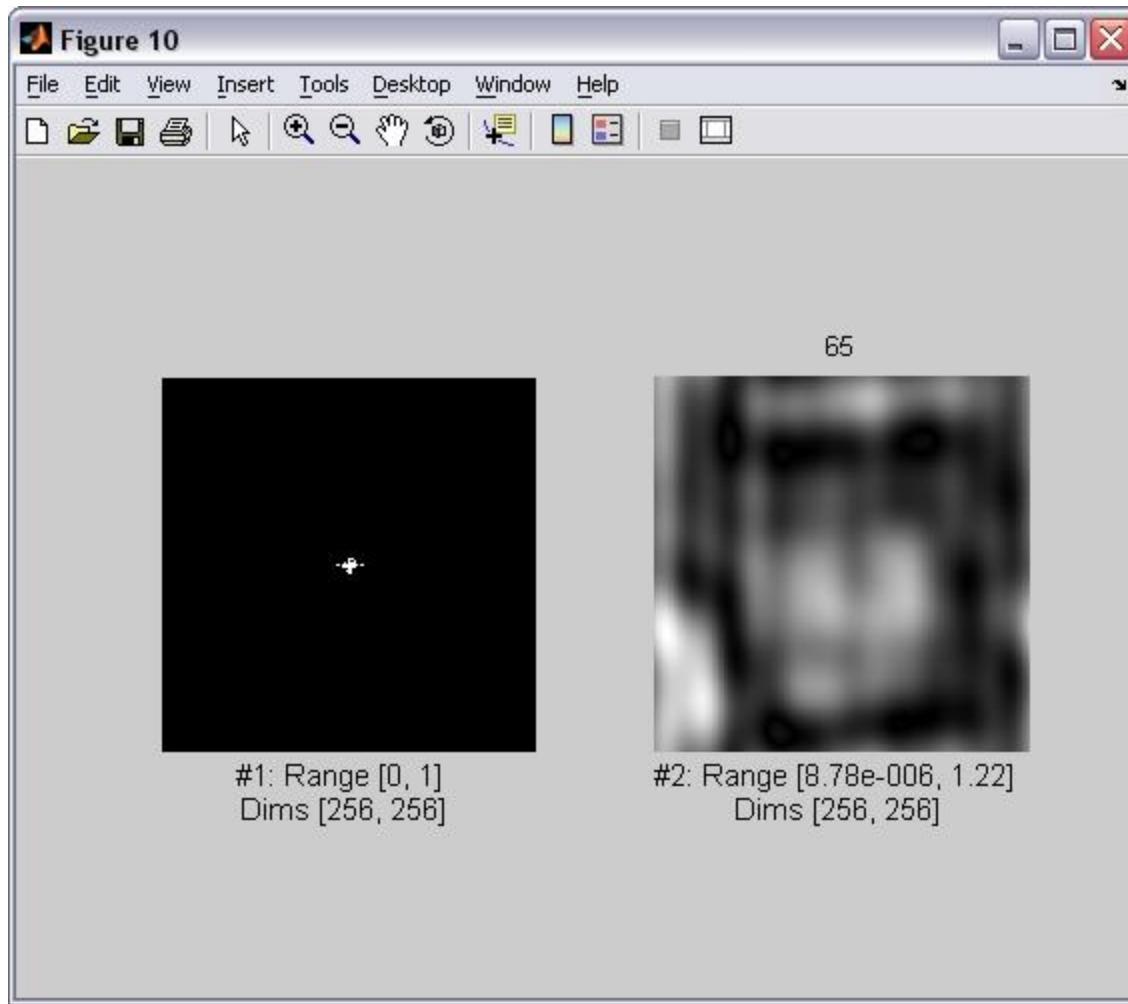
# 17



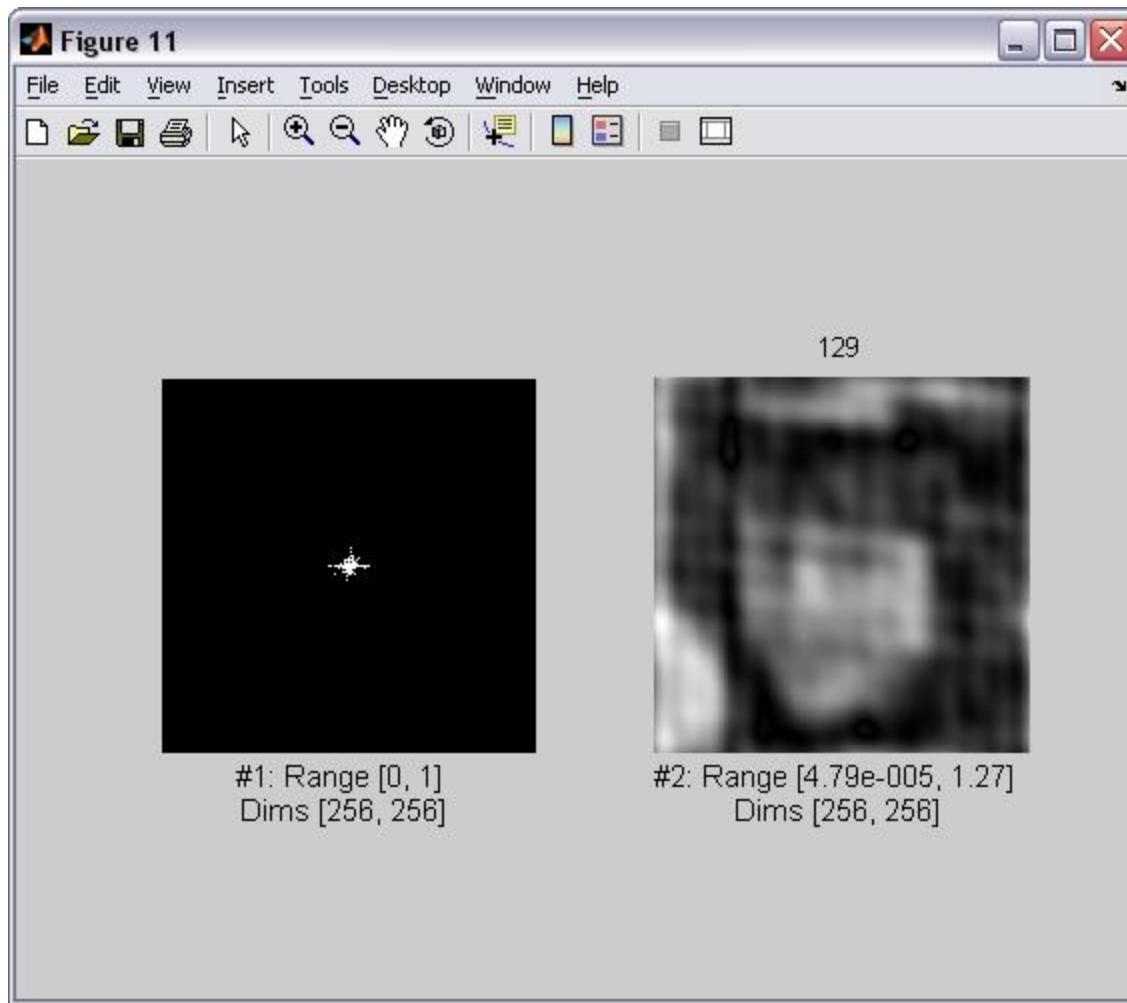
33



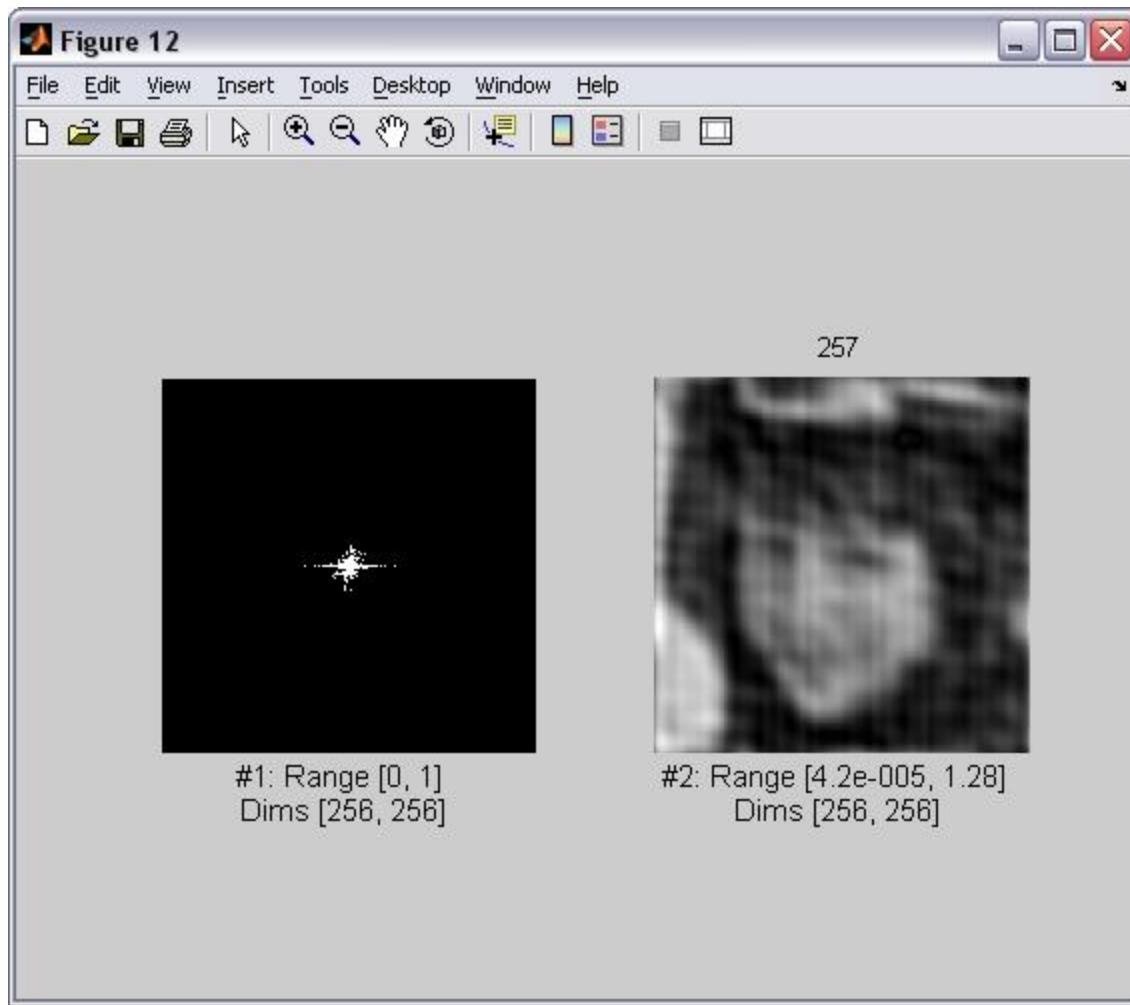
# 65



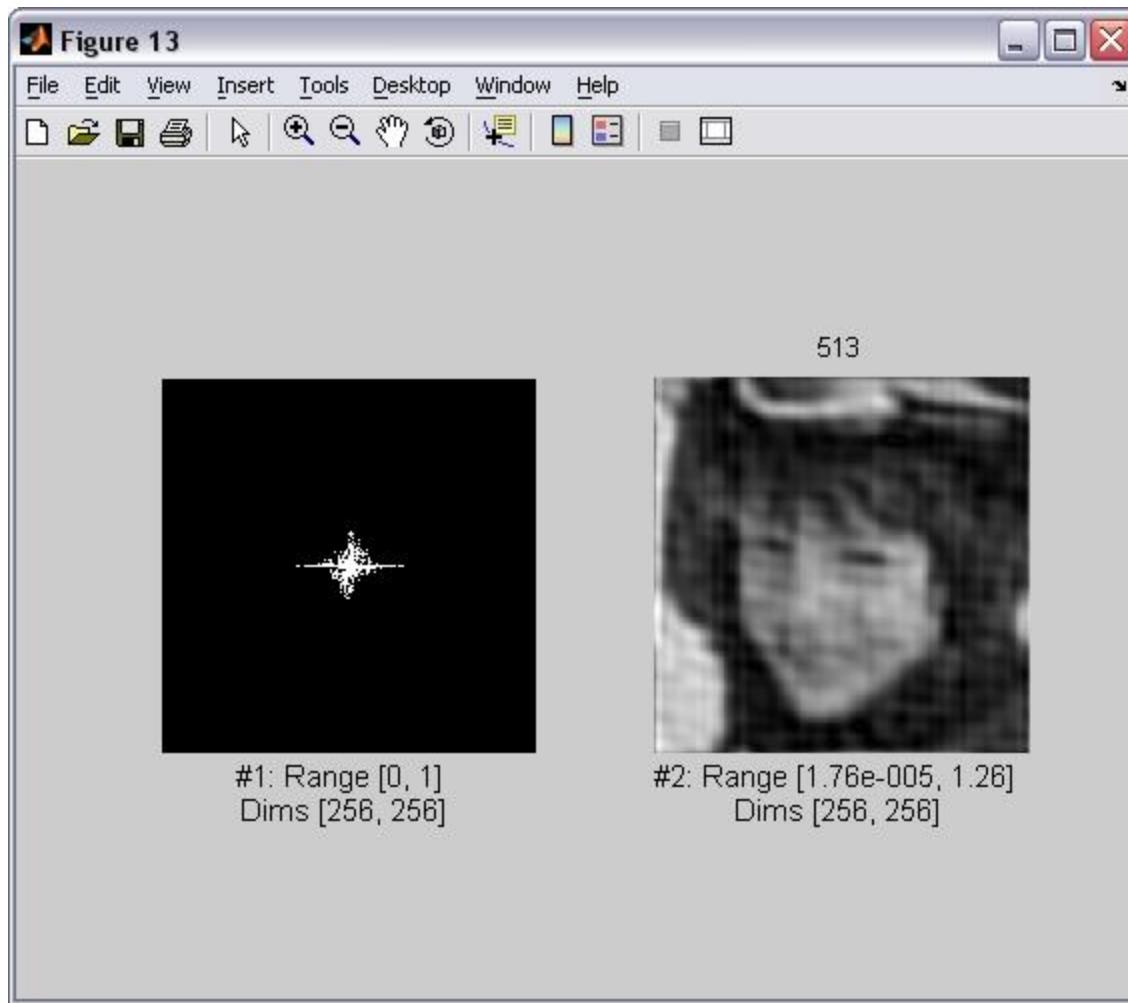
# 129



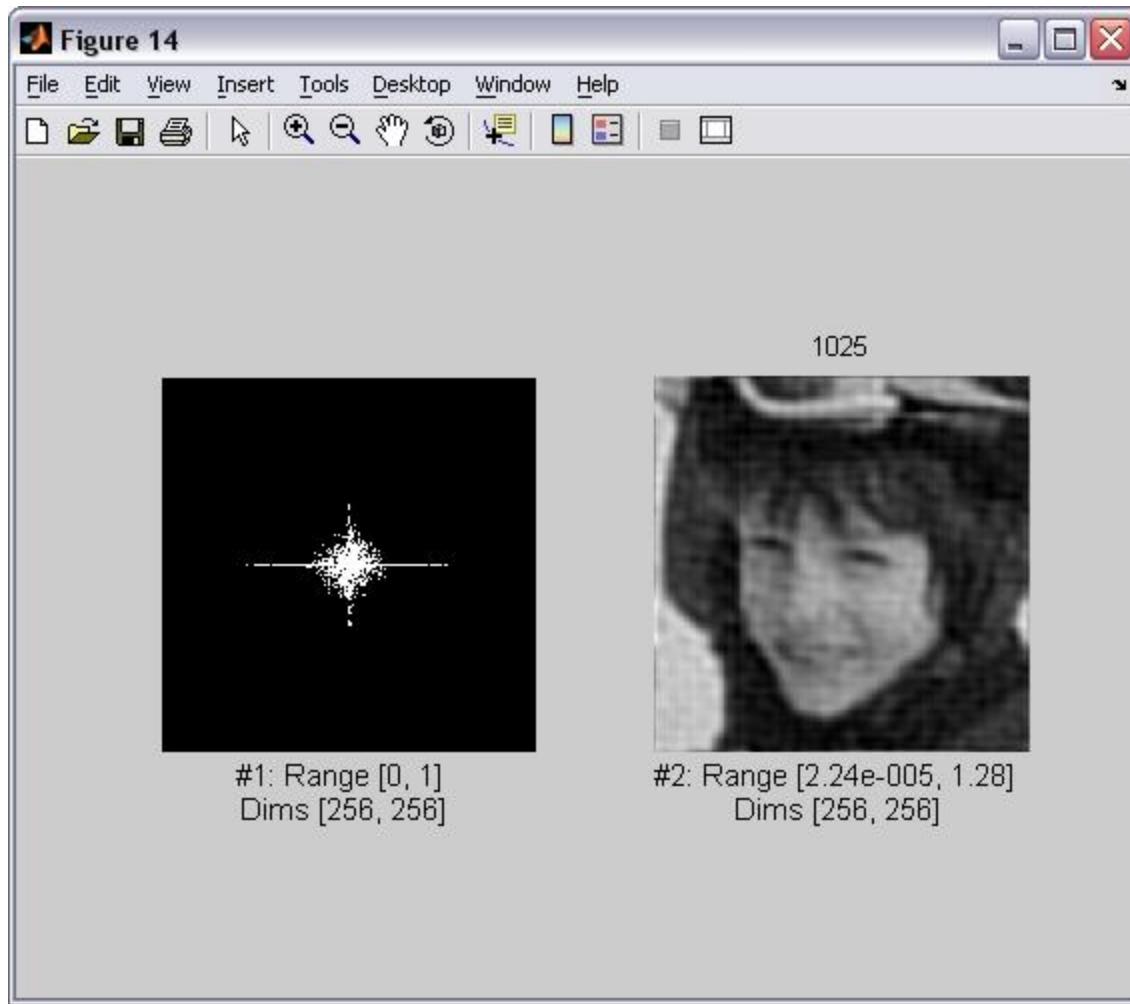
# 257



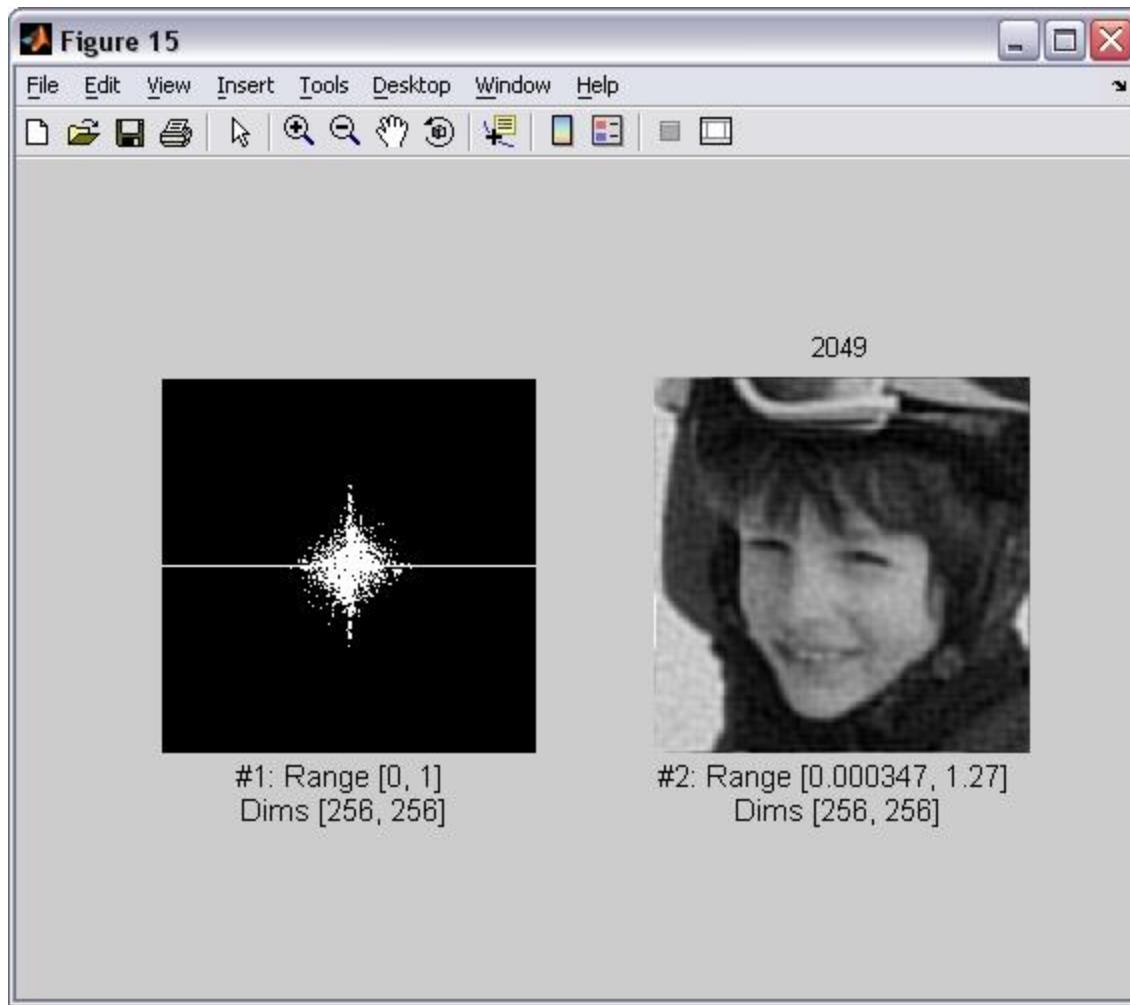
# 513



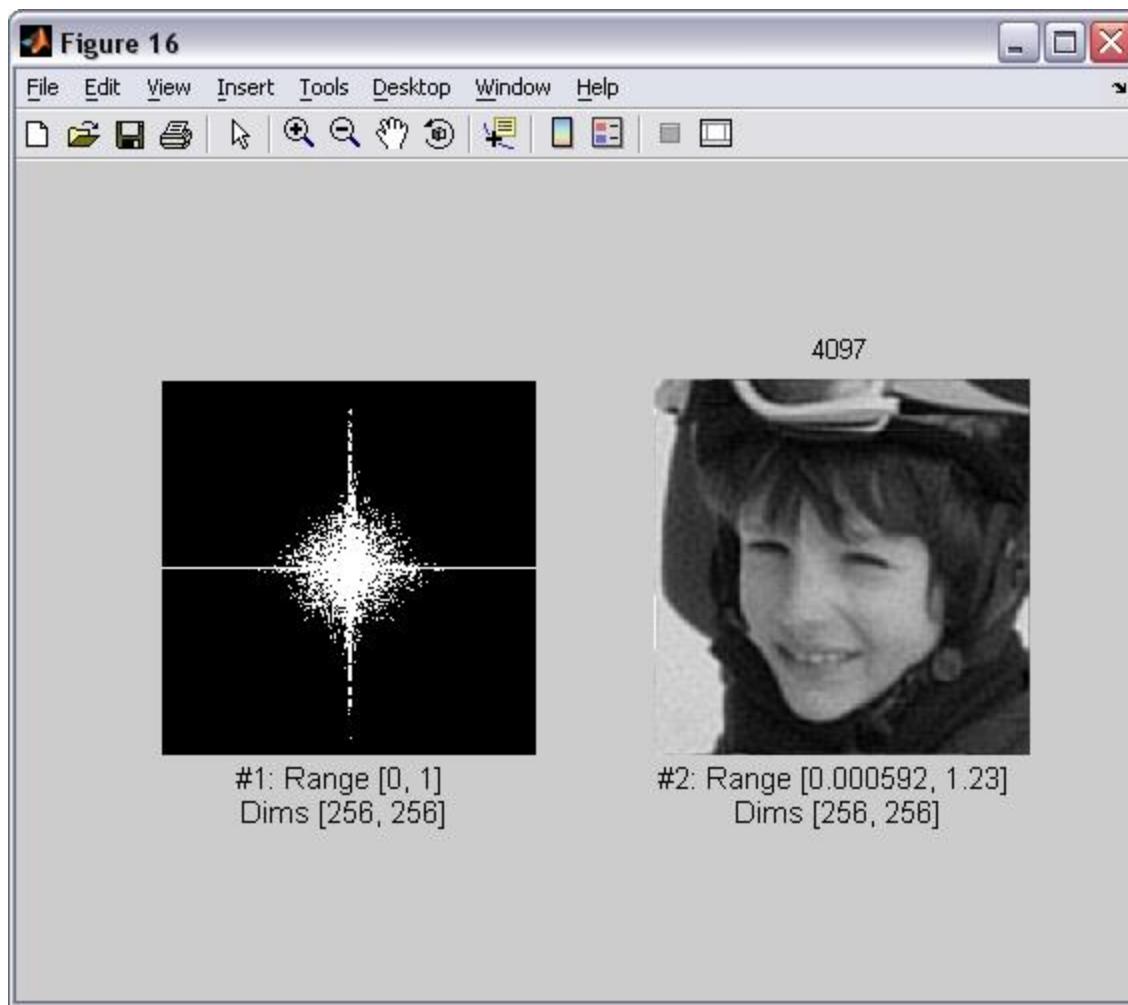
# 1025



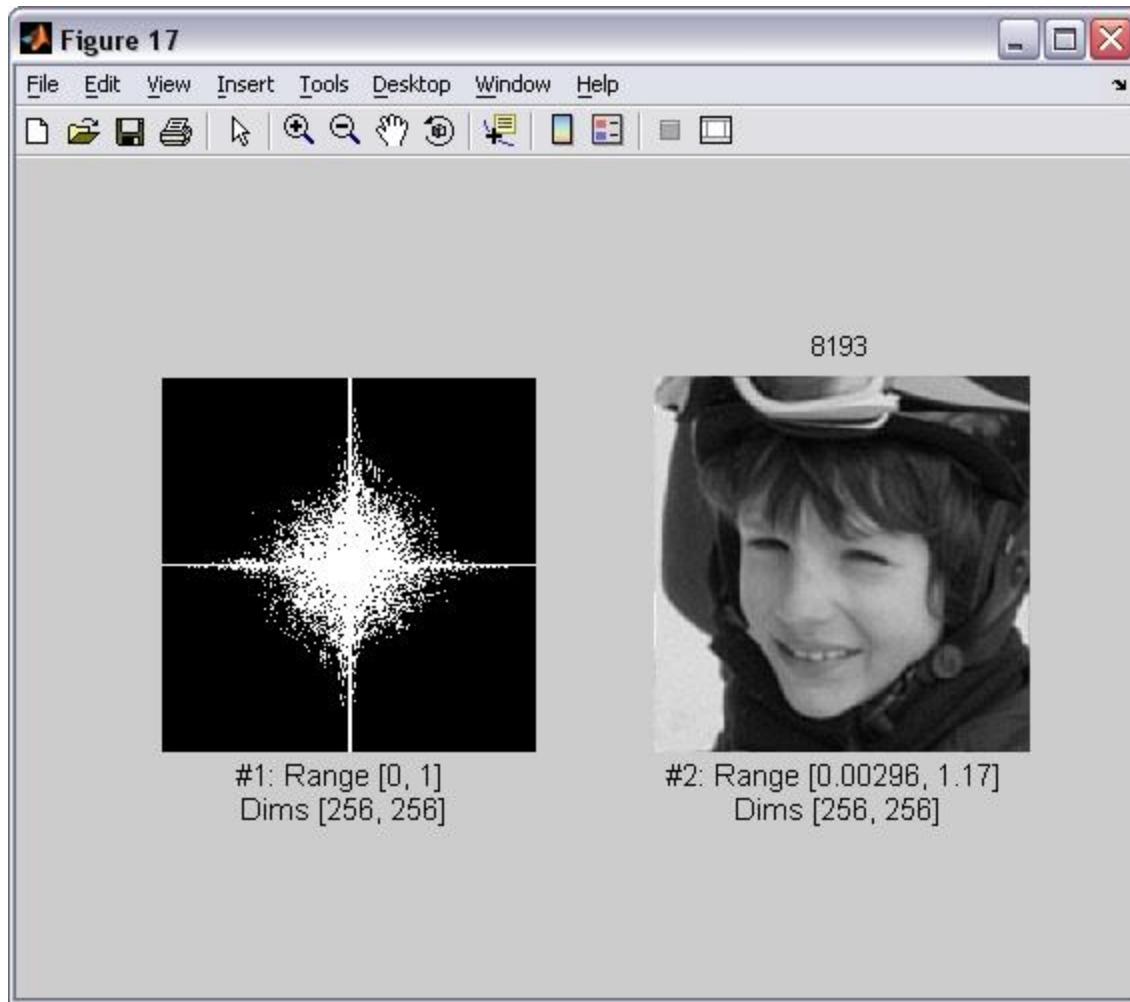
# 2049



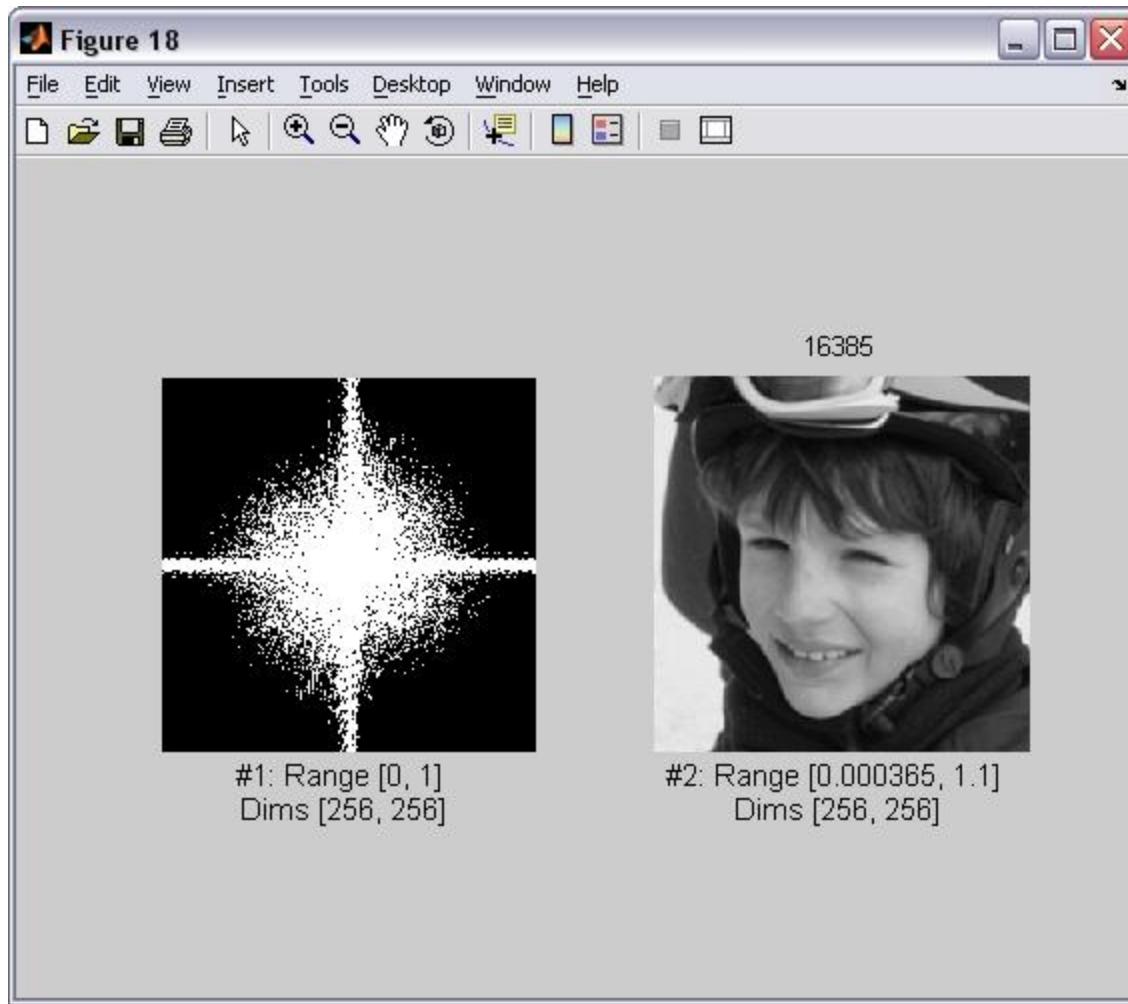
# 4097



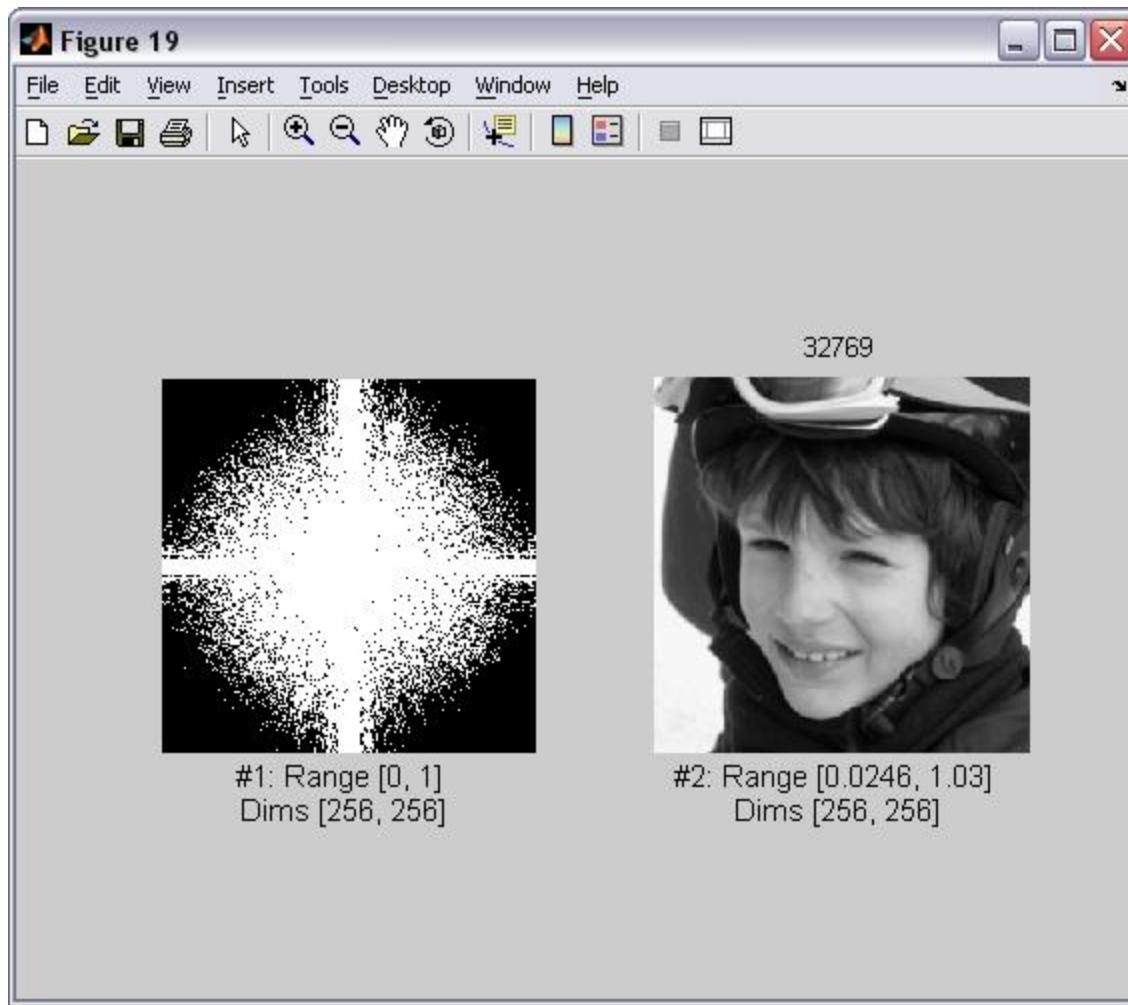
# 8193



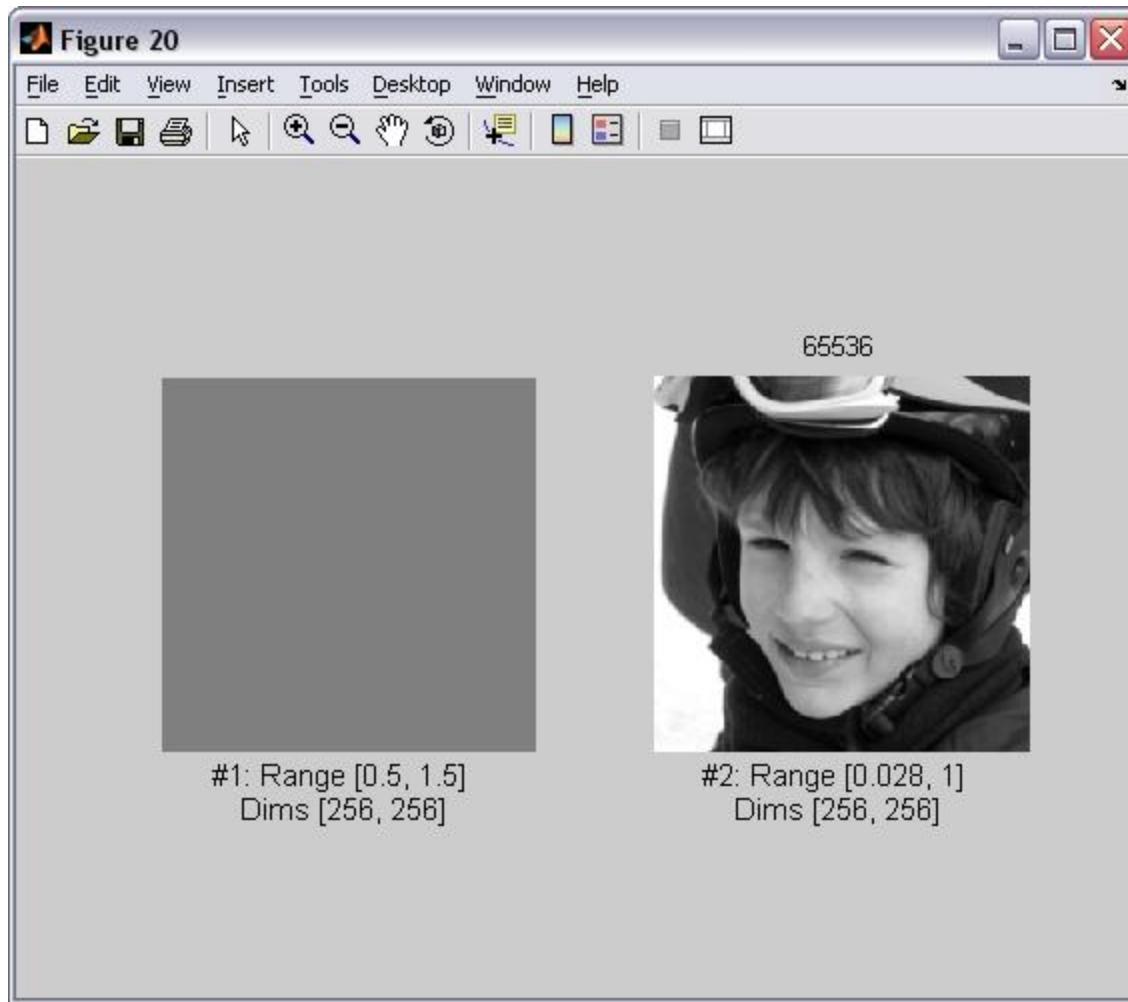
# 16385



# 32769

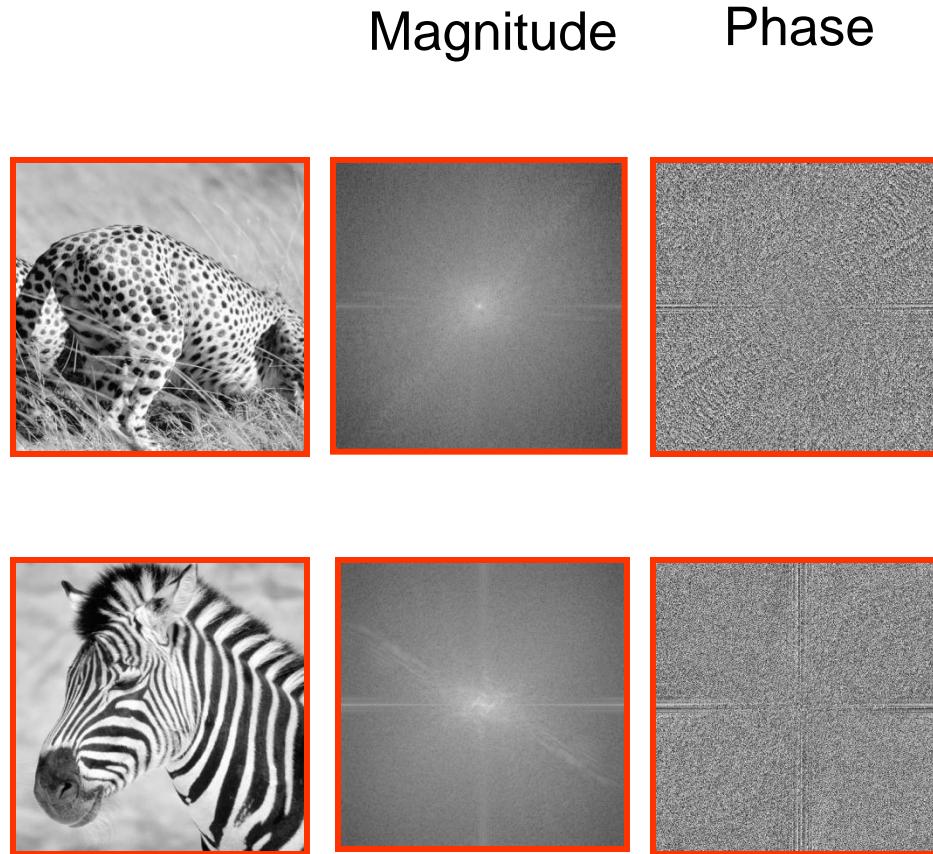


# 65536



# Fourier Transform

- Fourier transform of a real function is complex
  - difficult to plot, visualize
  - instead, we can think of the phase and magnitude of the transform
- Phase is the phase of the complex transform
- Magnitude is the magnitude of the complex transform

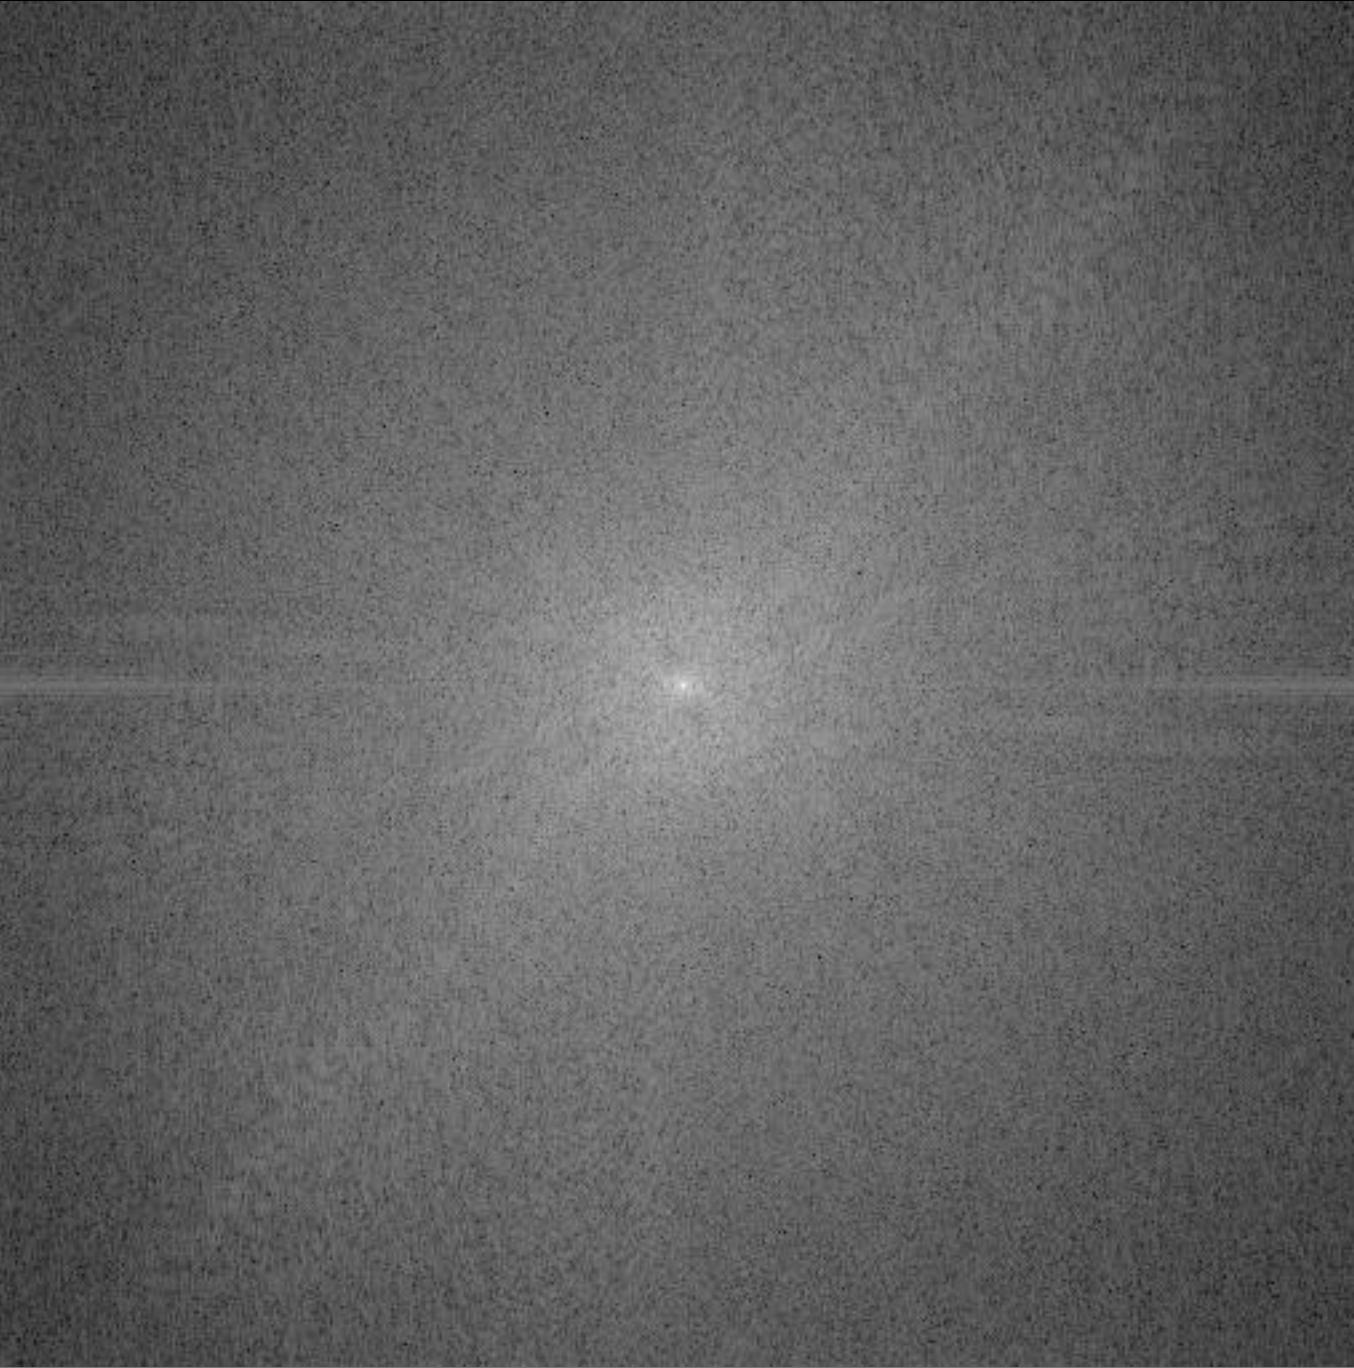


# Phase and Magnitude

- Curious fact
  - all natural images have about the same magnitude transform
  - hence, phase seems to matter, but magnitude largely doesn't
- Demonstration
  - Take two pictures, swap the phase transforms, compute the inverse - what does the result look like?



This is the  
magnitude  
transform of  
the cheetah  
pic

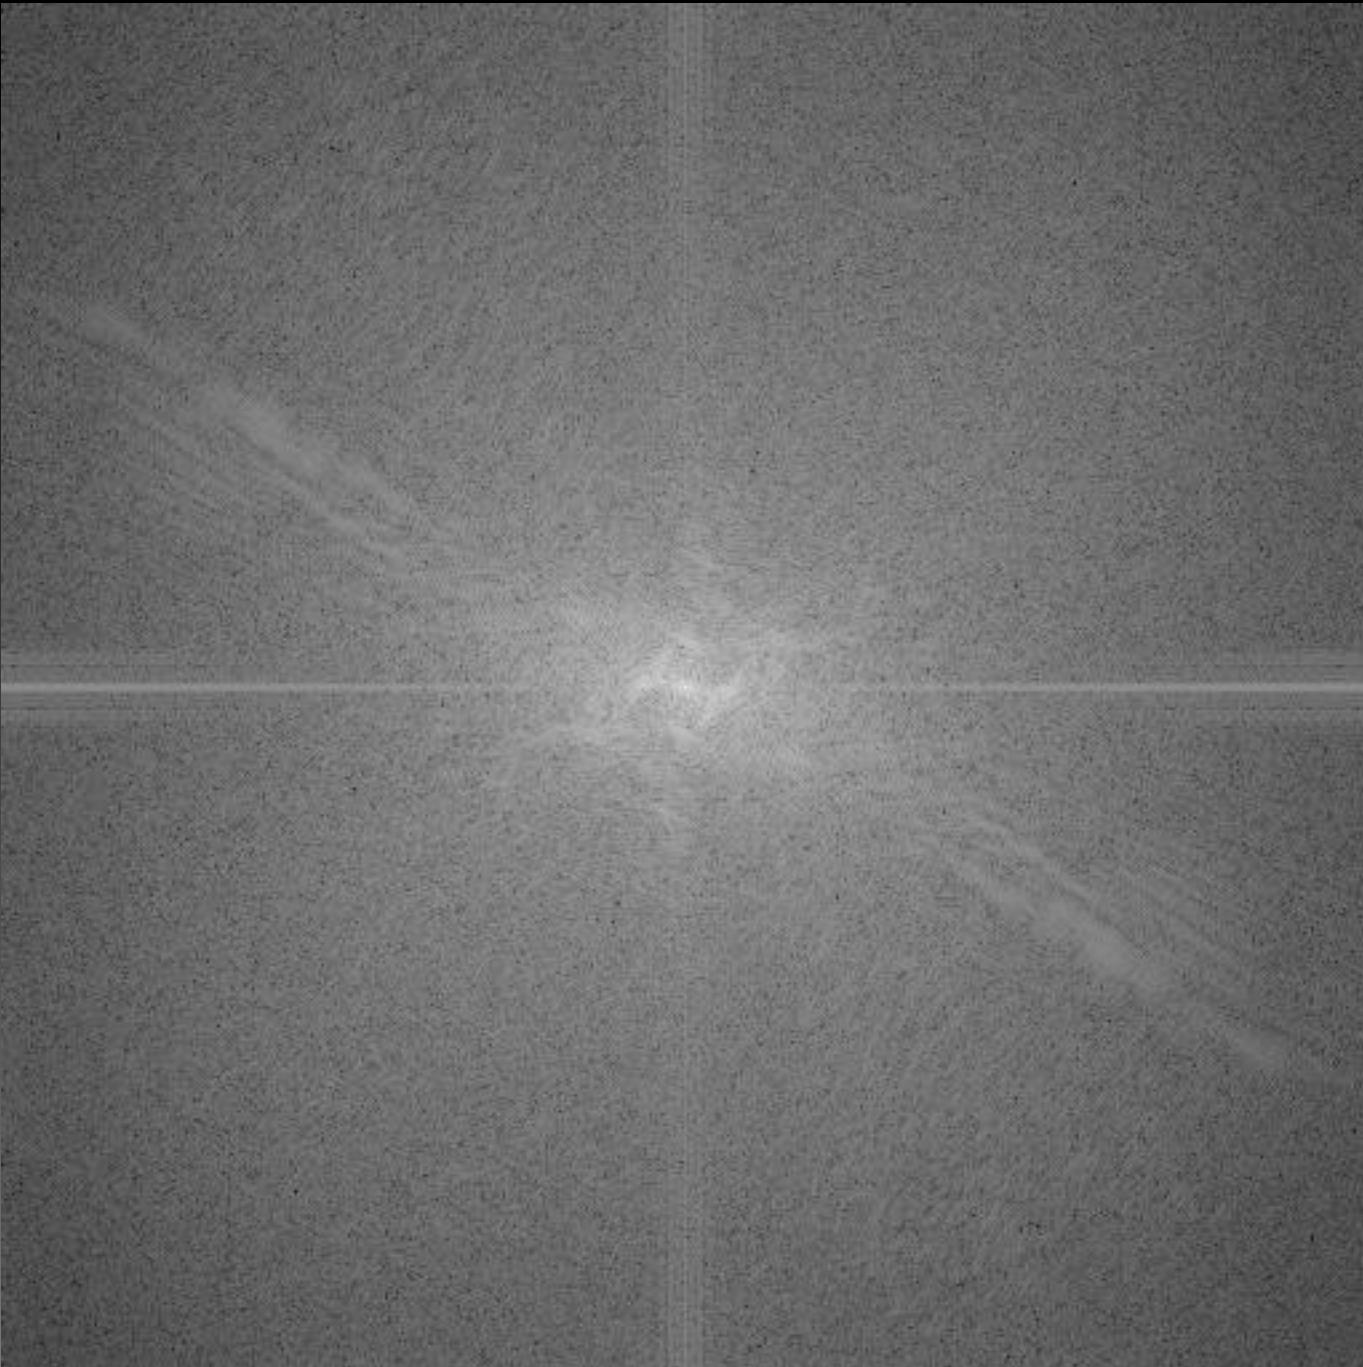


This is the  
phase  
transform  
of the  
cheetah pic

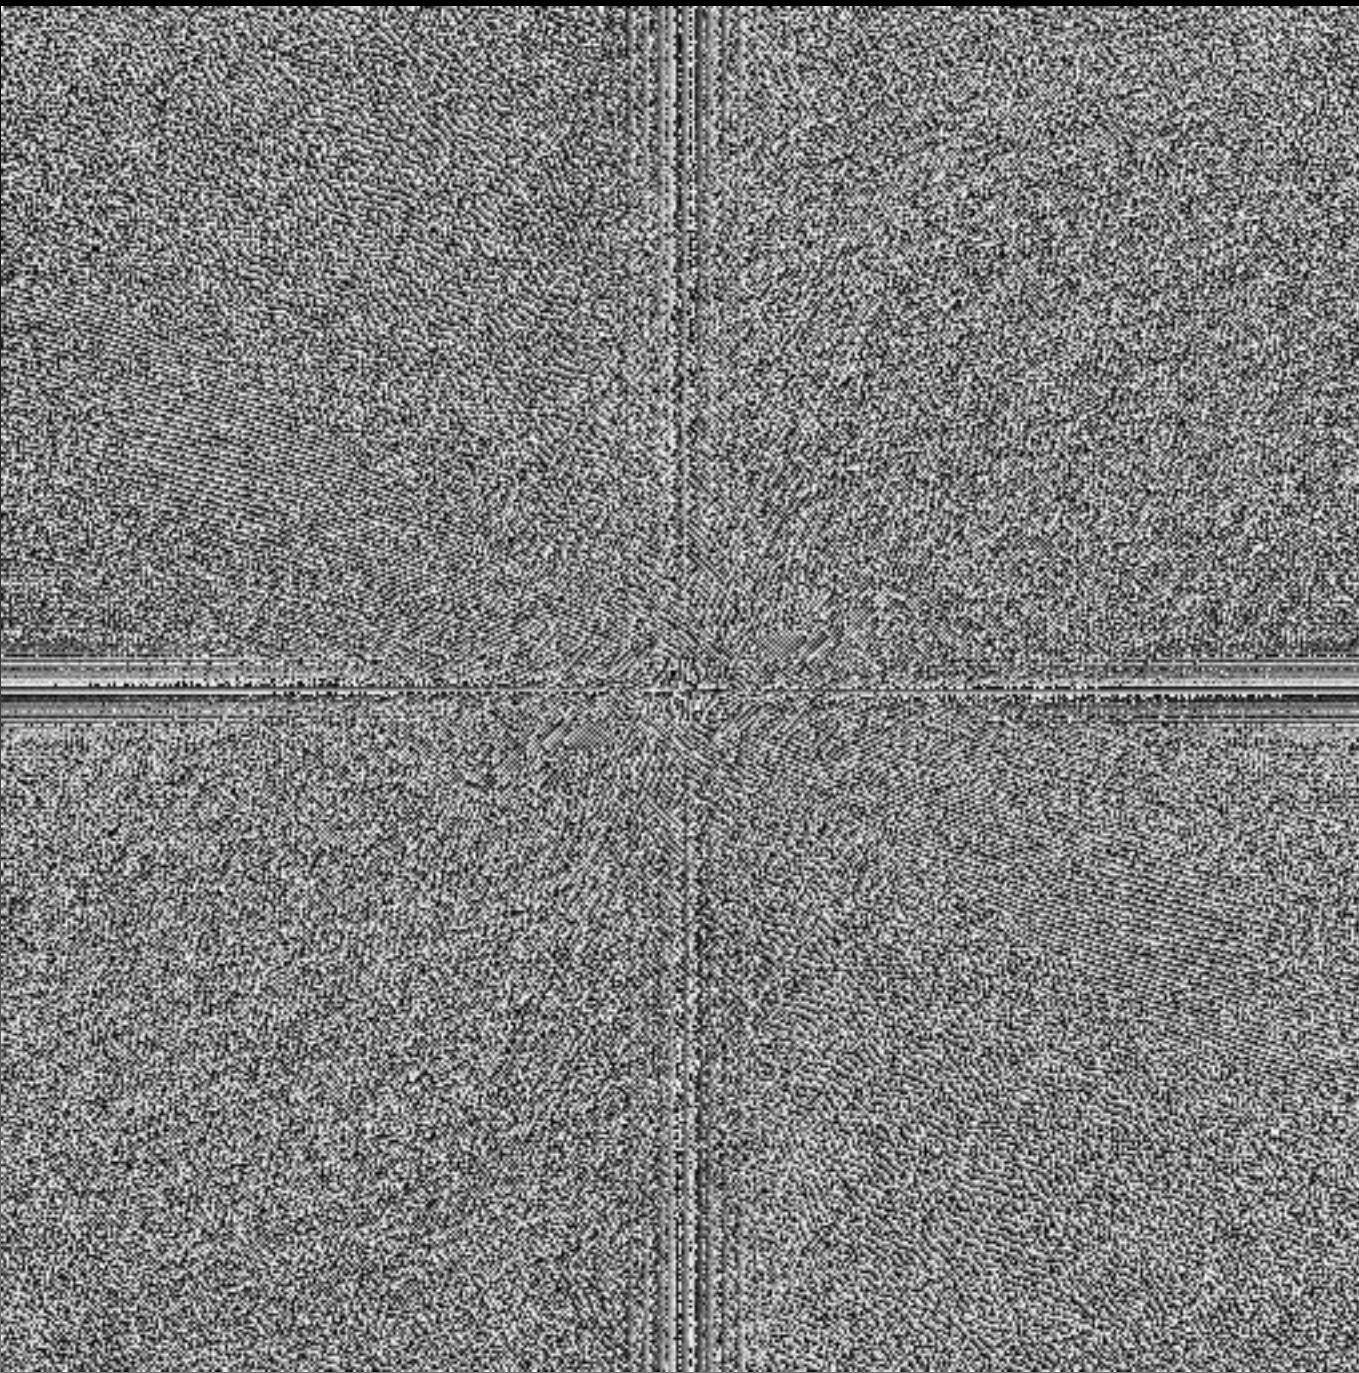




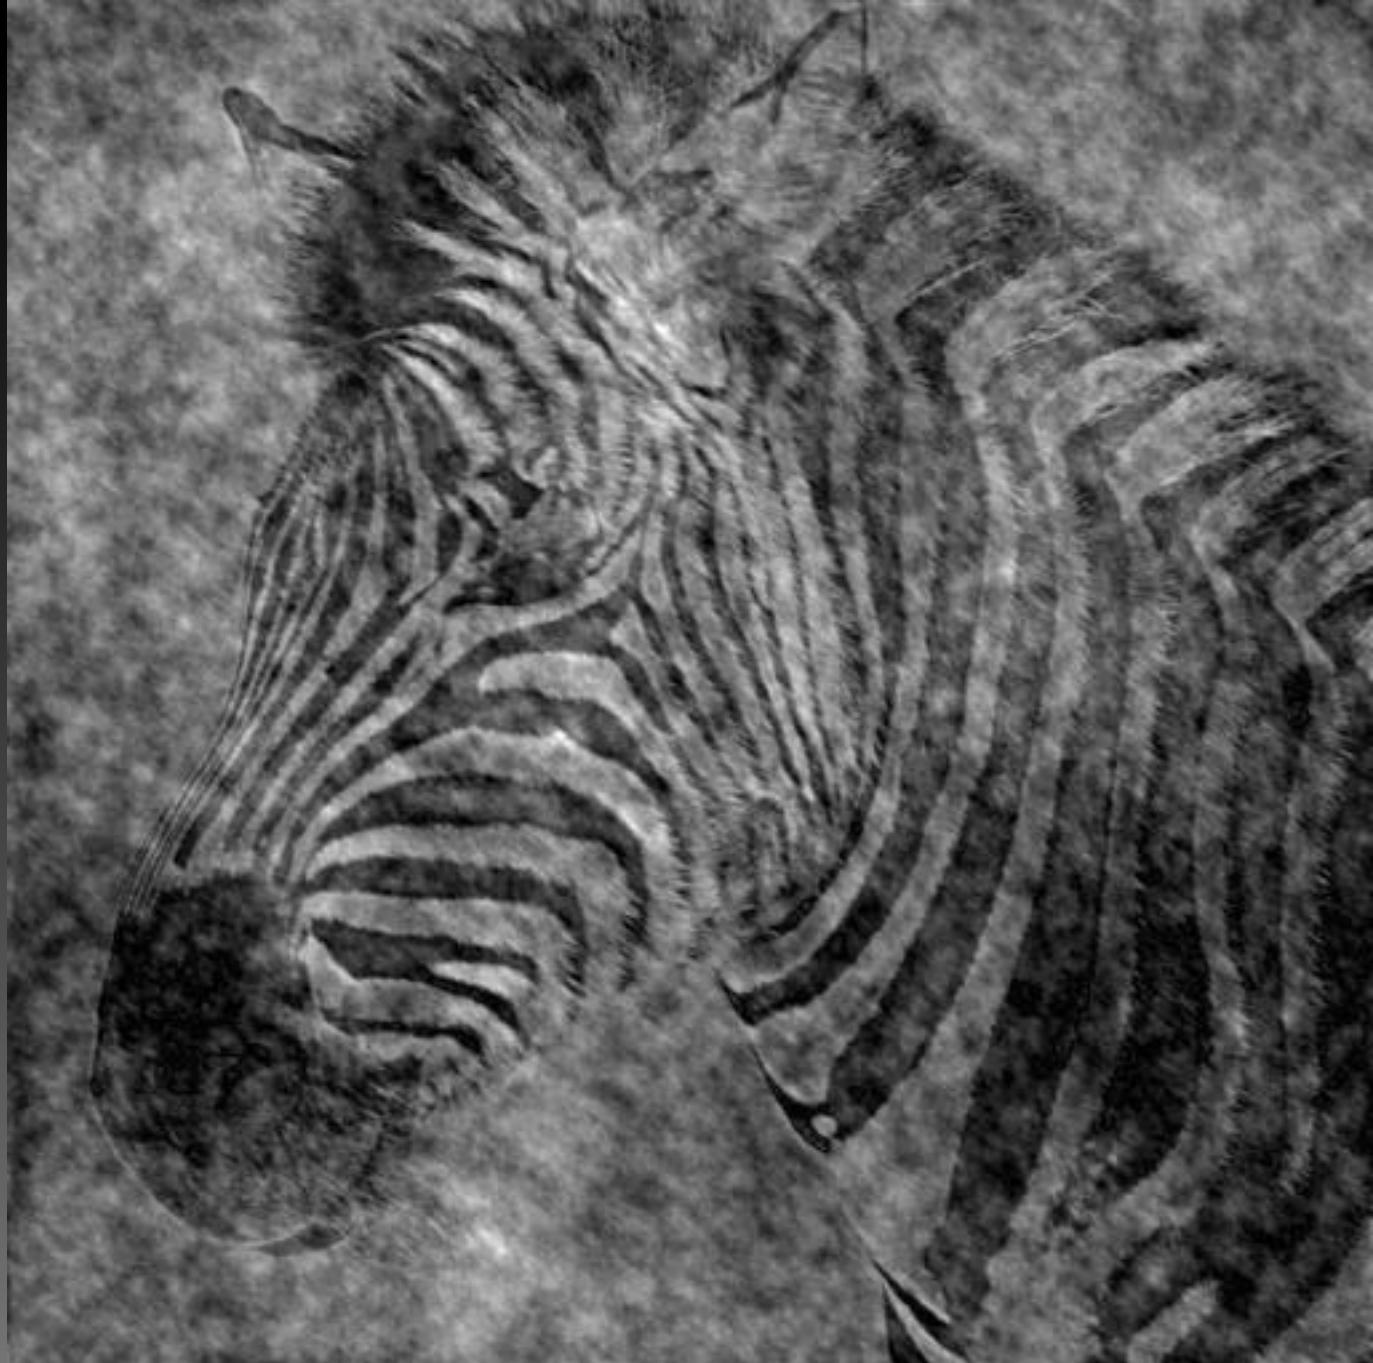
This is the  
magnitude  
transform  
of the zebra  
pic



This is the  
phase  
transform  
of the zebra  
pic



Reconstruction  
with zebra  
phase, cheetah  
magnitude



Reconstruction  
with cheetah  
phase, zebra  
magnitude



# Phase and Magnitude



- The magnitude of natural images can often be quite similar, one to another. But magnitude **encodes statistics of orientation at all spatial scales**.
- The phase carry the information of **where the image contours are**, by specifying how the phases of the sinusoids must line up in order to create the observed contours and edges.

Image with cheetah phase  
(and zebra magnitude)

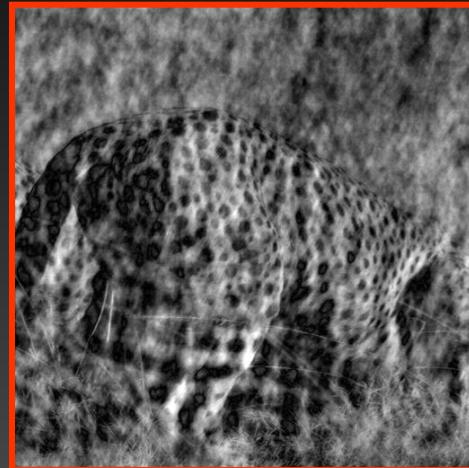
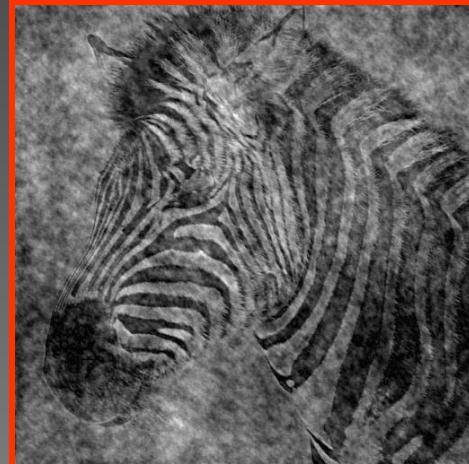
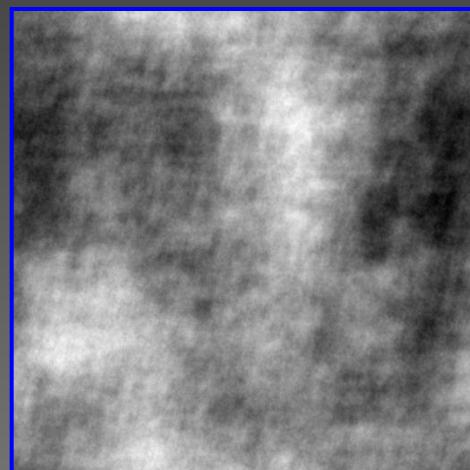
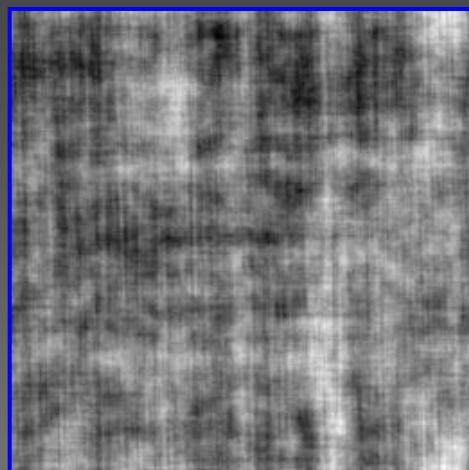
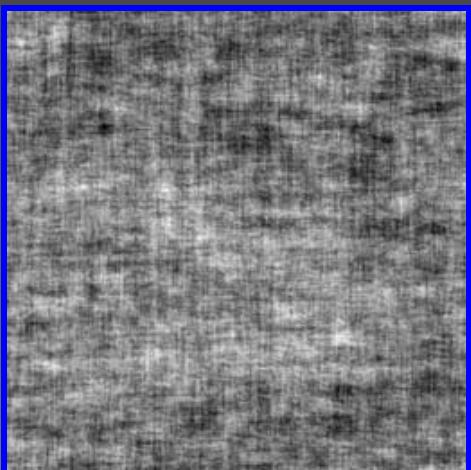


Image with zebra phase  
(and cheetah magnitude)

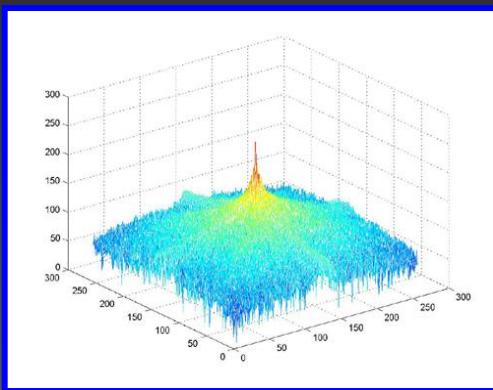


# Randomizing the phase



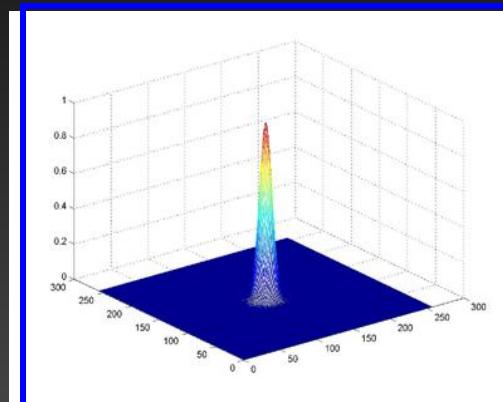
# Use of Fourier Spectrum : Filtering

Fourier space

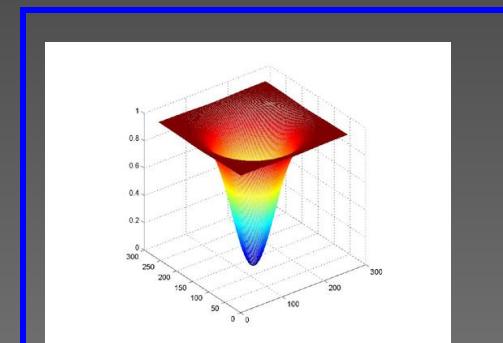


\*

Low pass Filter

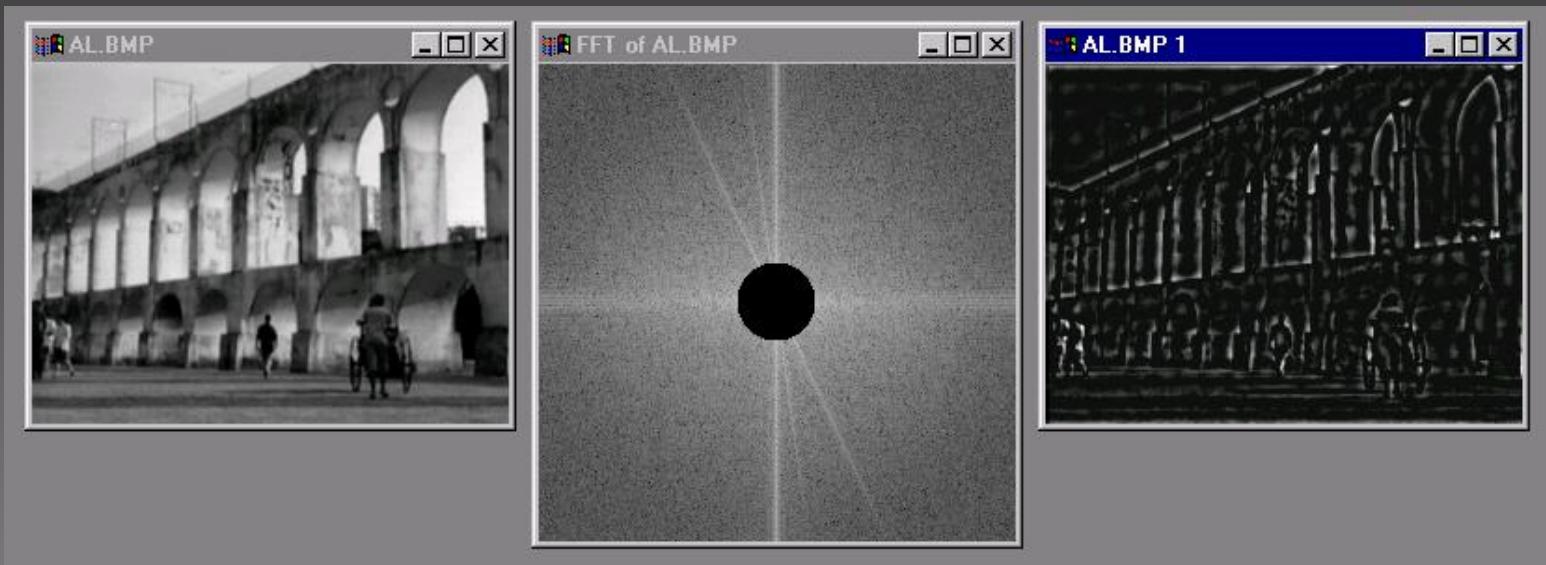


High pass filter



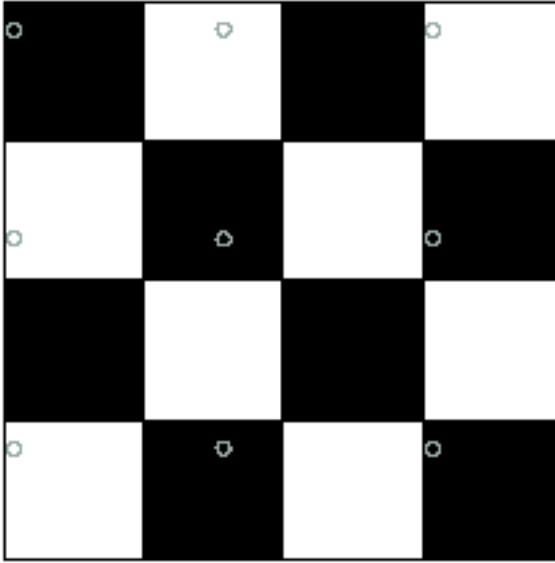
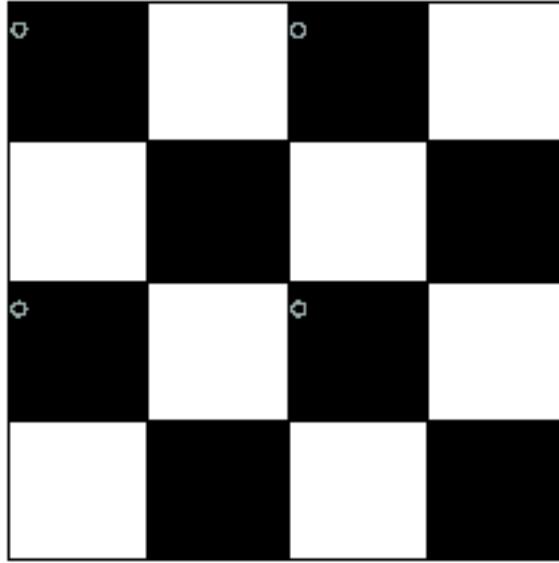
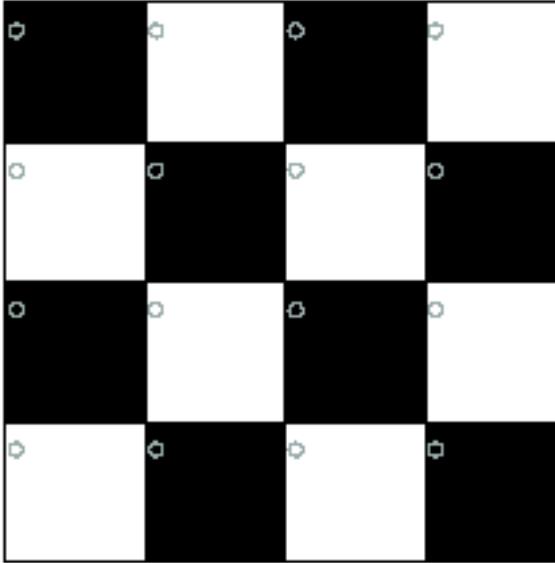
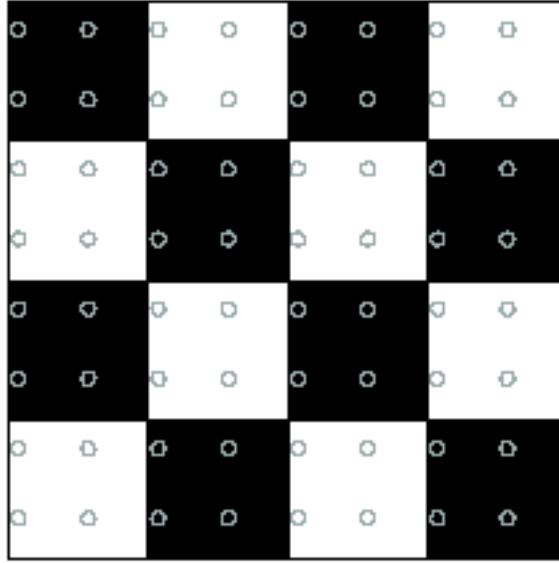
\*

# Low and High Pass filtering



# The Fourier transform of a sampled signal

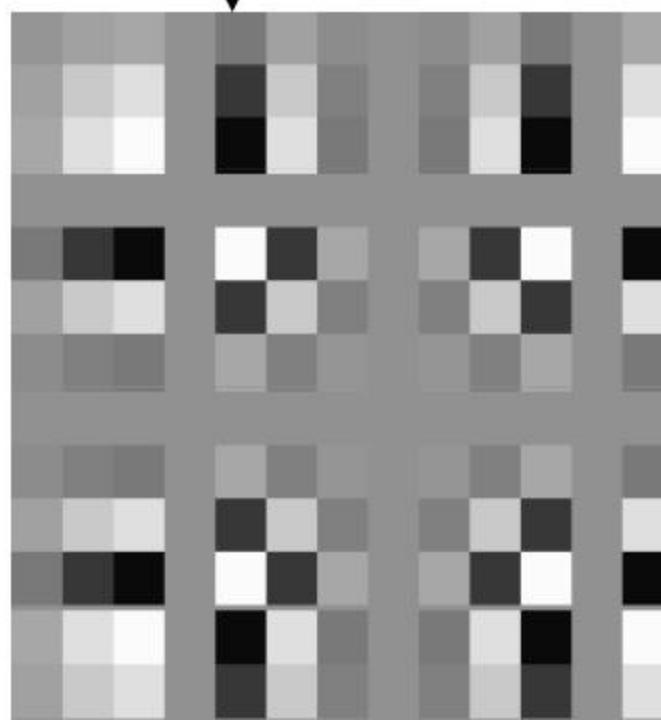
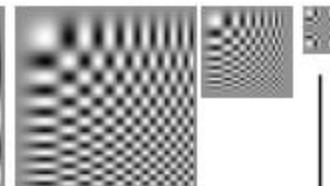
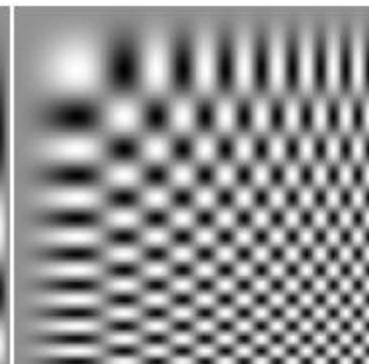
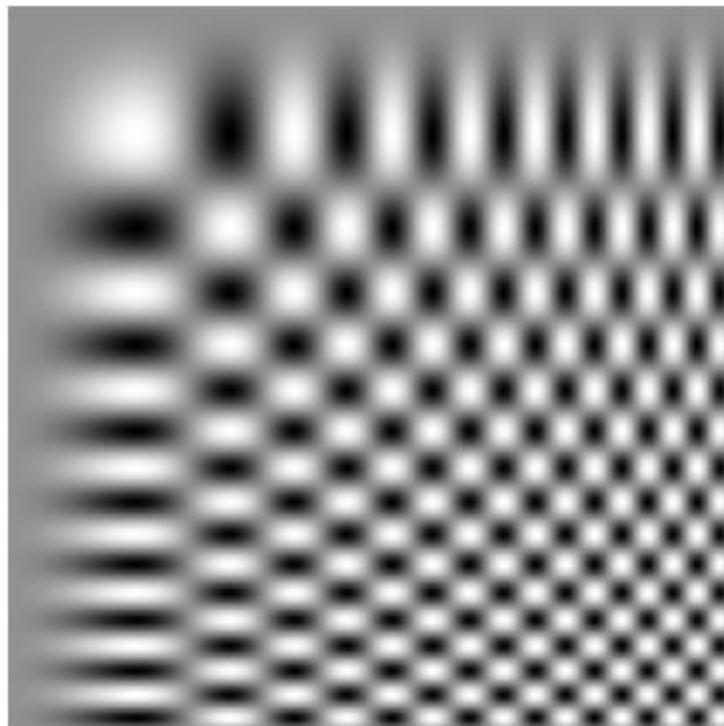
$$\begin{aligned} F(\text{Sample}_{2D}(f(x,y))) &= F\left(f(x,y) \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} \delta(x-i, y-j)\right) \\ &= F(f(x,y))^{**} F\left(\sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} \delta(x-i, y-j)\right) \\ &= \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} F(u-i, v-j) \end{aligned}$$



Resample the checkerboard by taking one sample at each circle. In the case of the top left board, new representation is reasonable.

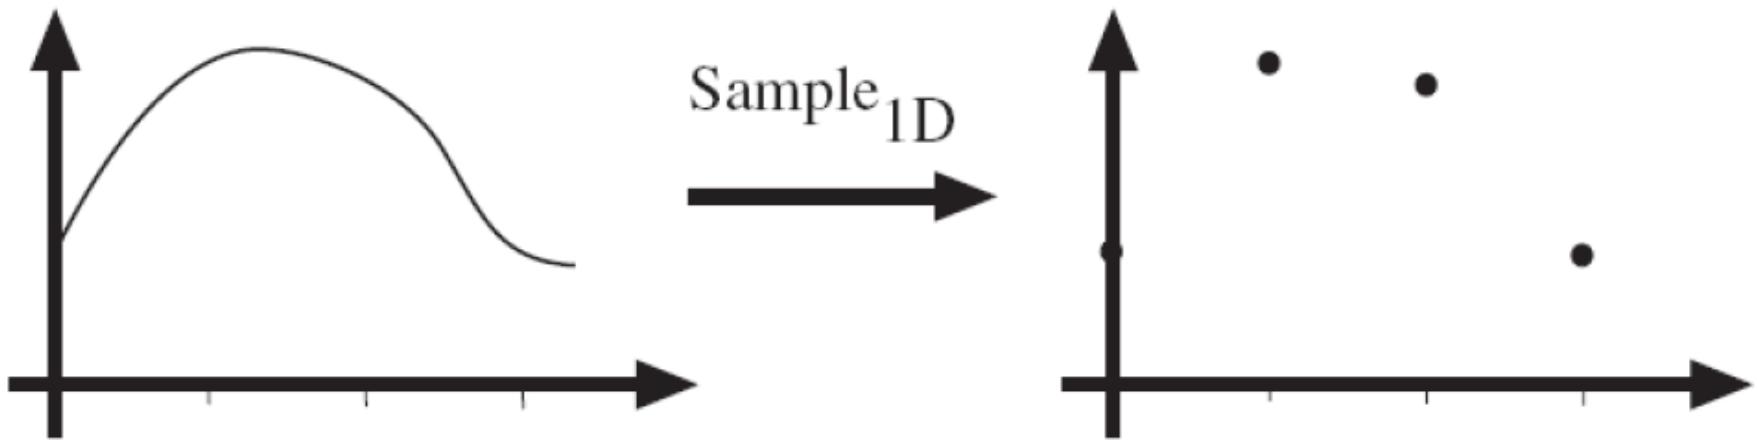
Top right also yields a reasonable representation. Bottom left is all black (dubious) and bottom right has checks that are too big.

# Sampling and aliasing



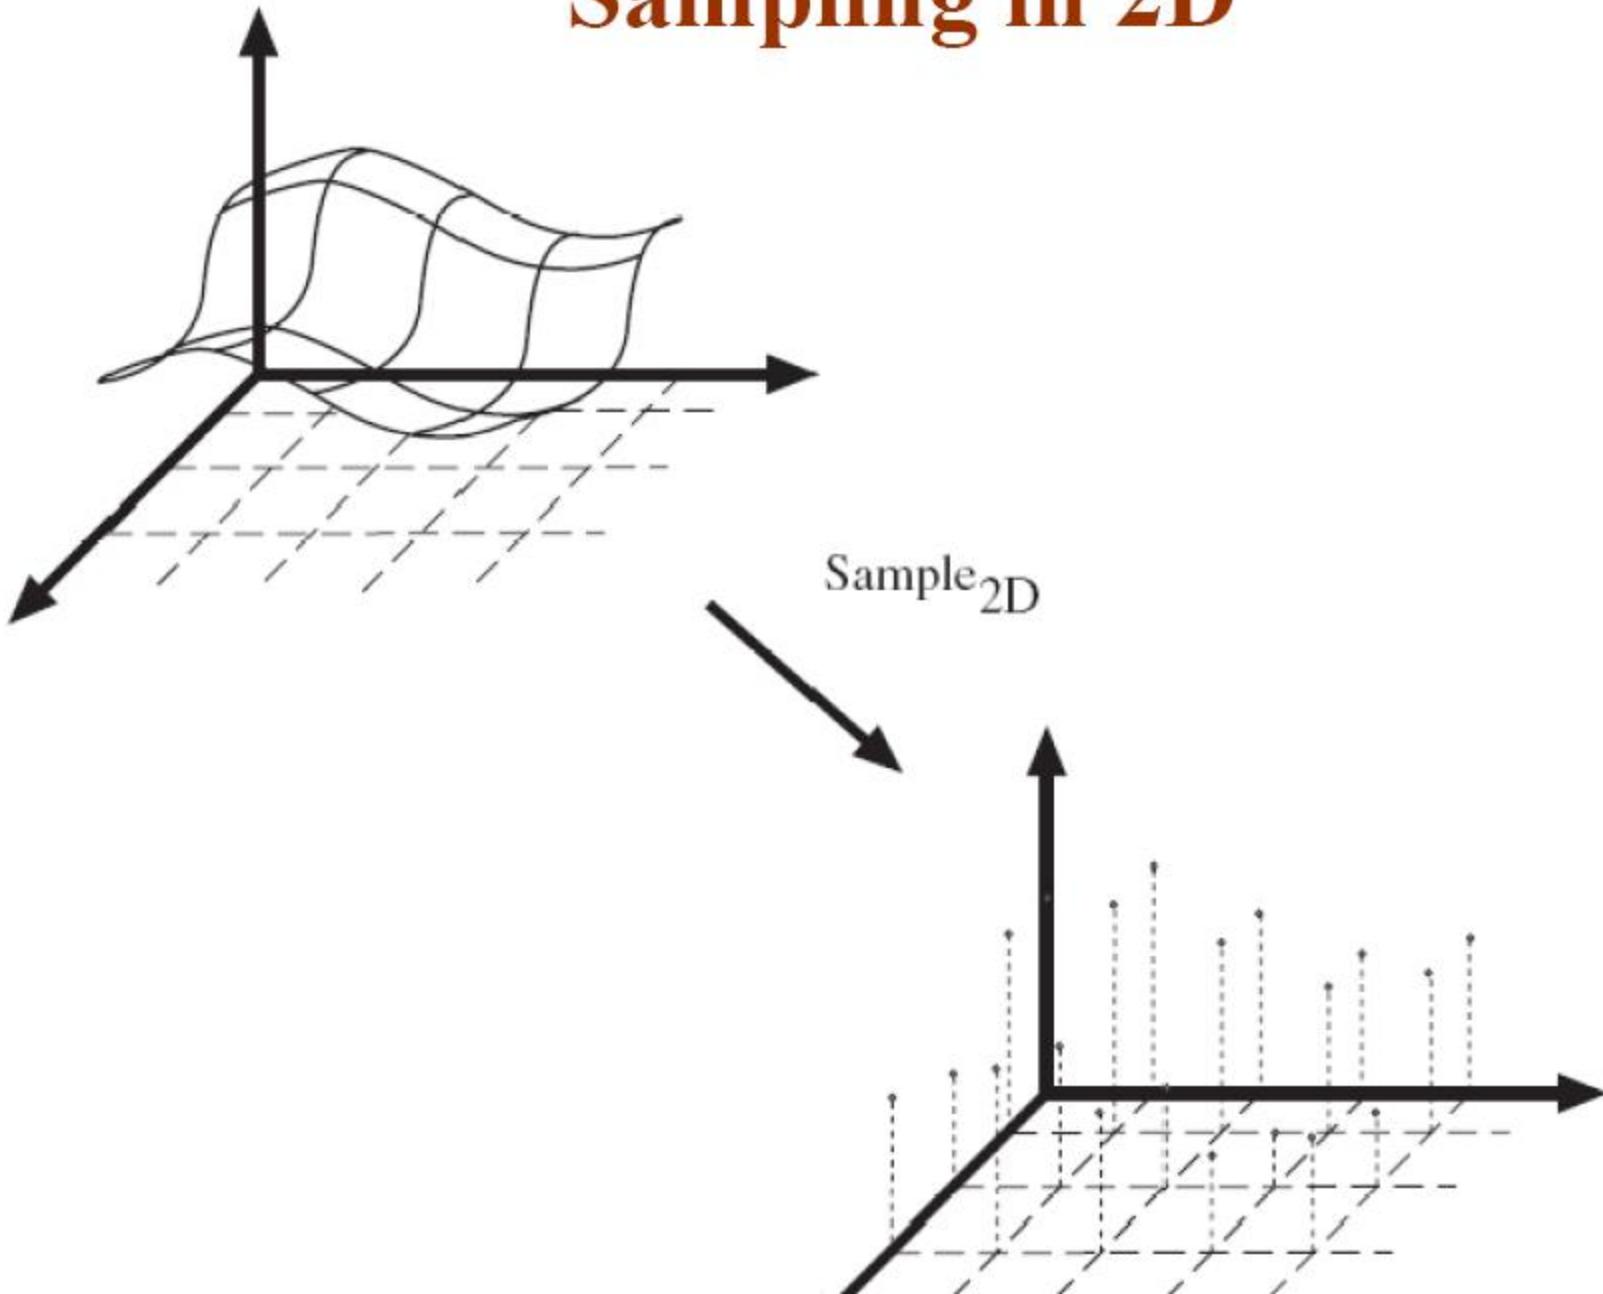
Constructing a pyramid by taking every second pixel leads to layers that badly misrepresent the top layer

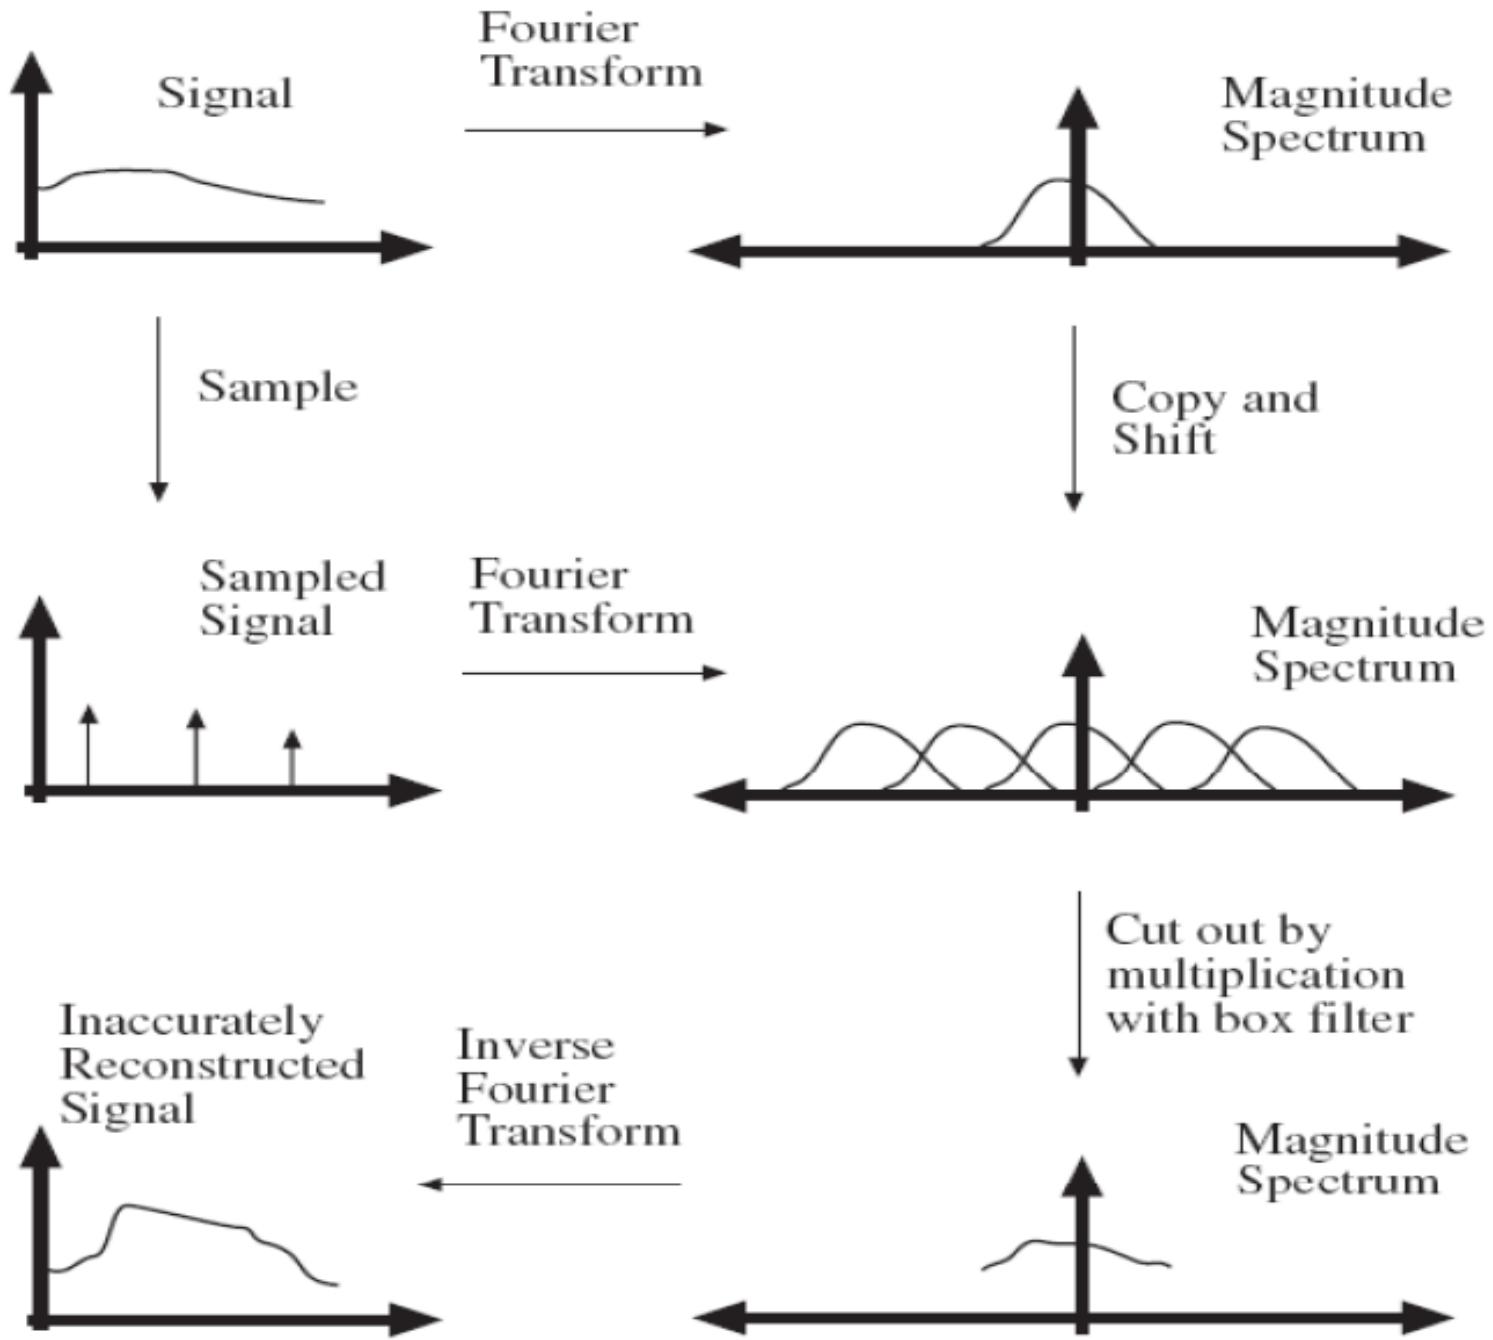
# Sampling in 1D

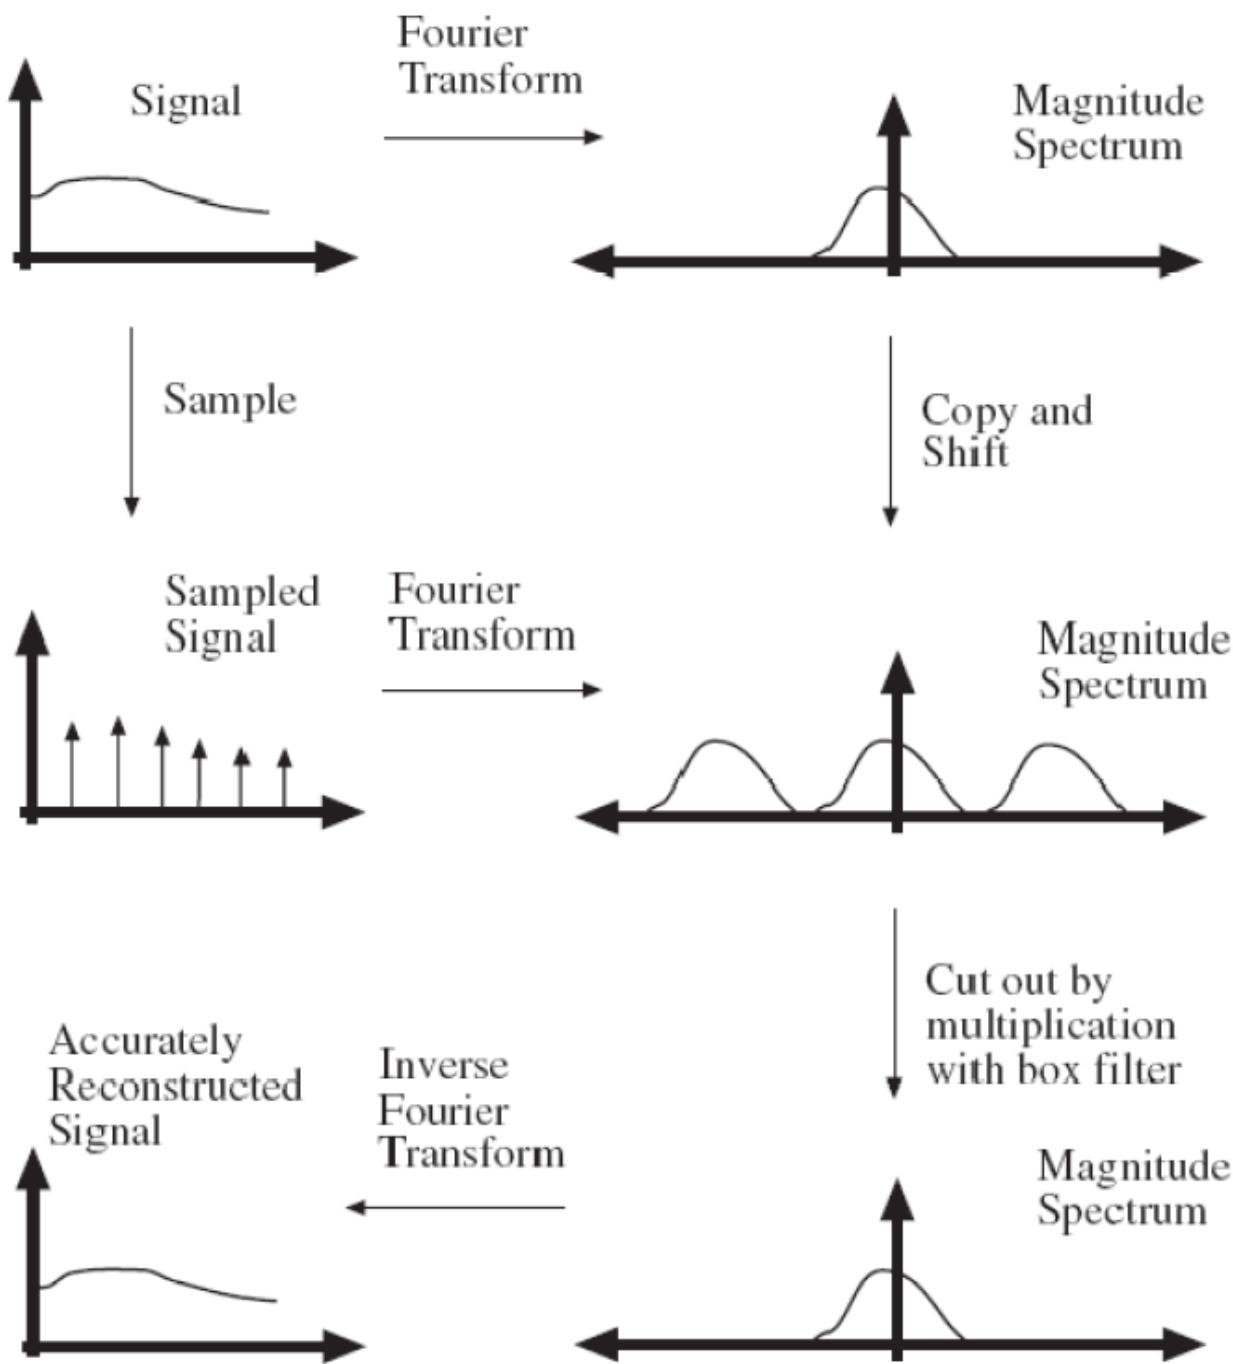


**FIGURE 7.8:** Sampling in 1D takes a function, and returns a vector whose elements are values of that function at the sample points, as the top figures show. For our purposes, it is enough that the sample points be integer values of the argument. We allow the vector to be infinite dimensional, and have negative as well as positive indices.

# Sampling in 2D





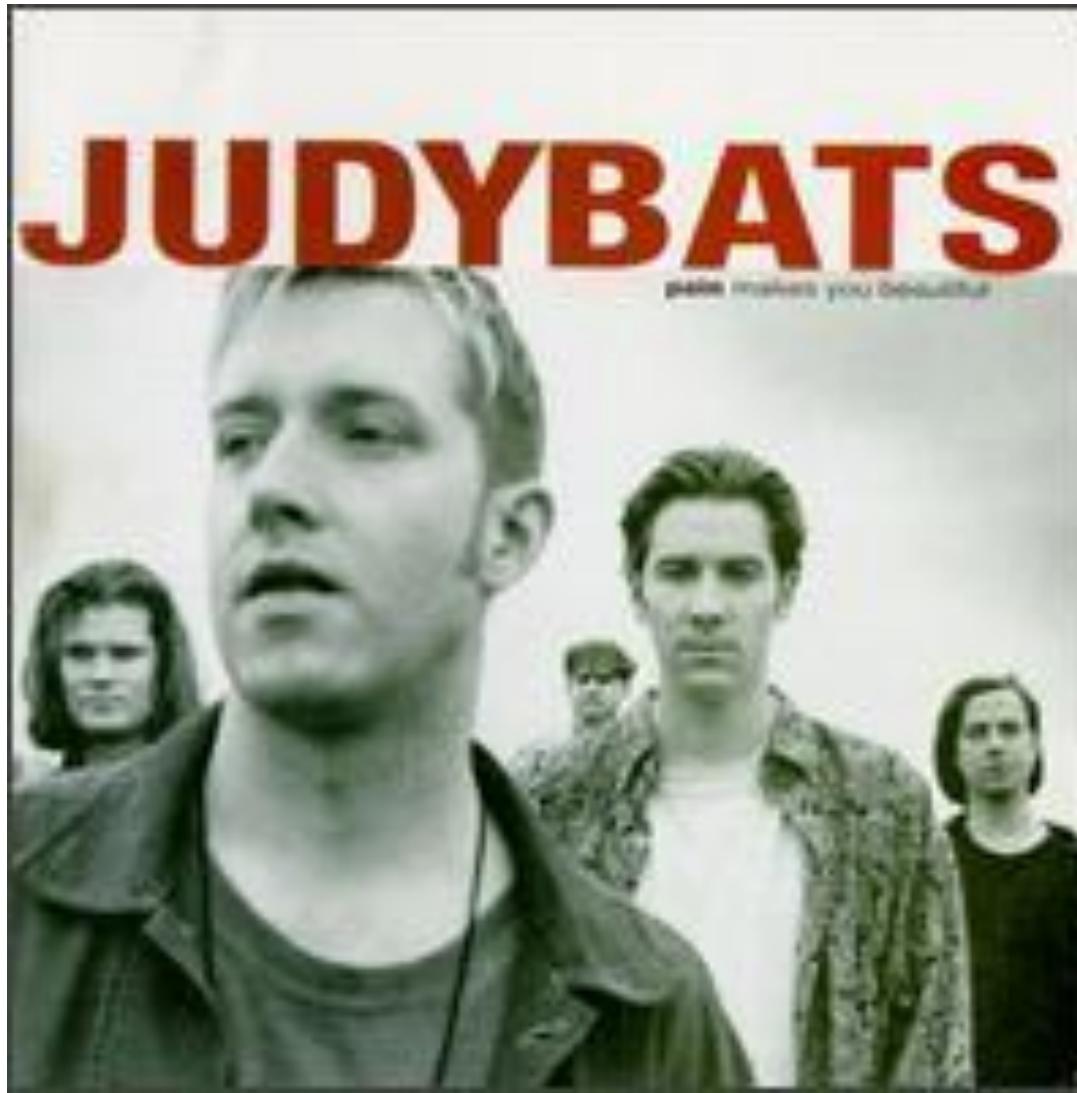


# Multiscale representations

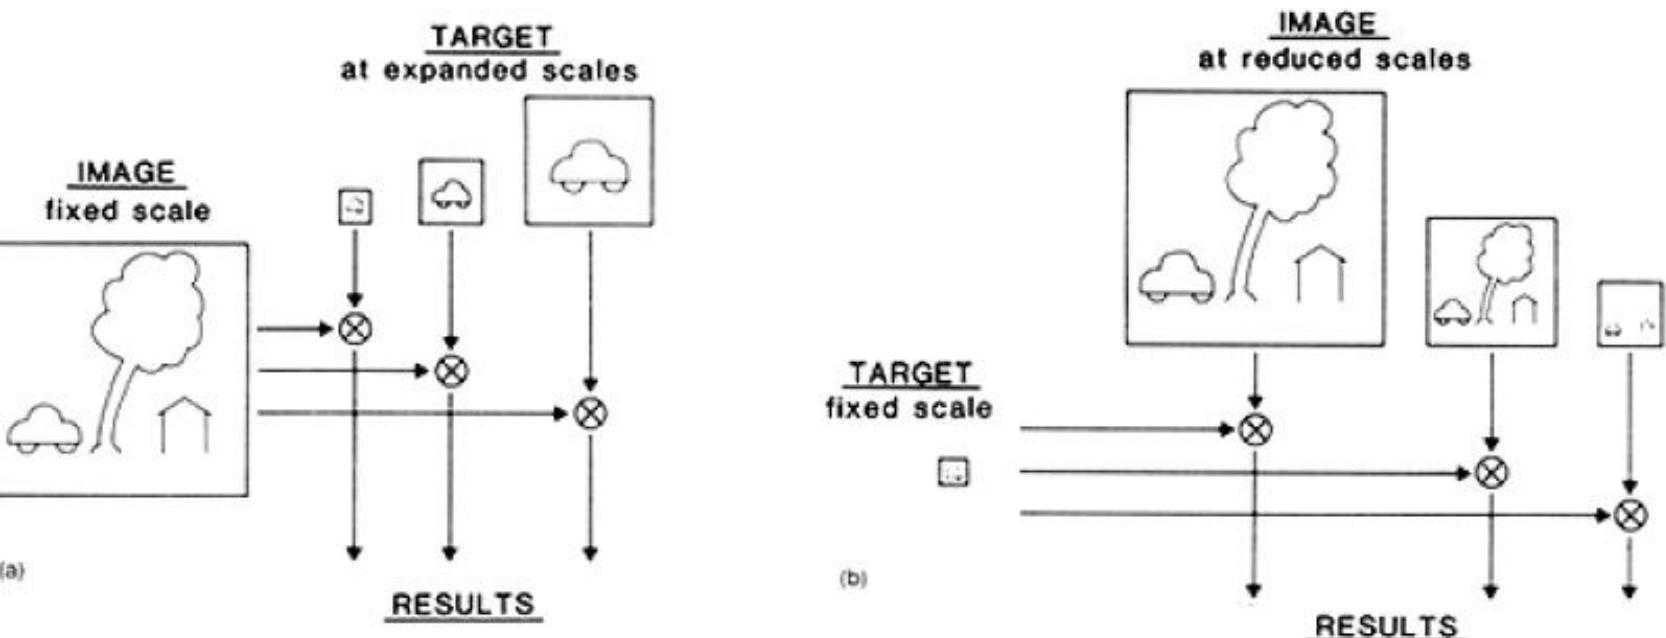
# Multiscale representations

- Find template matches at all scales
  - e.g., when finding hands or faces, we don't know what size they will be in a particular image
  - Template size is constant, but image size changes
- Efficient search for correspondence
  - look at coarse scales, then refine with finer scales
  - much less cost, but may miss best match
- Examining all levels of detail
  - Find edges with different amounts of blur
  - Find textures with different spatial frequencies (levels of detail)

# Image information occurs at all spatial scales



# Scale invariance



**Fig. 1.** Two methods of searching for a target pattern over many scales. In the first approach, (a), copies of the target pattern are constructed at several expanded scales, and each is convolved with the original image. In the second approach, (b), a single copy of the target is convolved with

copies of the image reduced in scale. The target should be just large enough to resolve critical details. The two approaches should give equivalent results, but the second is more efficient by the fourth power of the scale factor (image convolutions are represented by 'O').



# The Gaussian pyramid

- Create each level from previous one:
  - smooth and sample
- Smooth with gaussians, because
  - a gaussian\*gaussian=another gaussian
- Gaussians are low pass filters, so representation is redundant.



## GAUSSIAN PYRAMID



0

1

2

3

4

5

Fig. 4. First six levels of the Gaussian pyramid for the "Lady" image. The original image, level 0, measures 257 by 257 pixels and each higher level array is roughly half the dimensions of its predecessor. Thus, level 5 measures just 9 by 9 pixels.



512

256

128

64

32

16

8



# Some desirable properties for a blur kernel

- **Positivity:**  $h(m) \geq 0$
- **Symmetry:**  $h(m) = h(-m)$
- **Unimodality:**  $h(m) \geq h(m+1)$  for  $m \geq 0$
- **Normalized:**  $\sum h(m) = 1$
- **Equal contribution:**  $\sum h(2m) = \sum h(2m+1)$

Some kernels that verify this are:

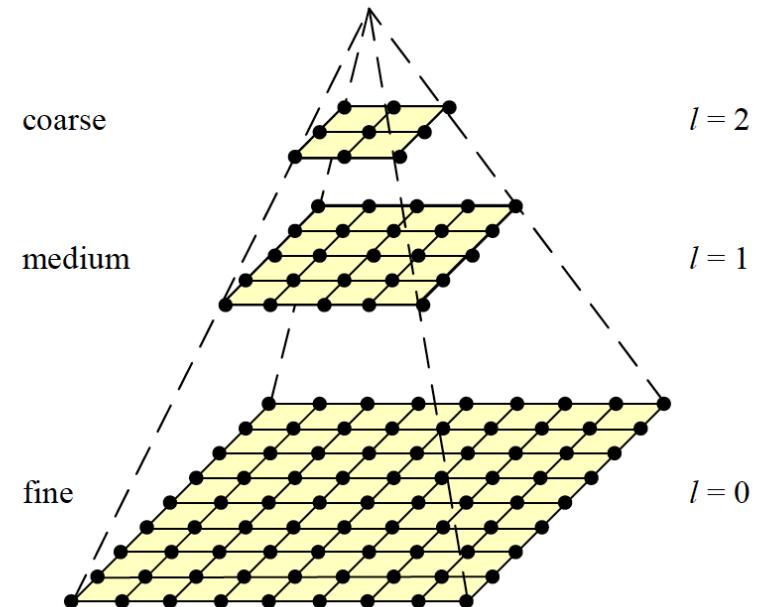
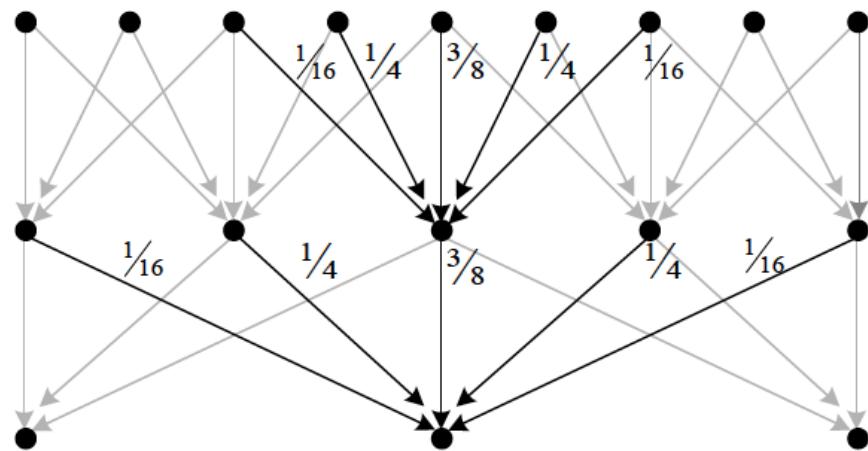
$$[\frac{1}{2}, \frac{1}{2}]$$

For length 5:

$$[\frac{1}{4}, \frac{1}{2}, \frac{1}{4}] \quad \frac{1}{4}-a/2, \quad \frac{1}{4}, \quad a, \quad \frac{1}{4}, \quad \frac{1}{4}-a/2]$$

The value of **a** that best approximates a Gaussian filter is 0.4

# Gaussian Pyramid



## Convolution and subsampling as a matrix multiply (1-d case)

$$x_2 = G_1 x_1$$

$$\begin{matrix} G_1 = \\ \begin{array}{cccccccccccccccccccc} 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \end{array} \end{matrix}$$

(Normalization constant of 1/16 omitted for visual clarity.)

# Next pyramid level

$$x_3 = G_2 x_2$$

$$G_2 = \begin{matrix} & 1 & 4 & 6 & 4 & 1 & 0 & 0 & 0 \\ & 0 & 0 & 1 & 4 & 6 & 4 & 1 & 0 \\ & 0 & 0 & 0 & 0 & 1 & 4 & 6 & 4 \\ & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 \end{matrix}$$

# The combined effect of the two pyramid levels

$$x_3 = G_2 G_1 x_1$$

$$G_2 G_1 = \begin{matrix} 1 & 4 & 10 & 20 & 31 & 40 & 44 & 40 & 31 & 20 & 10 & 4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 4 & 10 & 20 & 31 & 40 & 44 & 40 & 40 & 31 & 20 & 10 & 4 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 10 & 20 & 31 & 40 & 44 & 40 & 30 & 16 & 4 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 4 & 10 & 20 & 25 & 16 & 4 & 0 \end{matrix}$$

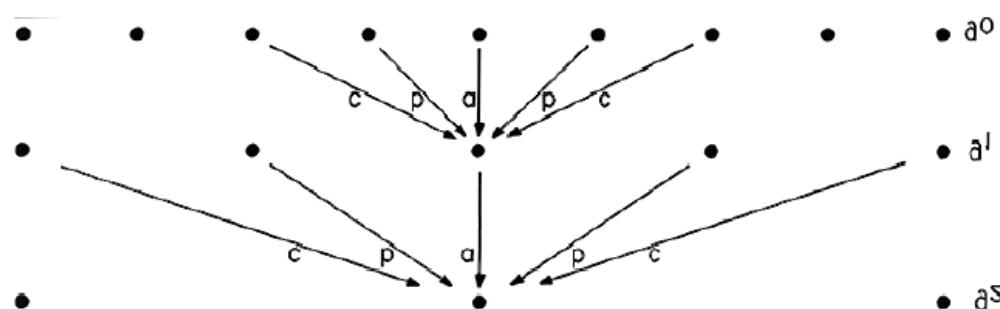
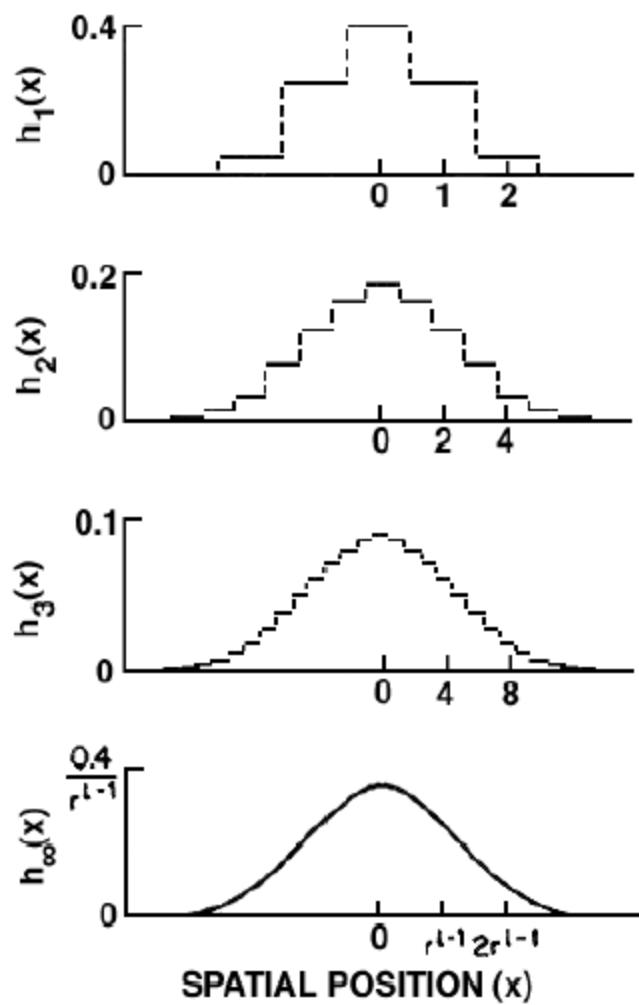


Fig. 2. The equivalent weighting functions  $h_i(x)$  for nodes in levels 1, 2, 3, and infinity of the Gaussian pyramid. Note that axis scales have been adjusted by factors of 2 to aid comparison. Here the parameter  $a$  of the generating kernel is 0.4, and the resulting equivalent weighting functions closely resemble the Gaussian probability density functions.

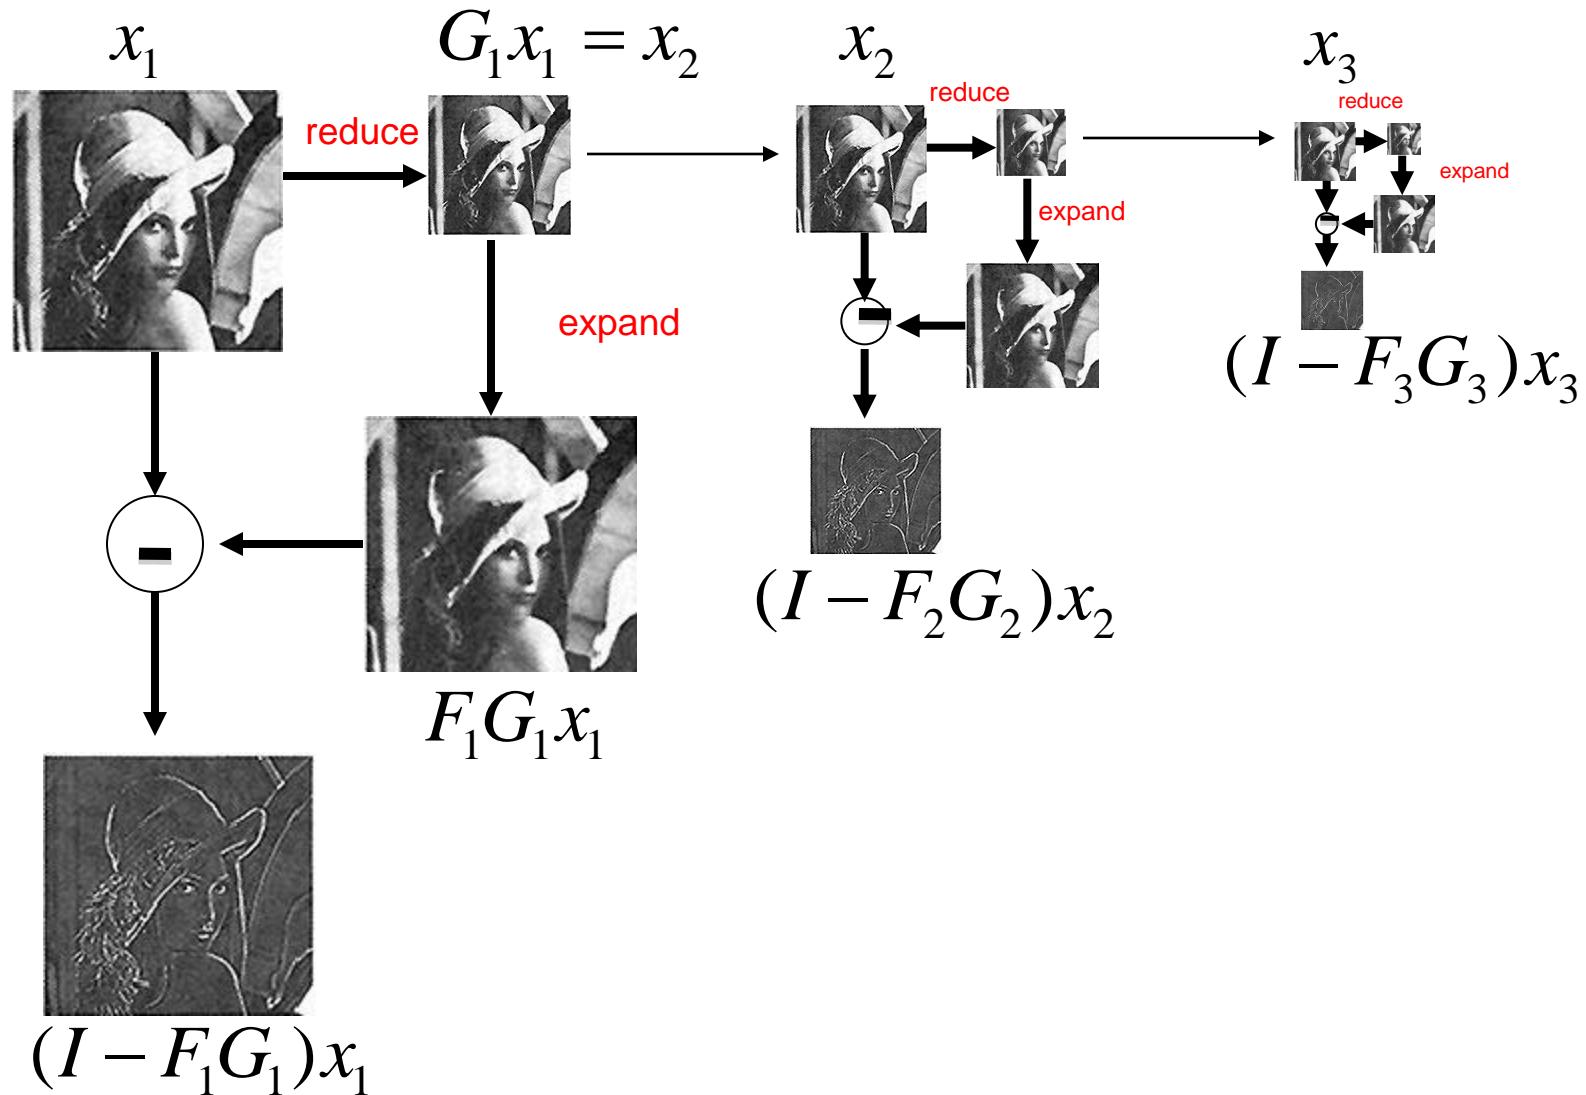
# Gaussian pyramids used for

- up- or down- sampling images.
- Multi-resolution image analysis
  - Look for an object over various spatial scales
  - Coarse-to-fine image processing: form blur estimate or the motion analysis on very low-resolution image, upsample and repeat. Often a successful strategy for avoiding local minima in complicated estimation tasks.

# The Laplacian Pyramid

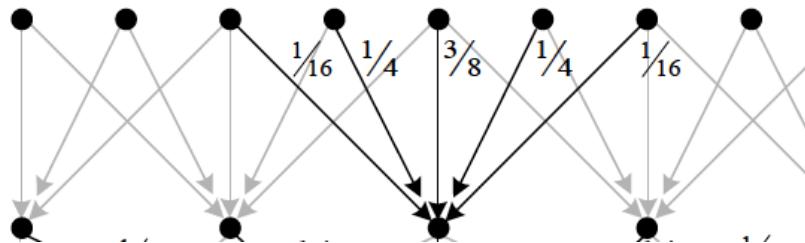
- **Synthesis**
  - Compute the difference between upsampled Gaussian pyramid level and Gaussian pyramid level.
  - band pass filter - each level represents spatial frequencies (largely) unrepresented at other level.
- **Applications:**
  - Texture synthesis; Image compression; Noise removal; Image blending

# Laplacian pyramid algorithm

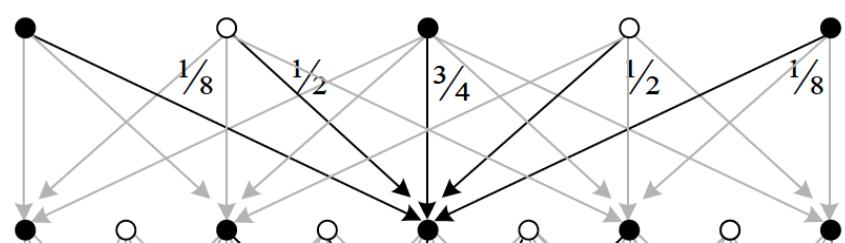


# Reduce and expand

Reduce



Expand



# Showing, at full resolution, the information captured at each level of a Gaussian (top) and Laplacian (bottom) pyramid.

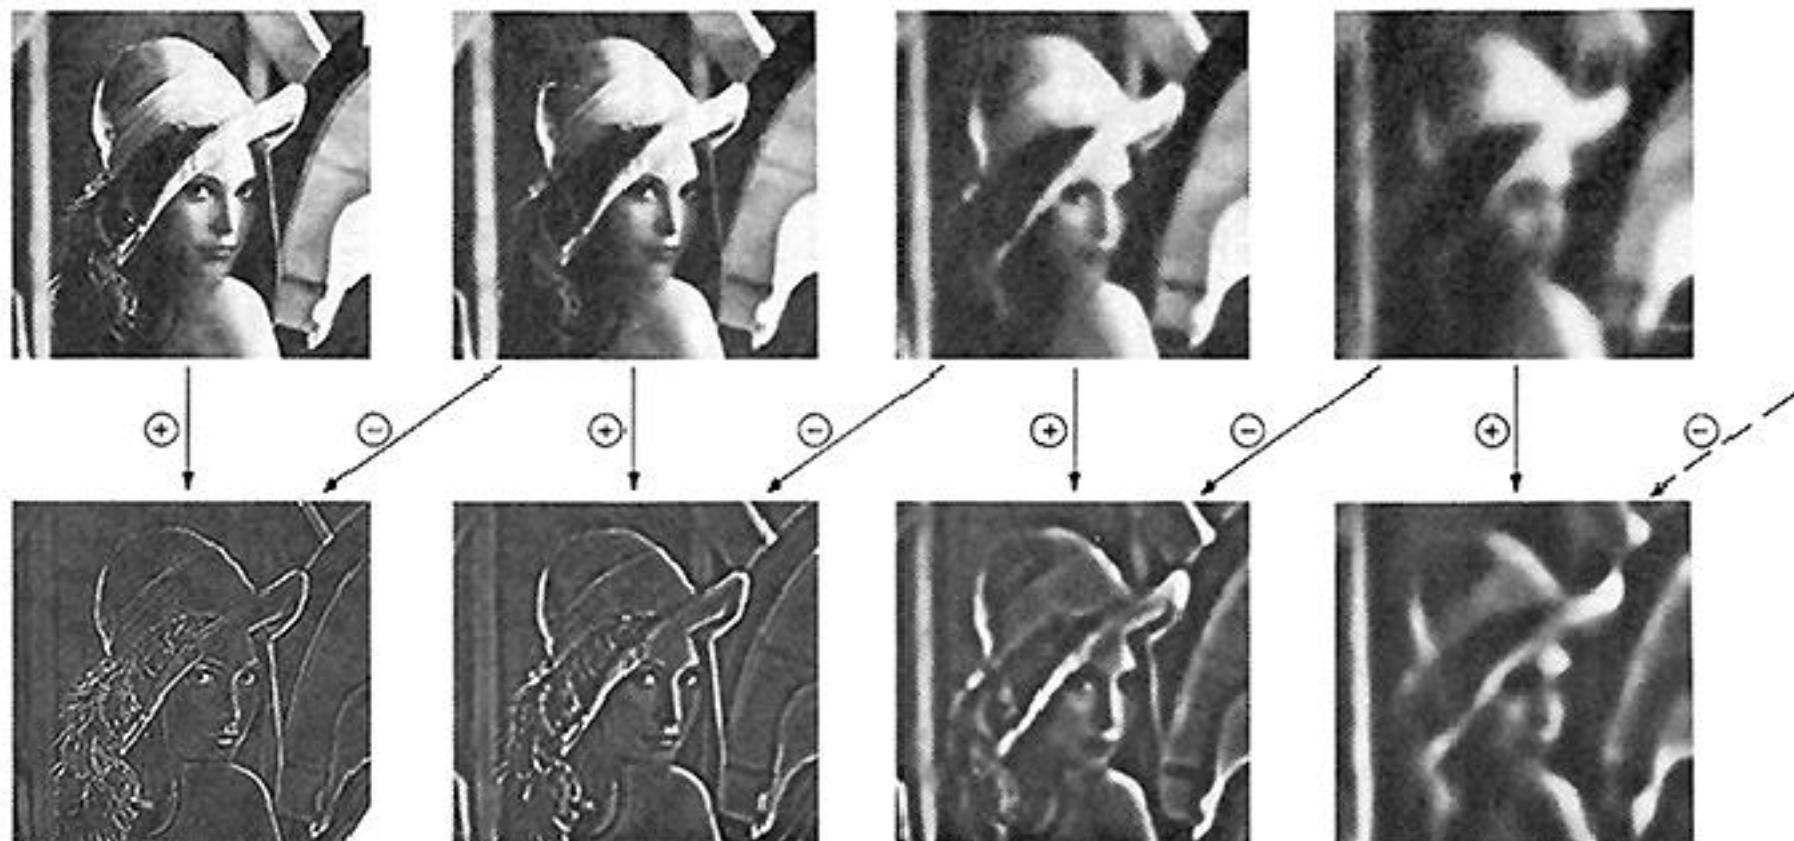


Fig. 5. First four levels of the Gaussian and Laplacian pyramid. Gaussian images, upper row, were obtained by expanding pyramid arrays (Fig. 4) through Gaussian interpolation. Each level of the Laplacian pyramid is the difference between the corresponding and next higher levels of the Gaussian pyramid.

# Laplacian pyramid reconstruction algorithm: recover $x_1$ from $L_1, L_2, L_3$ and $x_4$

$G\#$  is the blur-and-downsample operator at pyramid level  $\#$

$F\#$  is the blur-and-upsample operator at pyramid level  $\#$

Laplacian pyramid elements:

$$L_1 = (I - F_1 G_1) x_1$$

$$L_2 = (I - F_2 G_2) x_2$$

$$L_3 = (I - F_3 G_3) x_3$$

$$x_2 = G_1 x_1$$

$$x_3 = G_2 x_2$$

$$x_4 = G_3 x_3$$

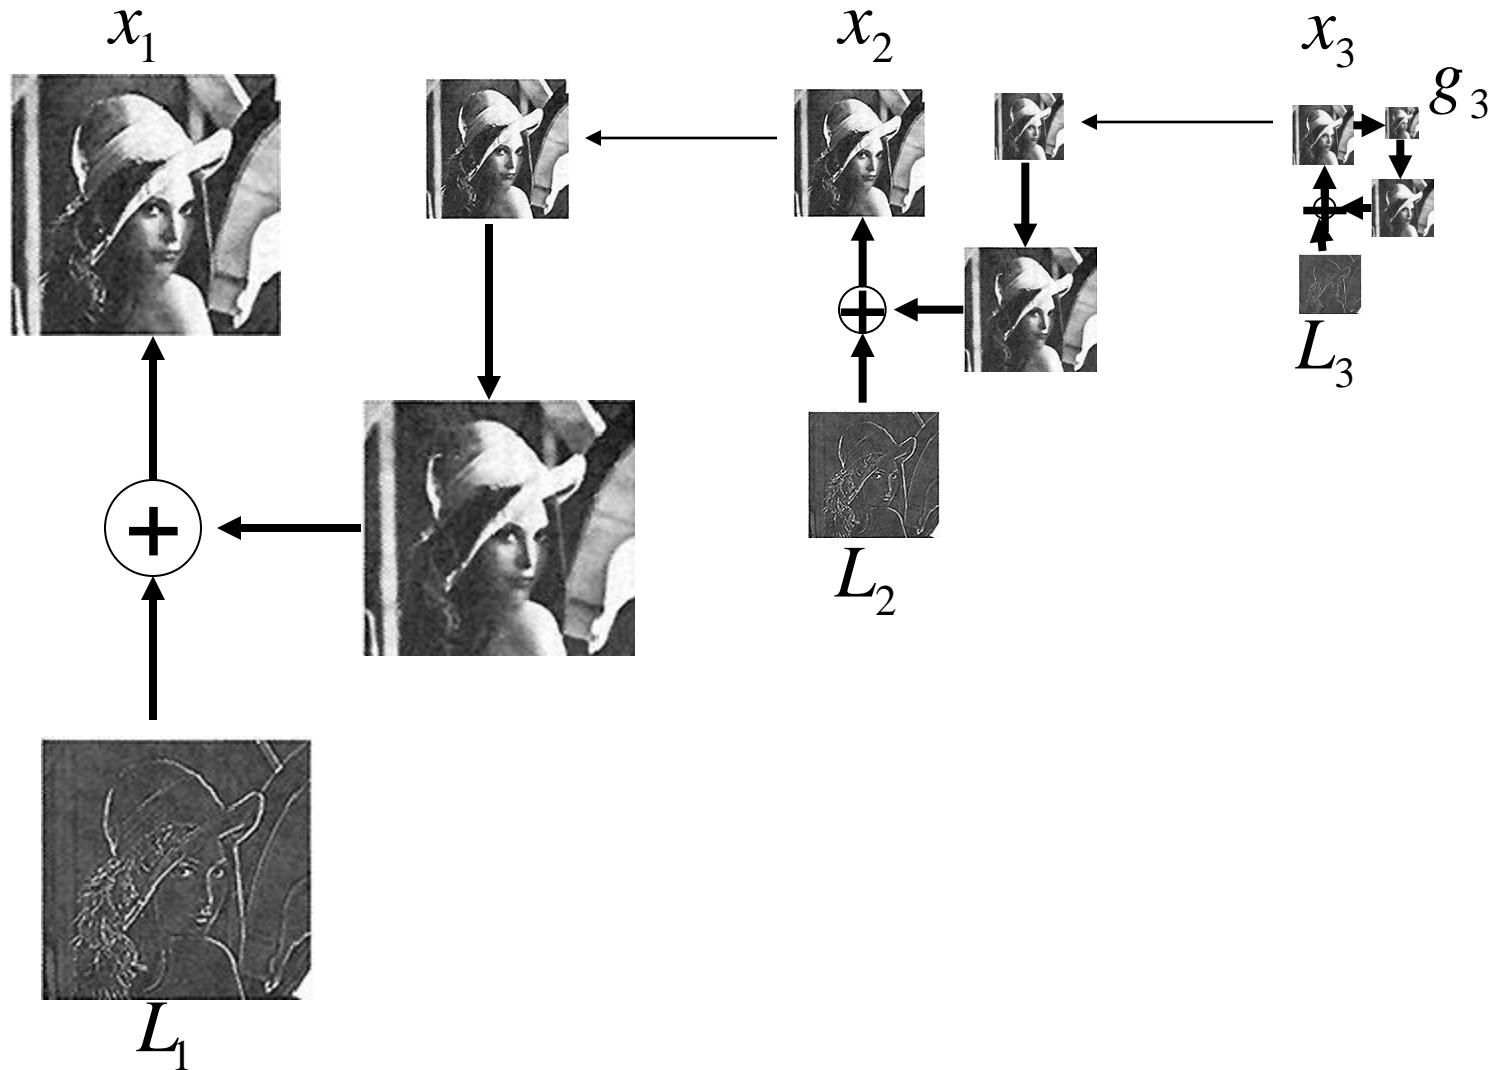
Reconstruction of original image ( $x_1$ ) from Laplacian pyramid elements:

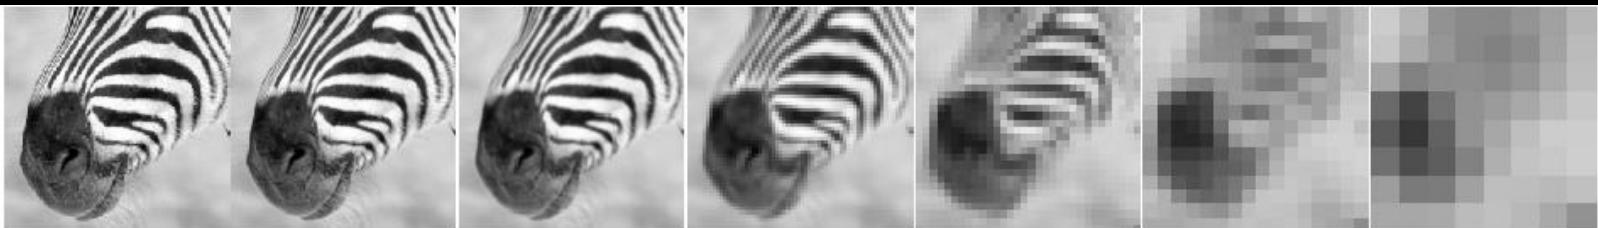
$$x_3 = L_3 + F_3 x_4$$

$$x_2 = L_2 + F_2 x_3$$

$$x_1 = L_1 + F_1 x_2$$

# Laplacian pyramid reconstruction algorithm: recover $x_1$ from $L_1$ , $L_2$ , $L_3$ and $g_3$





512

256

128

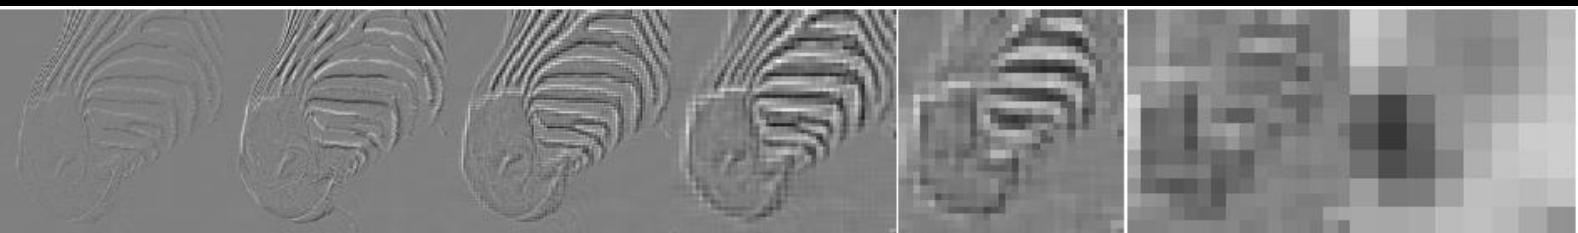
64

32

16

8





512

256

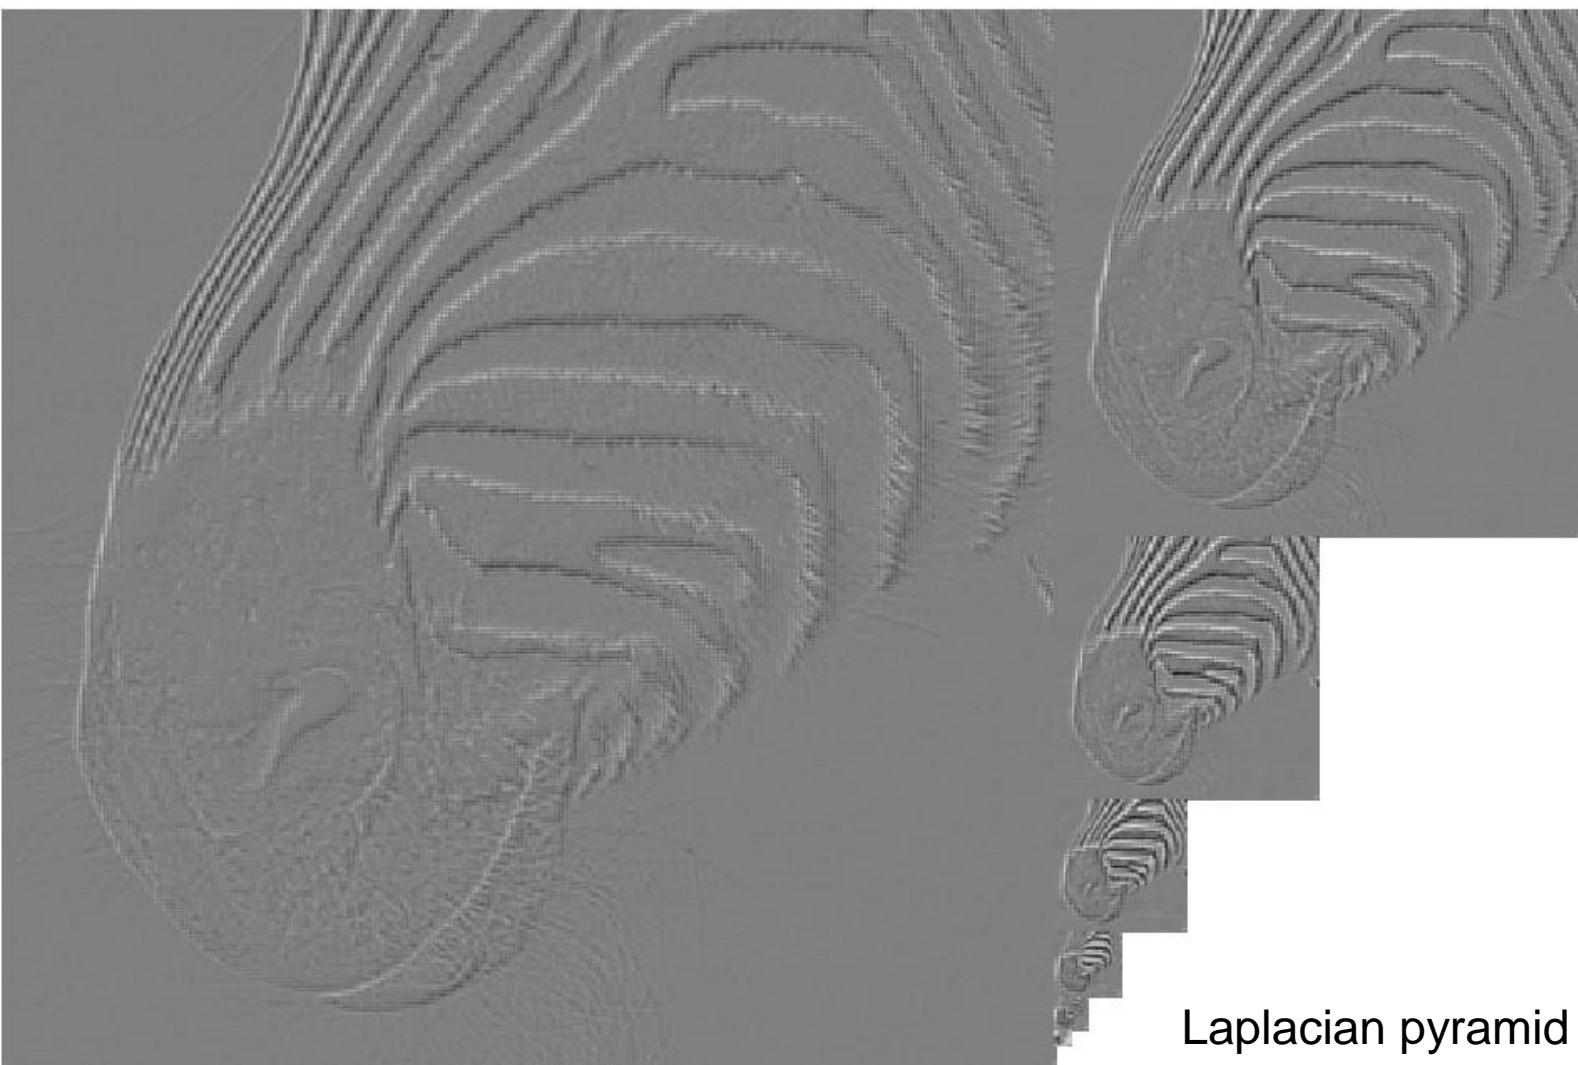
128

64

32

16

8



Laplacian pyramid

# Why use these representations?

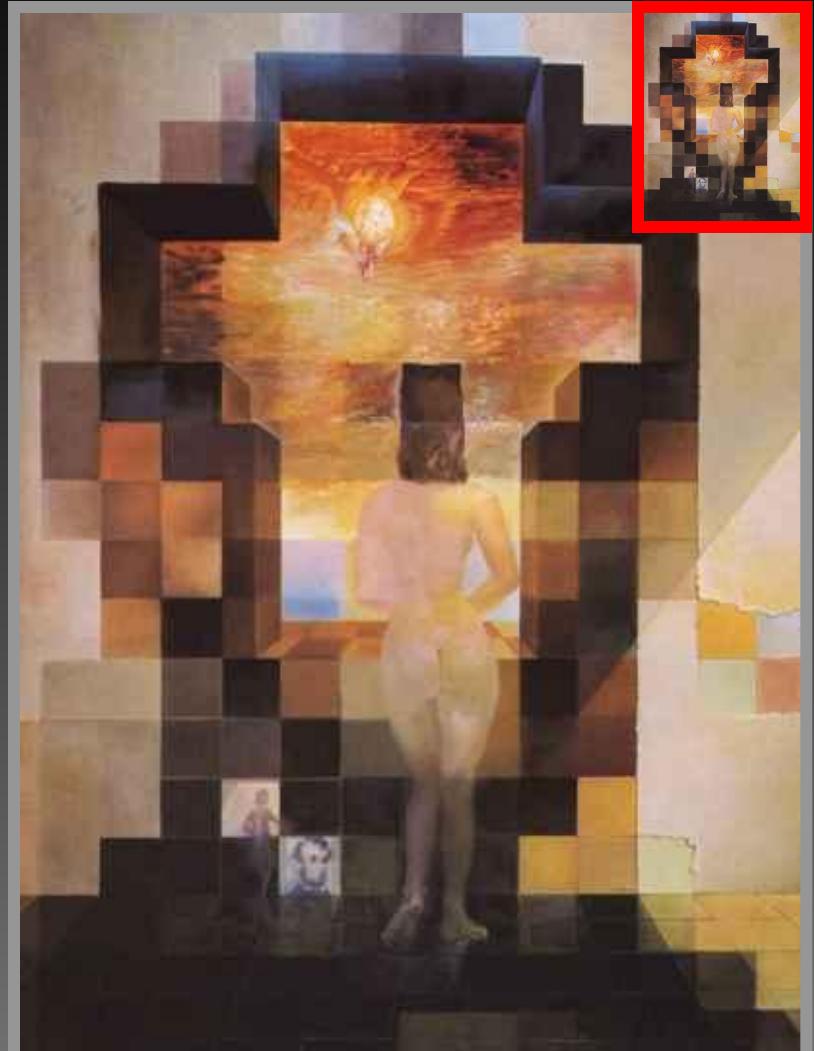
- Handle real-world size variations with a constant-size vision algorithm.
- Remove noise
- Analyze texture
- Recognize objects
- Label image features
- Image priors can be specified naturally in terms of wavelet pyramids.

# Visual Perception

For computer vision systems to “understand art” other properties of human visual perception must be taken into account

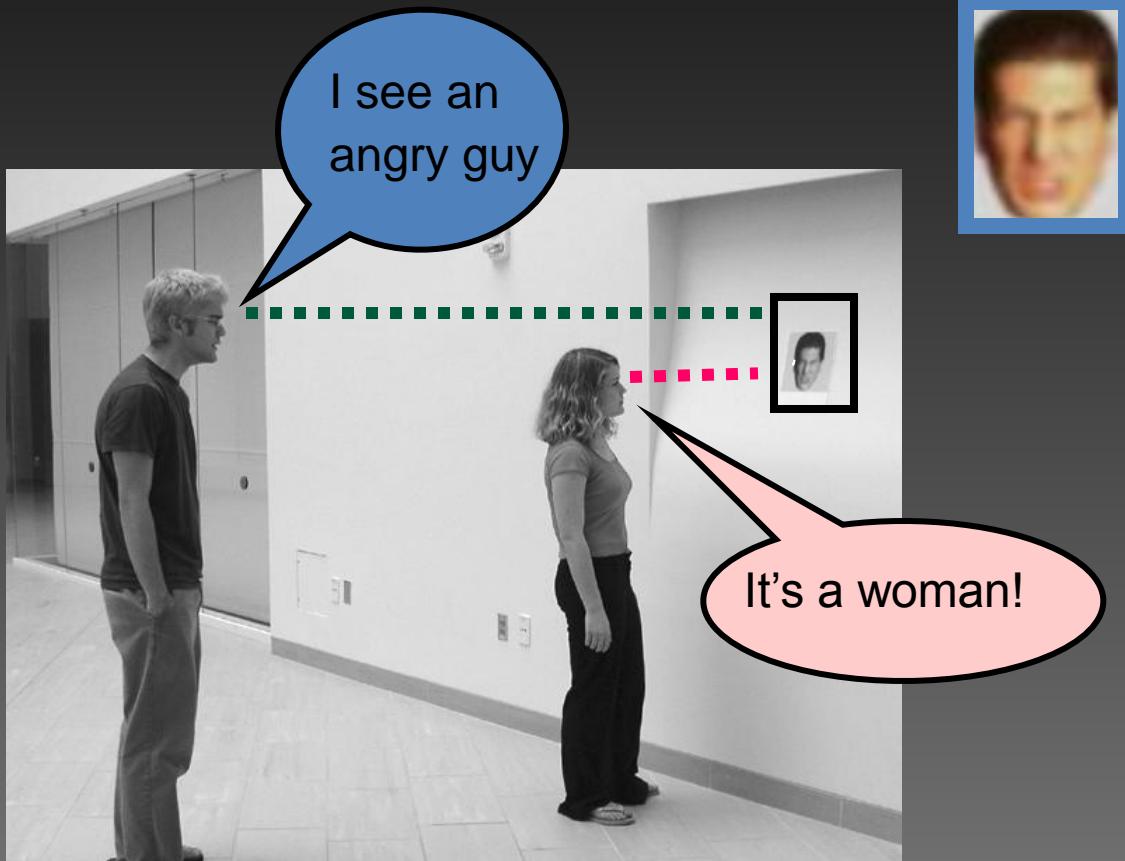


Salvador Dali



# Hybrid Images

What you see...



From Far Away



Up Close



# *Advances* in multi-resolution perception

- We tend to constantly reinterpret the representation if we are gaining information
- We tend to stick with our first grouping interpretation if we are losing information. The hysteresis effect is modulated by the strength of perceptual grouping across scales.

The visual system is designed to integrate across multi resolution in a way that provides **the most accurate interpretation** of the world as we move through it.

# Readings

- Chapter 4