


计算机视觉

Computer Vision

Lecture 8: Homography and Image Alignment

张 超

信息科学技术学院 智能科学系



Today


- **Image Alignment**
 - **Fitting a 2D transformation**
 - Affine, Homography
 - **Computing an image mosaic**
 - **2D image warping**
 - **Image Blending**

Planar Projective transformations

a.k.a. Homographies

$$\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} u \\ v \\ w \end{bmatrix} \quad \begin{aligned} x' &= u/w \\ y' &= v/w \end{aligned}$$

“keystone” distortions





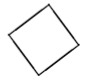
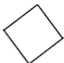

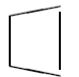
Finding the transformation



- How can we find the transformation between these images?
- How many corresponding points do we need to solve?

2D image transformations



Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation + ...	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths + ...	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles + ...	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism + ...	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

These transformations are a nested set of groups

- Closed under composition and inverse is a member

Finding the transformation


Translation	=	2 degrees of freedom
Similarity	=	4 degrees of freedom
Affine	=	6 degrees of freedom
Homography	=	8 degrees of freedom

How many corresponding points do we need to solve?

Homography


- How to relate two images from the same camera center?
 - how to map a pixel from PP1 to PP2?
- Think of it as a 2D image warp from one image to another.
- A projective transform is a mapping between any two PPs with the same center of projection
 - rectangle should map to arbitrary quadrilateral
 - parallel lines aren't
 - but must preserve straight lines
- called Homography

$$\begin{bmatrix} wx' \\ wy' \\ w \\ \mathbf{p}' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ \mathbf{H} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \\ \mathbf{p} \end{bmatrix}$$



Homography

camera rotation




Homography

camera A




camera B



Analysing patterns and shapes

What is the shape of the b/w floor pattern?



Homography




The floor (enlarged)



Automatically
rectified floor

Analysing patterns and shapes

Automatic rectification



From Martin Kemp *The Science of Art*
(*manual reconstruction*)

Image mosaics



How do we build panorama?

- We need to match (align) images




Matching with Features

- Detect feature points in both images



Matching with Features

- Detect feature points in both images
- Find corresponding pairs



Aligning the Images

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images




Building panorama

- Detect feature points in both images
- Find corresponding pairs
- Use these pairs to align images



Mosaics



...




image from S. Seitz


Obtain a wider angle view by combining multiple images.

How to stitch together a panorama (a.k.a. mosaic)?

- **Basic Procedure**
 - Take a sequence of images from the same position
 - Rotate the camera about its optical center
 - Compute transformation between second image and first
 - Transform the second image to overlap with the first
 - Blend the two together to create a mosaic
 - (If there are more images, repeat)
- ...but wait, why should this work at all?
 - What about the 3D geometry of the scene?
 - Why aren't we using it?


Pinhole camera

- Pinhole camera is a simple model to approximate imaging process, perspective projection.



If we treat pinhole as a point, only one ray from any given point can enter the camera.

Mosaics




...



image from S. Seitz



Obtain a wider angle view by combining multiple images.

A pencil of rays contains all views



Can generate any synthetic camera view
as long as it has **the same center of projection!**

Projecting images onto a common plane



each image is warped
with a homography **H**

Can't create a 360 panorama this way...

mosaic PP


Homography

- **Projective – mapping between any two projection planes with the same center of projection**
- **represented as 3x3 matrix in homogenous coordinates**

$$\begin{bmatrix} wx' \\ wy' \\ w, \\ \mathbf{p}' \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ \mathbf{H} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \\ \mathbf{p} \end{bmatrix}$$

To apply a homography \mathbf{H}

- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates (divide by w)



Projective Transformations aka Homographies aka Planar Perspective Maps

$$\mathbf{H} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & 1 \end{bmatrix}$$

Called a *homography*
(or *planar perspective map*)




Image warping with homographies







image plane in front



black area
where no pixel
maps to



Image rectification



To unwarp (rectify) an image

- solve for homography \mathbf{H} given \mathbf{p} and \mathbf{p}'
- solve equations of the form: $w\mathbf{p}' = \mathbf{H}\mathbf{p}$
 - linear in unknowns: w and coefficients of \mathbf{H}
 - \mathbf{H} is defined up to an arbitrary scale factor

Solving for homographies

$$\begin{bmatrix} {}^w x'_i \\ {}^w y'_i \\ w \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x'_i = \frac{h_{00}x_i + h_{01}y_i + h_{02}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$y'_i = \frac{h_{10}x_i + h_{11}y_i + h_{12}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$x'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{00}x_i + h_{01}y_i + h_{02}$$

$$y'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{10}x_i + h_{11}y_i + h_{12}$$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x'_i x_i & -x'_i y_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -y'_i x_i & -y'_i y_i & -y'_i \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Solving for homographies

$$\begin{array}{ccc}
 \begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x'_1 x_1 & -x'_1 y_1 & -x'_1 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1 x_1 & -y'_1 y_1 & -y'_1 \\ & & & & & \vdots & & & \\ x_n & y_n & 1 & 0 & 0 & 0 & -x'_n x_n & -x'_n y_n & -x'_n \\ 0 & 0 & 0 & x_n & y_n & 1 & -y'_n x_n & -y'_n y_n & -y'_n \end{bmatrix} & \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} & = \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix} \\
 \mathbf{A} & \mathbf{h} & \mathbf{0} \\
 2n \times 9 & 9 & 2n
 \end{array}$$

Defines a least squares problem:

$$\text{minimize } \|\mathbf{A}\mathbf{h} - \mathbf{0}\|^2$$

- Since \mathbf{h} is only defined up to scale, solve for unit vector $\hat{\mathbf{h}}$
- Solution: $\hat{\mathbf{h}}$ = eigenvector of $\mathbf{A}^T \mathbf{A}$ with smallest eigenvalue
- Works with 4 or more points

Fun with homographies

Original image



St.Petersburg
photo by A. Tikhonov

Virtual camera rotations




Image Alignment Algorithm

Given images A and B


- 1. Compute image features for A and B**
- 2. Match features between A and B**
- 3. Compute homography between A and B using least squares on set of matches**

What could go wrong?

RANSAC for Homography



RANSAC for Homography



RANSAC for Homography



Probabilistic model for verification






Image Warp


Image Warping

- Given a coordinate xform $(x',y') = T(x,y)$ and a source image $f(x,y)$, how do we compute an xformed image $g(x',y') = f(T(x,y))$?




Forward Warping

- Send each pixel $f(x)$ to its corresponding location $(x',y') = T(x,y)$ in $g(x',y')$
- What if pixel lands “between” two pixels?




Forward Warping

- Send each pixel $f(x,y)$ to its corresponding location $x' = h(x,y)$ in $g(x',y')$
 - What if pixel lands “between” two pixels?
 - Answer: add “contribution” to several pixels, normalize later (*splatting*)
 - Can still result in holes




Inverse Warping

- Get each pixel $g(x',y')$ from its corresponding location $(x,y) = T^{-1}(x',y')$ in $f(x,y)$
 - Requires taking the inverse of the transform
 - What if pixel comes from “between” two pixels?



Inverse Warping

- Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$
- What if pixel comes from “between” two pixels?
- Answer: *resample* color value from *interpolated* (*prefiltered*) source image



Interpolation

- Possible interpolation filters:

- nearest neighbor
- bilinear
- bicubic (interpolating)
- sinc

- Needed to prevent “jaggies”
and “texture crawl”

(with prefiltering)



Blending

Blending

- We've aligned the images – now what?






Blending


- **Want to seamlessly blend them together**






Image Blending





Feathering




Effect of window size



Effect of window size






Good window size




“Optimal” window: smooth but not ghosted

- Doesn't always work...


Pyramid blending



(d)



(h)




(l)

Create a Laplacian pyramid, blend each level

- Burt, P. J. and Adelson, E. H., [A multiresolution spline with applications to image mosaics](#), ACM Transactions on Graphics, 42(4), October 1983, 217-236.

Alpha Blending



Optional: see Blinn (CGA, 1994) for details:

<http://ieeexplore.ieee.org/iel1/38/7531/00310740.pdf?isNumber=7531&prod=JNL&arnumber=310740&arSt=83&ared=87&arAuthor=Blinn%2C+J.F.>

Encoding blend weights: $I(x,y) = (\alpha R, \alpha G, \alpha B, \alpha)$

color at $p = \frac{(\alpha_1 R_1, \alpha_1 G_1, \alpha_1 B_1) + (\alpha_2 R_2, \alpha_2 G_2, \alpha_2 B_2) + (\alpha_3 R_3, \alpha_3 G_3, \alpha_3 B_3)}{\alpha_1 + \alpha_2 + \alpha_3}$

Implement this in two steps:

1. accumulate: add up the (α premultiplied) $RGB\alpha$ values at each pixel
2. normalize: divide each pixel's accumulated RGB by its α value

Q: what if $\alpha = 0$?

Poisson Image Editing



sources/destinations



cloning



seamless cloning

- For more info: Perez et al, SIGGRAPH 2003

– http://research.microsoft.com/vision/cambridge/papers/perez_siggraph03.pdf

Readings

- **F & P 12.1.2 12.1.3**
- **Szeliski, *CVAA*:**
 - Chapter 3.5: Image warping
 - Chapter 6.1: 2D and 3D Feature-based alignment
 - Chapter 9.1: Motion models
 - Chapter 9.2: Global alignment
 - Chapter 9.3: Compositing
- **Recognizing Panoramas, Brown & Lowe, ICCV'2003**
- **Szeliski & Shum, SIGGRAPH'97**